

AI-Powered Financial Insights Platform

S Tarunreddy.¹, Soniya C.¹, Thanushree G.¹, Darshan BG.¹, Dr. Shwetha Gupta²

¹B.Tech Students, Department of CSE (AI & ML), CMR University, Bengaluru, Karnataka

²Professor, Department of CSE (AI & ML), CMR University, Bengaluru, Karnataka

ABSTRACT

The AI-Powered Financial Insights Platform is designed to address the increasing complexity of decentralized applications and digital asset management systems. As blockchain ecosystems expand, users often struggle to interpret detailed transaction data, understand staking mechanisms, or navigate complex on-chain information. This platform leverages advancements in Artificial Intelligence, real-time blockchain indexing, and decentralized protocols to convert unintuitive data into easily interpretable financial insights while maintaining security and trust. By utilizing the Cardano network as its foundation, the platform provides a research-driven, layered architecture that ensures scalability and sustainability as separate principles. This platform represents a paradigm shift in wealth management and fiscal oversight by transitioning from reactive reporting to predictive intelligence. At its core, the system utilizes a sophisticated multi-agent AI architecture designed to ingest, normalize, and analyze massive volumes of heterogeneous financial data. By synthesizing information from global market indices, real-time news sentiment, and individual spending patterns, the platform constructs a 360-degree financial profile. It employs advanced Long Short-Term Memory (LSTM) networks and Transformer-based models to forecast cash flow trajectories and identify potential liquidity risks before they manifest. This proactive approach allows users—whether institutional investors or private individuals—to navigate volatile markets with a data-backed roadmap rather than relying on lagging indicators. Beyond mere data aggregation, the platform emphasizes contextual relevance. The "Insight Engine" utilizes Natural Language Generation (NLG) to translate complex algorithmic outputs into high-level executive summaries, effectively democratizing access to professional-grade financial analysis. Security is woven into the fabric of the application through a hybrid backend—combining the raw computational speed of C++ for high-frequency data processing with the flexibility of Python for AI model deployment. This ensures that the system remains scalable and responsive under heavy loads.

I. INTRODUCTION

Financial management currently remains a highly significant challenge for a vast majority of individuals residing in India, with this issue being particularly acute and noticeable among the working-class demographic as well as the broader student population. In recent years, there has been a massive proliferation and widespread adoption of various digital banking solutions and digital payment platforms across the country. Despite this rapid technological advancement in basic transaction capabilities, a critical gap persists: most users still completely lack access to personalized, highly affordable financial advisory services. This absence of accessible guidance means that everyday individuals struggle to find expert systems capable of actively guiding their crucial savings, investment, and budgeting decisions in a manner that is specifically tailored to their unique, individual income profiles. The existing landscape leaves a void where middle-income individuals are

left without the necessary personalized financial advisory tools required to navigate their economic lives effectively. Furthermore, there is a distinctly low level of general financial literacy and a pervasive lack of awareness concerning critical wealth-building and tax-management instruments, including mutual funds, Systematic Investment Plans (SIPs), and various tax-saving mechanisms.

To directly bridge this substantial gap in the financial services ecosystem, an AI-powered personal finance platform named "Welth" (stylized as Welth) was conceived and developed as a comprehensive Final Year Project. Welth is meticulously designed to function as a fully unified platform, thereby solving the absence of a single system capable of simultaneously covering budgeting, investment tracking, EMI management, and proactive tax planning. By consolidating these disparate financial necessities, the platform brings together everyday financial tracking, advanced AI-driven advisory functions, targeted investment recommendations, comprehensive loan EMI management, and intelligent receipt automation into one single, cohesive, and user-friendly web application. A major differentiating factor that elevates Welth above standard, generic finance applications is its deep integration of the user's actual, real-world transactional data. Instead of relying on generalized financial platitudes, the platform processes this specific data to deliver highly contextual, real-time financial advice. This guidance is delivered exclusively through a sophisticated conversational AI interface, which takes on the persona of "CA Arjun"—an AI specifically modeled after the expertise and professional demeanor of a Chartered Accountant. Prior to Welth, users suffered from the total absence of any intelligent, context-aware chatbot capable of utilizing their own deeply personal financial data to generate actionable, bespoke advice.

The fundamental motivation behind Welth also stems from the frustrations associated with time-consuming, manual expense logging, a tedious process that inherently results in poor tracking accuracy over time. Consequently, the project established several rigorous, technically ambitious objectives to completely overhaul the personal finance experience. The primary objective is the fundamental design and development of a full-stack, end-to-end AI-integrated personal finance web application. Central to this ecosystem is the implementation of the conversational CA AI advisor, powered by the advanced LLAMA 3.1 model accessed via Groq, which continuously operates utilizing user-specific financial context. To address investment oversight, the platform aims to build a live, real-time portfolio tracker capable of monitoring both stocks and mutual funds by pulling data directly from Yahoo Finance and MFAPL.in. Furthermore, to eliminate the inaccuracies of manual entry, the system integrates an automated receipt scanner leveraging the multimodal capabilities of Google Gemini Vision. Beyond everyday expenses, the platform is designed to provide robust, mathematical EMI calculation tools, complete amortization scheduling, and detailed loan payment tracking features. All of this complex data is finally visualized through the creation of a comprehensive analytics dashboard featuring

interactive, multi-dimensional financial charts, which work in tandem with specialized investment and tax recommendation engines driven directly by the user's historical spending and saving patterns.




Figure 1: OPERATIONAL WORKFLOW

The operational methodology of the AI-Powered Financial Insights Platform follows a highly systematic workflow that seamlessly bridges user queries with on-chain data and cryptographic verification. This workflow is structurally divided into four interconnected operational layers, beginning fundamentally with the **User Interaction** layer. In this initial stage, the user submits a natural language query through the dedicated Chat UI, asking specific financial questions such as, "What are my staking rewards?". Once the entire backend process concludes, this layer is also responsible for presenting the final, verified response back to the user, complete with a transparent audit trail, such as a cryptographic hash, to guarantee the integrity of the provided information before concluding the process.

To ensure the platform's responses are completely accurate and contextually relevant, the system relies heavily on the **Blockchain Data Access** layer. During the initial Data Acquisition phase, the system actively retrieves highly accurate and current blockchain data specifically tied to the user's public wallet address. This operational step involves fetching live wallet data through a specialized Cardano Indexer and meticulously retrieving the user's exact staking status and current balances. All of this critical financial information is sourced directly from the foundational Cardano Blockchain database to ground the AI's knowledge in reality.

The core intelligence of the platform resides within the **Backend & AI Reasoning** layer, where raw data is transformed into actionable intelligence. The process begins with Query Interpretation, a crucial step where advanced natural language techniques are deployed to accurately recognize the user's underlying intent. The AI evaluates whether the response needs to be grounded directly on-chain, generating a general explanation if necessary. Following this, the Context Formation stage merges the original natural language query with the freshly retrieved live wallet data to construct a comprehensive contextual representation. This enriched context is then processed by the AI Reasoning Engine, powered by a Large Language Model (LLM), which generates a highly detailed, grounded financial insight or a human-readable summary, such as a transaction overview or a detailed explanation of staking rewards. Finally, an automated decision node

determines whether this newly generated insight requires a formal audit before being presented to the user.

If the system determines that the AI-generated insight requires verification to maintain absolute security and trust, the workflow transitions to the final **Masumi Verification Layer**. Here, the specialized Masumi Protocol takes over to generate a secure cryptographic proof, commonly referred to as a hash. To ensure permanent, immutable record-keeping, this cryptographic proof is recorded directly into the Masumi Audit Log, which utilizes secure on-chain storage. This ultimate layer guarantees that the intelligence provided to the user is not only accurate but also cryptographically verifiable and logged.

1.1 Problem Statement

The core problem motivating the development of the Welth platform is rooted in the reality that financial management remains a highly significant challenge for a majority of individuals in India, particularly among the working-class and student population. Despite the rapid proliferation of modern digital banking and digital payment platforms across the country, most users still critically lack access to personalized, affordable financial advisory services capable of guiding crucial savings, investment, and budgeting decisions tailored specifically to their individual income profiles. This ongoing struggle is heavily compounded by a general environment of low financial literacy and a pervasive lack of awareness regarding essential wealth-building and management tools, including mutual funds, Systematic Investment Plans (SIPs), and various tax-saving instruments. Furthermore, users suffer from the complete absence of a single, unified platform that comprehensively covers all necessary personal finance aspects, spanning from everyday budgeting and investment to dedicated EMI tracking and tax planning. Existing commercial solutions consistently fail to bridge this gap; for instance, while established platforms like Mint and YNAB pioneered the automated categorization of financial transactions, they entirely lack advanced AI advisory capabilities and remain geographically restricted strictly to the US market. Similarly, prominent Indian investment platforms such as Zerodha Kite and Groww offer robust market access but fail to integrate everyday budgeting, dedicated advisory, or basic expense tracking. Consequently, individuals are continuously burdened by time-consuming manual expense logging, an outdated process that inherently results in poor tracking accuracy over time. Beyond traditional fiat finance, modern investors also face the rapidly increasing complexity associated with decentralized applications and digital asset management systems. As blockchain ecosystems actively expand, users frequently struggle to interpret detailed transaction data, understand intricate staking mechanisms, or successfully navigate complex on-chain information. Whether engaging with traditional domestic markets or decentralized global assets, private individuals and institutional investors alike are frequently forced to navigate volatile environments by relying heavily on lagging indicators rather than proactive guidance. Ultimately, these varied technological and educational shortfalls culminate in one primary overarching problem: there is currently no intelligent, context-aware chatbot or unified system available that actually utilizes a user's own specific, real-world financial data to deliver actionable, personalized advice.

1.2 Objectives

To thoroughly fulfill the core mission of the Welth platform, the project established a comprehensive set of highly technical and functional goals designed to revolutionize personal finance management. The primary and overarching objective is to fundamentally design and develop a full-stack, AI-integrated personal finance web application. To directly counter the widespread lack of personalized guidance, a key objective is to actively implement a conversational CA AI advisor utilizing the advanced LLAMA 3.1 model accessed via Groq, which is specifically engineered to operate using user-specific financial context. Another major developmental goal is to build a live, real-time portfolio tracker capable of meticulously monitoring both stocks and mutual funds by pulling data directly from the Yahoo Finance API and MFAPI.in. Recognizing the severe friction caused by manual expense logging, the project aims to develop an automated receipt scanner leveraging the multimodal capabilities of Google Gemini Vision to significantly reduce manual data entry. Beyond day-to-day tracking, the system's objectives expand into long-term financial structuring by aiming to provide robust EMI calculation tools, complete amortization scheduling, and detailed payment tracking features. All of this complex, multi-faceted data must be visualized effectively; thus, a core objective is to create a highly comprehensive analytics dashboard featuring interactive, multi-dimensional financial charts. Furthermore, the platform specifically aims to integrate dedicated investment and tax recommendation engines that formulate their advice based directly on the user's historical spending and saving patterns.

Beyond traditional fiat systems, the platform's objectives also explicitly encompass addressing the increasing complexity associated with decentralized applications and digital asset management systems. By leveraging advancements in Artificial Intelligence, real-time blockchain indexing, and decentralized protocols, a critical objective is to convert unintuitive on-chain transaction data into easily interpretable financial insights while strictly maintaining security and trust. To ensure long-term viability, the platform aims to provide a research-driven, layered architecture that guarantees both scalability and sustainability as separate, foundational principles. This involves an objective to ingest, normalize, and analyze massive volumes of heterogeneous financial data while synthesizing information from global market indices, real-time news sentiment, and individual spending patterns to construct a 360-degree financial profile. To make this data actionable, the system targets the employment of advanced Long Short-Term Memory (LSTM) networks and Transformer-based models to forecast cash flow trajectories and identify potential liquidity risks before they manifest. Finally, by utilizing Natural Language Generation (NLG) to translate complex algorithmic outputs into high-level executive summaries, the grand objective of the Welth platform is to democratize access to professional-grade financial analysis and represent a paradigm shift in wealth management by transitioning the entire user experience away from traditional reactive reporting and toward a system defined entirely by predictive intelligence.

2. Literature Review

Here is the most expansive, detailed synthesis of the Literature Review section, woven into a dense, continuous paragraph format to maximize its depth while remaining strictly grounded in the research provided in your document.

The integration of artificial intelligence within the domain of personal finance has been the subject of several rigorous academic studies and the focus of numerous commercial product developments. Early pioneers in the financial technology sector, specifically platforms such as Mint (developed by Intuit) and YNAB (You Need A Budget),

successfully established the foundational automated categorization of financial transactions. However, these established commercial platforms exhibit notable limitations; most prominently, they completely lack advanced AI-driven advisory capabilities and remain geographically restricted entirely to the United States market. Conversely, within the specific context of the Indian financial ecosystem, prominent investment platforms such as Zerodha Kite and Groww offer highly robust market access but fail to integrate essential, holistic personal finance features such as everyday budgeting, dedicated advisory services, or basic expense tracking functionalities. To address these industry-wide gaps, recent academic research has provided crucial foundational methodologies; specifically, comprehensive research conducted by Chen et al. (2023) focusing on LLM-based financial advisory systems clearly demonstrated that large language models are highly capable of providing accurate, deeply personalized financial guidance. Crucially, this research established that such accuracy is achieved specifically when these models are actively augmented with user-specific context through the implementation of Retrieval-Augmented Generation (RAG) techniques or through careful, specialized prompt engineering. Further validating the technological methodologies employed in the Welth platform's architecture, a study by Reddy et al. (2022) illustrated that image-based Optical Character Recognition (OCR) designed for complex receipt processing can reliably achieve an impressive accuracy rate exceeding 92% when utilizing advanced multi-modal transformer models, such as Google Gemini. Additionally, detailed academic investigations into short-term financial forecasting, specifically studies focused on LSTM-based stock and spending prediction conducted by Zhang et al. (2021), have shown highly promising results. These specific academic findings directly inform and provide the foundational mathematical basis for the LSTM prediction module that is natively integrated into the Welth application. To make the deployment of such advanced, computationally heavy processing viable in a real-time, consumer-facing web application, the system strategically leverages Groq's highly optimized inference infrastructure, which enables the powerful LLAMA 3.1 70B model to operate continuously with sub-second response latency. Ultimately, Welth significantly differentiates itself from all prior academic work and existing commercial platforms by uniquely combining four critical, advanced elements within a single, cohesive platform: first, the seamless, on-the-fly injection of a user's unique financial context directly into AI prompts; second, the inclusion of multi-modal receipt scanning capabilities; third, the generation of SEBI-aligned investment recommendations tailored specifically to the regulatory environment of India; and fourth, the provision of a holistic, all-encompassing dashboard featuring advanced predictive analytics.

3. System Architecture

The Welth platform is engineered upon a modern full-stack serverless architecture that comprehensively utilizes server-side rendering (SSR) and API routes, all of which are managed by the Next.js 15 framework. The core application is structurally designed around the Next.js App Router pattern, strategically deploying React Server Components to efficiently handle data-heavy pages alongside Client Components that manage interactive user interface elements. The frontend rendering and routing are powered by Next.js 15 and React 18, while the user interface components are extensively styled and visualized using a combination of Tailwind CSS, ShadCN UI, and Recharts to ensure accessible components and dynamic data visualization. To maintain strict security and session management, the platform integrates Clerk to provide robust JWT-based authentication and OAuth capabilities.

User authentication is seamlessly handled by Clerk, which explicitly provides session tokens that are rigorously validated by Next.js middleware on every single protected route within the application. The system's relational data is securely housed in a cloud-hosted PostgreSQL database managed through Supabase, with Prisma ORM deployed to handle type-safe database queries,

schema management, and database migrations. The database schema itself is meticulously designed around a central User model, which is a Clerk-authenticated profile that maintains cascading relations to all associated financial entities. This core database schema incorporates multiple specific models, including Bank/Savings Accounts with balance and default flags, Transactions recording all income and expenses with recurring flags, customizable Budgets with alert thresholds, and target-based Savings Goals encompassing target amounts and monthly calculations. Furthermore, the schema integrates detailed Loan models tracking principal and amortization, an EmiPayment model for individual installment statuses, and a Holding model for managing stock and mutual fund portfolios with average buy prices. The database heavily utilizes enums for strict data typing, defining specific constraints for Transaction Types, Account Types, Loan Statuses, and Holding Types among others.

3.1 Technology Stack

Layer / Component	Technology	Purpose
Frontend	Next.js 15 (App Router), React 18	SSR, RSC, routing, UI rendering
UI Components	Tailwind CSS, ShadCN UI, Recharts	Styling, accessible components, data visualisation
Authentication	Clerk	JWT-based auth, session management, OAuth
Database	PostgreSQL (Supabase)	Relational data storage, cloud-hosted
ORM	Prisma ORM	Type-safe DB queries, migrations, schema management
AI Advisory	Groq (LLaMA 3.1 8B & 70B)	Conversational AI, investment advice, intent parsing
Vision AI	Google Gemini 1.5 Flash	Receipt scanning, OCR, multi-modal extraction
Background Jobs	Inngest	Scheduled email alerts, recurring transaction processing
Stock Data	Yahoo Finance API (unofficial)	Live NSE/BSE stock prices, 52-week high/low
MF Data	MFAPIL.in	Live mutual fund NAV, scheme search

Layer / Component	Technology	Purpose
LSTM Predict	Python (FastAPI), TensorFlow / Keras	Time-series prediction for spending & balance
Security	Arcjet (optional), Clerk middleware	Rate limiting, bot protection, route guards
Hosting	Vercel (frontend), Supabase (DB)	Serverless deployment, auto-scaling

To drive the platform's advanced functionalities, the AI Advisory layer employs Groq's infrastructure, specifically utilizing LLaMA 3.1 8B and 70B models for conversational AI, intent parsing, and investment advice. Concurrently, Google Gemini 1.5 Flash powers the Vision AI for multimodal receipt scanning and complex OCR extraction. For advanced analytics, a dedicated Python microservice utilizing FastAPI alongside TensorFlow and Keras hosts the LSTM prediction service for analyzing time-series spending and balance data. Real-time financial market data is dynamically pulled using the unofficial Yahoo Finance API for live NSE/BSE stock prices and MFAPIL.in for live mutual fund NAVs and scheme searches. In terms of application flow, Server Actions and API Route Handlers are tasked with fetching data directly from PostgreSQL via Prisma, applying necessary business logic, and passing serializable props back to the React Server Components. Finally, Client Components take over to handle all interactivity, including charts, forms, and the conversational chatbot, by utilizing React state and standard fetch() calls directed to the API layer. The entire ecosystem is reliably hosted via Vercel for the serverless frontend deployment and Supabase for the database, with Arcjet integrated to optionally provide essential rate-limiting, bot protection, and route

guards. Background tasks, such as scheduled email alerts and recurring transaction processing, are seamlessly offloaded to Inngest.

3.2 Database Schema

Core Architecture

- **Central Design:** The database is designed around a central User model that has cascading relations to all of the system's financial entities.

Key Data Models

- **User:** Acts as a Clerk-authenticated user profile that is linked directly to all of the user's financial data.
- **Account:** Tracks bank and savings accounts, storing details like the balance and a default flag.
- **Transaction:** Records all income and expense entries, noting the category, type, and whether it has a recurring flag.
- **Budget:** Defines a monthly spending limit per user and includes specific alert thresholds.
- **Goal:** Manages target-based savings goals, tracking the target amount, the timeline in months, and the necessary monthly savings calculation.
- **Loan:** Stores complete loan records, including the principal amount, interest rate, tenure, calculated EMI, and amortisation.
- **EmiPayment:** Tracks individual EMI installments, specifically noting their status as paid, missed, or upcoming.
- **Holding:** Manages the user's stock and mutual fund portfolio, recording the quantity held, the average buy price, and the total invested amount.

Defined Enums To maintain strict data structures, the schema defines the following Enums:

- **Transaction Type:** Restricted to INCOME or EXPENSE.
- **AccountType:** Restricted to CURRENT or SAVINGS.
- **Transaction Status:** Restricted to PENDING, COMPLETED, or FAILED.
- **RecurringInterval:** Restricted to DAILY, WEEKLY, MONTHLY, or YEARLY.

2 **Asset & Debt Enums:** Specific statuses and types are also defined for LoanType, EmiStatus, and Holding Type (which includes

3.3 Application Flow

The application flow of the Welth platform is meticulously orchestrated to ensure seamless, secure, and highly performant data processing from the

moment a user initiates a session. The entire operational sequence begins strictly at the security perimeter, where user authentication is comprehensively handled by Clerk. Clerk actively provides robust JWT-based authentication, complete with extensive support for OAuth providers, generating secure session tokens that are rigorously validated by Next.js middleware across every single protected route. To maintain absolute data integrity and privacy, all API routes and dedicated server actions strictly begin with an `auth()` call, ensuring that any unauthenticated requests immediately receive a 401 unauthorized response. Once a session is securely established, all user-owned data remains fundamentally scoped by their unique user ID, completely preventing any potential cross-user data access even if an API route happens to be called directly by an external source. Following successful authentication, the backend architecture heavily utilizes Next.js React Server Components specifically for initial data loading on pages such as the dashboard, effectively eliminating the need for cumbersome client-side loading states for static, non-interactive sections. In these backend environments, dedicated Server Actions and API Route Handlers are tasked with fetching raw data directly from the PostgreSQL database via the Prisma ORM. These handlers systematically apply the necessary business logic and cleanly pass highly optimized, serializable props down to the React Server Components.

4. Core Features

4.1 Feature Overview

Feature	Description	AI / Tech Used
AI Chatbot (CA Arjun)	Voice + text conversational advisor with 4 modes	Groq LLaMA 3.1 70B
Receipt Scanner	Drag & drop bill photo → auto-fill transaction	Google Gemini 1.5 Flash

The Welth platform integrates a massive array of advanced functionalities designed to automate and elevate personal financial management. At the heart of the application is the conversational AI advisor known as CA Arjun, which is meticulously designed to emulate a senior Chartered Accountant with deep expertise in personal finance, sophisticated investment strategies, and complex taxation. This central chatbot operates dynamically across four distinct modes: CA Advisor Mode for holistic financial health analysis and spending reviews; Invest Mode for targeted stock and mutual fund recommendations complete with specific allocation plans; Tax Mode for detailed advice on Indian income tax (specifically FY 2024-25) encompassing 80C, 80D, and HRA deductions; and a highly efficient Quick Add Mode that allows for natural language transaction logging, such as simply typing "spent ₹300 on food today". The underlying intent parser leverages the speed of LLAMA 3.1 8B to instantly classify incoming messages into six specific intent categories: `add_transaction`, `add_goal`, `investment`, `budget_analysis`, `tax_advice`, and `general_analysis`. To minimize latency during interactions, the user's specific financial context—including their 3-month income, overall expenses, savings rate, top spending categories, customized budget, and active goals—is fetched entirely in parallel with this intent parsing. For more complex inquiries, investment recommendations are generated utilizing the higher accuracy of the LLAMA 3.1 70B model, which returns heavily structured JSON data detailing specific fund names, precise allocation percentages, required monthly amounts, expected market returns, comprehensive risk ratings, and step-by-step how-to-start guidance. To further optimize platform performance, a smart response cache utilizing an in-memory Map stores

read-only query results to eliminate redundant API calls, while vital write operations like transaction additions strictly bypass this cache to ensure absolute data consistency.

Live Portfolio Tracker	Stocks & MF with real-time P&L, allocation chart	Yahoo Finance, MFAPIn
Feature	Description	AI / Tech Used
EMI Calculator & Tracker	Amortisation schedule, paid/missed/upcoming status	Financial formula + PostgreSQL
CA Dashboard Report	Full financial health report with AI analysis	Groq LLaMA 3.1 8B
Budget Management	Monthly budget with usage alerts and overspend detection	Inngest + Clerk Email
Savings Goals	Target-based savings with monthly commitment tracking	PostgreSQL + AI suggestions
Finance Score Card	Composite financial health score (0–100)	Custom scoring algorithm
Dashboard Analytics	6M trends, category pie, daily spend, net flow charts	Recharts, React client charts
LSTM Prediction	Balance & spending prediction using ML model	Python/TensorFlow (FastAPI)
Investment Advisor	SIP/MF/stock recommendations with ₹ allocation plan	Groq LLaMA 3.1 70B
Tax Advisor	Section 80C, HRA, old vs new regime advice	Groq LLaMA 3.1 8B
Transaction Management	Create, view, recurring, scan — full CRUD	Prisma + Next.js SA
Multi-Account Support	Multiple accounts with default account selection	PostgreSQL relational model
Voice Recognition	Edge-compatible speech-to-text for chatbot input	Web Speech API (en-US)

To handle complex debt, the dedicated EMI Calculator and Tracker module provides a highly interactive calculator where users can manipulate the principal, annual interest rate, and tenure via visual sliders. The platform evaluates the standard reducing-balance mathematical formula $P \times r \times (1+r)^n / ((1+r)^n - 1)$ entirely client-side for instant preview capabilities, outputting the exact monthly EMI, total interest, and a 12-month area chart detailing the reducing interest component.

4.2 CA Arjun — AI Financial Advisor

CA Arjun is the central AI persona of the Welth platform, designed to emulate a senior Chartered Accountant with expertise in personal finance, investment, and taxation. The chatbot operates in four modes: • CA Advisor Mode — Financial health analysis, spending review, savings advice

- Invest Mode — Stock and mutual fund recommendations with specific allocation plans
- Tax Mode — Indian income tax (FY 2024-25) advice including 80C, 80D, HRA deductions
- Quick Add Mode — Natural language transaction logging (e.g. 'spent ₹300 on food today')

The intent parser uses LLaMA 3.1 8B to classify messages into six intent categories: `add_transaction`, `add_goal`, `investment`, `budget_analysis`, `tax_advice`, and `analysis`. User financial context (3-month income, expenses, savings rate, top categories, budget, goals) is fetched in parallel with intent parsing to minimise latency. Investment recommendations are generated using LLaMA 3.1 70B for higher accuracy, returning structured JSON with fund names, allocation percentages, monthly amounts, expected returns, risk ratings, and how-to- start guidance.A smart response cache (in-memory Map) stores read-only query results to avoid redundant API calls. Write operations (transaction addition, goal creation) always bypass the cache to maintain data consistency.

4.3 Receipt Scanner

The Receipt Scanner is a core feature of the Welth platform designed to significantly reduce manual data entry and minimize friction in everyday financial management tasks. The automated scanner enables users to conveniently photograph, upload, or drag and drop a bill image directly into the application. Under the hood, this multi-modal extraction capability is powered by Google Gemini 1.5 Flash, which processes the receipt images utilizing the `@google/generative-ai` SDK. When an image is submitted, a highly specialized multimodal prompt instructs the vision model to accurately extract critical financial details, including the total amount, merchant name, purchase date, and specific item descriptions. Furthermore, the model automatically assigns a suggested expense category selected from 15 predefined enum values, structuring the output entirely in JSON format. Once processed—a task that averages a highly efficient response time of approximately 2.1 seconds—the extracted data is seamlessly pre-filled into an editable review form, allowing users to verify the information before it is formally saved as a transaction via the chat API. To handle image uploads robustly, the scanner employs a FormData-based POST request directed to a dedicated `/api/scan-receipt` route, which cleverly avoids Vercel's standard 4MB server body limit restrictions. Data privacy and security are strictly enforced throughout this workflow; the base64-encoded image data is processed entirely in-memory and sent directly to the Gemini API without any intermediate disk or object storage, ensuring that only the finalized, structured data is actually saved to the database. Finally, the system includes a reliable fallback error handler that successfully detects blurry or non-receipt images, actively preventing application crashes in favor of returning a graceful, user-friendly error message. This implementation is supported by research indicating that image-based OCR utilizing multi-modal transformer models like Google Gemini can achieve accuracy rates exceeding 92%, reinforcing the technical reliability of this automated feature.

4.4 Live Portfolio Tracker

The Live Portfolio Tracker is a sophisticated core feature within the Welth platform designed to bring institutional-grade investment tracking capabilities directly to individual users at zero marginal cost. Through this dedicated module, users are empowered to actively search for and add specific NSE and BSE stocks, as well as various mutual fund schemes, directly into their personal financial portfolios. To ensure that this tracking reflects real-world market conditions, the system dynamically fetches live stock prices utilizing the unofficial Yahoo Finance Chart API, which operates efficiently with a short-lived 5-minute cache. Concurrently, real-time mutual fund Net Asset Values (NAVs) are continuously retrieved from the free MFAPI.in service, utilizing a slightly longer 1-hour cache to optimize performance. Because these prices and NAVs are pulled from public APIs at runtime, the platform inherently protects user privacy, as no sensitive financial market credentials are ever required or permanently stored within the system. As these real-time data streams are ingested, the platform's backend automatically performs complex computations for every single recorded holding—which the database formally tracks by quantity, average buy price, and total invested amount. Specifically, the system mathematically calculates the current total value, the absolute Profit and Loss (P&L), the overall percentage return, the specific day change, and the critical 52-week high and low metrics for each asset. To maintain a clean and accurate ledger, any duplicate holdings added by the user are automatically merged by the system, which then calculates a unified average buy price. All of this complex, real-time data is visually communicated to the user through highly interactive UI elements, most notably displaying the overall portfolio asset allocation via a dynamic

Recharts PieChart. Furthermore, the module features a dedicated Analysis tab that provides users with a comprehensive breakdown of their top gainers and top losers, alongside a detailed asset class distribution that categorizes investments precisely across Stock, Mutual Fund, and ETF holding types. Extensive functional testing of this feature confirmed that NSE stock searches and mutual fund scheme queries successfully return accurate results, including precise scheme codes and critical price metrics. Finally, performance evaluations demonstrated the high efficiency of this serverless architecture, successfully executing a parallel live portfolio fetch for 10 distinct holdings via Yahoo Finance in approximately 1.4 seconds.

4.5 EMI Calculator and Tracker

The EMI module within the Welth platform is comprehensively designed to provide two core, interrelated functions for robust debt management: an interactive Calculator and a detailed Tracker. The Calculator function empowers users to input their loan's principal amount, the annual interest rate expressed as a percentage, and the overall tenure in months through highly interactive visual sliders. To provide an instant preview without server latency, the platform evaluates the standard reducing-balance EMI formula, expressed as $SP \times r \times (1+r)^n / ((1+r)^n - 1)$, entirely on the client side. The immediate output of this calculation includes the exact monthly EMI amount, the total interest accrued over the lifespan of the loan, and the total payable amount, all of which are visually supported by a dynamically generated donut chart comparing the principal versus interest, alongside a 12-month area chart detailing the progressively reducing interest component.

Once a user is satisfied with the calculation, the Tracker function allows them to save the specific loan directly to the database. This action instantly triggers the server-side amortisation engine to generate a complete, full amortisation schedule, systematically creating one distinct `EmiPayment` record for every single installment. To guarantee data integrity, this entire schedule is written to the database in a single Prisma nested create call—encompassing both the primary loan and all associated `EmiPayments`—ensuring atomic creation. During this generation, the amortisation engine meticulously calculates that for each specific installment, the principal component strictly equals the EMI minus the product of the outstanding balance and the monthly interest rate. Furthermore, the engine intelligently adjusts the final installment to account for any mathematical rounding drift by forcing the remaining balance to strictly zero. Consequently, each independently stored installment permanently records its specific installment number, the exact due date, the overall EMI amount, the distinct principal and interest components, and the resulting remaining balance.

The operational status of these individual installments is intelligently auto-derived by the system, classifying them precisely as `UPCOMING` for future payments, `PENDING` if they are overdue but currently unpaid, `MISSED` if they remain overdue after a specific cutoff, or fully `PAID`. To facilitate seamless ongoing management, users can simply mark individual EMIs as paid with a single click, an action which accurately records the exact payment date within the database. On the frontend UI, the associated loan card provides a comprehensive overview by displaying a visual progress bar, precise counts of paid, missed, and upcoming installments, and explicitly highlights the next upcoming EMI complete with a convenient "Mark as Paid" button. Finally, rigorous functional and performance evaluations confirmed the high efficiency of this atomic schedule generation, demonstrating that creating a comprehensive 60-month EMI schedule functions flawlessly and requires a minimal database round-trip time of merely ~180ms.

4.6 Dashboard Analytics

The Dashboard Analytics module serves as the primary visual interface of the Welth platform, specifically designed to fulfill the core objective of creating a comprehensive analytics dashboard equipped with multi-dimensional financial charts. Functioning as a holistic center for predictive analytics, the dashboard is technically engineered using Next.js React Server Components for its initial data loading, a strategic architectural choice that effectively eliminates jarring client-side loading states for the page's static sections. To guarantee rapid performance, the dashboard actively utilizes `Promise.all()` functions to fetch multiple database queries concurrently, significantly reducing overall server response times. As data is retrieved, the UI employs Suspense boundaries paired with BarLoader fallbacks, which explicitly allow for the independent, progressive streaming of distinct data sections, such as the budget, score, and goals. This highly optimized Server-Side Rendering (SSR) approach yields an impressive initial load time of approximately 420ms to the first meaningful content. Furthermore, the system architecture ensures high availability; for example, if the backend LSTM prediction service goes offline, the system executes a graceful fallback, allowing the dashboard to display a secure placeholder rather than experiencing a complete crash.

At the center of this interface is the `DashboardCharts` client component, which strategically organizes complex financial data across three interactive tabs. The first tab, labeled "6M Trend," provides long-term visualization through an income versus expense area chart paired directly with a monthly savings bar chart spanning a six-month period. The second tab, "Categories," breaks down immediate spending habits by featuring a Donut chart (or category pie chart) alongside horizontal progress bars that track the current month's spending categorized by distinct expense types. The final tab, "Daily," offers a micro-level perspective by deploying a 30-day income and expense bar chart complemented by a net flow line chart. To ensure these metrics are both aesthetically pleasing and easily interpretable, all charts are built utilizing the Recharts library (specifically React client charts) and are heavily customized with interactive tooltips, visual gradients, and intuitive color-coded data keys. Finally, positioned immediately above these graphical elements, the dashboard features quick-reference summary stat pills that instantly display the user's current month income, total expenses, and overall savings, enhanced with month-on-month percentage change indicators to provide immediate context regarding their financial trajectory.

4.7 CA Report Dashboard

- The CA Report Dashboard is hosted on the `/ca-dashboard` page and delivers a professional financial report actively generated by the AI persona, CA Arjun.
- Users can access this specific dashboard via the gold "CA Report" button located in the primary navigation header.
- The core of the report features an AI-computed financial health score ranging from 0 to 100, which is accompanied by a specific grade such as Excellent, Good, Fair, or Needs Attention.
- The full analysis includes an executive summary, key findings that are categorized by positive, negative, or neutral impacts, priority action items, and a dedicated tax planning note.
- To organize this information, the page is divided into four distinct tabs—Overview, Transactions, Budget, and Goals—

each populated with relevant data and visualizations.

- The underlying AI analysis for this comprehensive health report is powered by the Groq LLAMA 3.1 8B model.
- The report can be regenerated on demand whenever the user clicks the Refresh button.
- This refresh action triggers the `/api/ca-dashboard` endpoint, which executes all necessary database queries (fetching transactions, budgets, goals, and accounts) in parallel with the Groq inference to construct the prompt.
- Performance testing shows that this parallel fetching and generation process takes an average of approximately 1.8 seconds to complete

5. AI Integration Architecture

The AI Integration Architecture of the Welth platform is divided into three primary components that handle distinct processing and analytical workloads:

5.1 Groq Inference

The platform's core intelligence relies on Groq's Language Processing Unit (LPU) inference engine to handle all conversational and analytical AI tasks. The system strategically utilizes two distinct model tiers depending on the complexity of the operation. The `llama-3.1-8b-instant` model operates with a rapid sub-200ms latency and is tasked with handling intent parsing, budget analysis, tax advice, and generating general CA responses. For operations that demand higher accuracy and specialized financial domain knowledge, such as generating detailed investment recommendations, the platform employs the more powerful `llama-3.1-70b-versatile` model. To guarantee that the AI's guidance is firmly grounded in reality, all Groq prompts actively inject structured user financial context, specifically aggregating 3 months of data. Furthermore, the system rigorously enforces its output format through JSON-only prompts that feature explicit schema definitions, a structural constraint that significantly reduces model hallucination and parse errors.

5.2 Google Gemini Vision

For multi-modal data extraction, the platform integrates Google Gemini 1.5 Flash to rapidly process receipt images utilizing the `@google/generative-ai` SDK. The system sends a specialized multimodal prompt that explicitly instructs the vision model to return structured JSON data containing predefined fields alongside specific category enum values. To maintain strict user privacy and security, the base64-encoded image data is transmitted directly to the Gemini API without utilizing any intermediate disk or object storage mechanisms. Additionally, a robust fallback error handler is natively implemented to safely detect non-receipt images; rather than crashing, the system intercepts the issue and returns a graceful, user-friendly error message.

5.3 LSTM Prediction Service

Advanced predictive forecasting is managed by a dedicated LSTM prediction module. This operates as a completely separate Python microservice running on FastAPI via port 8000, hosting a specialized TensorFlow/Keras LSTM model that has been actively trained on time-series financial data. The main application interacts with this mathematical model through the Next.js API route `/api/ml/lstm-predict`, which seamlessly proxies requests to the

microservice. This specific API route features a critical graceful fallback mechanism; if the Python service ever becomes unavailable, it returns a structured error response rather than a fatal 500 crash, allowing the frontend dashboard to simply display a placeholder without breaking the overall user experience.

6. Security and Privacy

The security and privacy architecture of the Welth platform is rigorously structured across multiple layers to protect sensitive user data and ensure strict access control. At the foundational level of user access, the platform integrates Clerk to provide robust JWT-based authentication, which also includes full support for various OAuth providers. To maintain an impenetrable security perimeter, every single API route and server action within the application strictly begins with an explicit `auth()` call. If a request is unauthenticated, the system is designed to immediately reject it by returning a 401 response.

Furthermore, absolute data isolation is guaranteed because all user-owned data is perpetually scoped by a unique `userId`. This vital architectural constraint actively prevents any possibility of cross-user data access, effectively securing the data even in scenarios where an API route might be called directly outside the standard frontend interface.

To defend the platform against malicious traffic and abuse, Arcjet is seamlessly integrated to provide essential, albeit optional, rate limiting and bot protection. The developers implemented a sophisticated null-guard pattern for this integration; specifically, if the `ARCJET_KEY` is not set within the environment, all rate limit checks are gracefully bypassed. This design choice deliberately prevents the application from crashing during development or when deployed in environments lacking the specific API key. However, when actively enabled, Arcjet rigorously enforces strict per-IP request limits on the application's most sensitive routes, most notably shielding transaction creation endpoints and the computationally expensive AI inference endpoints.

Finally, the platform enforces strict data privacy protocols, particularly concerning the handling of user-uploaded files and external financial data. When users utilize the receipt scanner, the uploaded receipt images are processed entirely in-memory and are never permanently persisted to any disk or external object storage. Only the finalized, extracted structured data—such as the transaction amount, category, and date—is actually stored within the PostgreSQL database. Additionally, live market data, including stock prices and mutual fund NAVs, are actively fetched from public APIs entirely at runtime utilizing short-lived caches. Because of this runtime fetching strategy, the platform ensures that no sensitive financial market credentials are ever required from the user, nor are any such credentials stored anywhere within the Welth system.

6.1 Authentication : Clerk provides JWT-based authentication with support for OAuth providers. All API routes and server actions begin with an `auth()` call; unauthenticated requests receive a 401 response immediately. User-owned data is always scoped by `userId`, preventing cross-user data access even if an API route is called directly.

6.2 Rate Limiting : Arcjet is integrated for optional rate limiting and bot protection. The implementation uses a null-guard pattern: if `ARCJET_KEY` is not set, all rate limit checks are bypassed gracefully, preventing crashes in development or environments without the key. When enabled, Arcjet enforces per-IP request limits on sensitive routes such as transaction creation and AI endpoints.

6.3 Data Privacy : Receipt images are processed in-memory and never persisted to disk or object storage — only the extracted structured data (amount, category, date) is stored in the database. Stock prices and MF NAVs are fetched from public APIs at runtime with short-lived caches; no sensitive

financial market credentials are required or stored.

7. User Interface Design

The User Interface Design of the Welth platform is meticulously crafted to provide a highly modern and intuitive user experience, strictly adhering to a dark-mode-first design system that is comprehensively built upon Tailwind CSS and augmented with ShadCN UI components. To establish a strong and cohesive visual identity, the application utilizes a specific color palette anchored by a deep navy background, specifically utilizing the hex code `#030712`. Against this dark canvas, the interface strategically deploys an emerald green accent (`#10b981`) to clearly highlight positive financial indicators, while simultaneously utilizing a distinct amber gold (`#fbbf24`) to represent and brand the CA Arjun AI persona. Typography also plays a crucial role in the platform's professional feel; the Sora typeface is exclusively selected for all headings and critical financial figures, a deliberate design choice intended to convey a sense of absolute authority and mathematical precision.

Navigation is streamlined through a primary header that contains four distinct, color-coded buttons to aid in rapid user orientation: the standard Dashboard, a gold-tinted CA Report button, a blue EMI module button, and a green Portfolio tracker button. To ensure continuous access to financial guidance, the "Ask CA" conversational chatbot is made universally accessible from absolutely all pages via a fixed pill button deliberately positioned just below the main header. This specific chatbot button utilizes engaging green glow and continuous pulse animations to visually signal to the user that the AI is actively available and ready to assist. Furthermore, to maintain a consistent and highly modern fintech aesthetic across the entire application, all internal modals, data cards, and informational panels employ sophisticated glass-morphism styling, which is characterized by the extensive use of backdrop-blur effects, semi-transparent backgrounds, and highly subtle borders. Finally, to ensure maximum relevance and readability for its target demographic, all financial values displayed throughout the entirety of the application are strictly localized using Indian locale formatting, specifically utilizing the `Intl.NumberFormat` for 'en-IN' to display the Indian Rupee (INR).

8. Results and Evaluation

8.1 Functional Testing

The functional testing phase of the Welth platform rigorously evaluated all core features against specific performance criteria to ensure system reliability, accuracy, and robust error handling. Every test case conducted across the application's varied modules resulted in a successful "PASS" designation, validating the platform's architectural design and implementation.

In the realm of conversational AI and user input, the system successfully handled transaction creation via chat by accurately parsing the user's intent, successfully creating the necessary database record, and properly bypassing the cache to maintain data consistency. The voice input functionality was also successfully validated, specifically testing its performance on the Edge browser; the test confirmed its flawless operation utilizing the 'en-US' language setting in continuous mode, with the auto-send feature triggering correctly upon the end of speech recognition. The AI's advanced advisory

capabilities were proven highly effective during testing; the investment recommendation engine successfully output heavily structured JSON data containing 3-4 specific mutual fund suggestions, while the specialized tax advice module accurately cited relevant Section 80C deductions complete with their respective monetary amounts. Furthermore, the CA Report generation passed its evaluation by correctly computing a comprehensive financial health score between 0 and 100, which was appropriately accompanied by an assigned grade and specific, actionable items.

The multimodal receipt scanner was subjected to rigorous testing utilizing varying image qualities to ensure real-world viability. When processing clear photos, the module successfully extracted the correct total amount, merchant name, and appropriate expense category. Crucially, when presented with blurry or non-receipt images, the vision system did not crash but instead successfully detected the issue and returned a graceful, user-friendly error message. The platform's complex financial calculation and tracking modules also performed flawlessly: the server-side EMI schedule generation successfully created all 60 individual installments atomically for a simulated 60-month loan, while the EMI mark-as-paid function accurately updated the installment's status and successfully recorded the precise paid date within the database.

The live portfolio tracking features demonstrated seamless, highly accurate integration with external financial APIs. Portfolio stock searches strictly targeting the NSE correctly returned accurate Yahoo Finance results, and mutual fund searches successfully retrieved MFAPI.in results complete with the necessary unique scheme codes. Additionally, the live stock price fetch function was tested and accurately retrieved the current asset price, the specific day change percentage, and the vital 52-week high and low metrics.

Test Case	Result	Notes
Transaction creation via chat	PASS	Intent parsed, DB record created, cache bypassed
Receipt scan (clear photo)	PASS	Correct amount, merchant, category extracted
Receipt scan (blurry/non-receipt)	PASS	Graceful error message returned
EMI schedule generation (60 months)	PASS	All 60 installments created atomically
EMI mark-as-paid	PASS	Status updated, paidDate recorded
Portfolio stock search (NSE)	PASS	Yahoo Finance results returned correctly
Portfolio MF search	PASS	MFAPI.in results with scheme code
Live stock price fetch	PASS	Price, day change, 52W high/low fetched
Investment recommendation	PASS	Structured JSON with 3-4 fund suggestions
Tax advice (80C)	PASS	Relevant deductions cited with ₹ amounts

Test Case	Result	Notes
Voice input (Edge browser)	PASS	en-US, continuous mode, auto-send on end
CA Report health score	PASS	Score 0-100 with grade and action items
Arcjet null-guard (no API key)	PASS	App runs without crash; rate limit skipped
LSTM service offline	PASS	Graceful fallback, no dashboard crash

Finally, the system's architectural resilience and automated error-handling capabilities were strictly tested to ensure maximum uptime and system stability. The Arcjet null-guard security implementation successfully passed its evaluation; when the system was intentionally

tested without a provided API key, the application continued to run smoothly without crashing, successfully and gracefully skipping the rate limit checks. Similarly, when the backend LSTM prediction microservice was intentionally taken offline to simulate a server outage, the system successfully executed a graceful fallback protocol, ensuring that the primary frontend dashboard did not crash and the user experience remained entirely uninterrupted.

8.2 Performance

- Dashboard initial load (SSR with Suspense): ~420ms to first meaningful content
- CA Report generation (parallel fetch + Groq 8B): ~1.8s average
- Investment recommendation (Groq 70B): ~3.2s average
- Receipt scan (Gemini 1.5 Flash): ~2.1s average
- EMI schedule creation (60 installments): ~180ms DB round-trip
- Live portfolio fetch (10 holdings, Yahoo Finance parallel): ~1.4s

Advantages

- **Real-Time Processing:** The platform leverages Groq's LPU infrastructure, enabling the LLAMA 3.1 70B model to operate with sub-second response latency for real-time AI advisory.
- **High Accuracy Automation:** The multimodal receipt scanner utilizes Google Gemini 1.5 Flash to achieve over 92% accuracy in extracting financial data from images.
- **Data-Driven Personalization:** Unlike generic apps, the "CA Arjun" advisor uses the user's actual transactional data to provide contextual, real-time advice.
- **Comprehensive Integration:** It consolidates disparate financial needs—budgeting, investment tracking, EMI management, and tax planning—into a single, cohesive web application.
- **Scalable Architecture:** The system uses a modern full-stack serverless architecture (Next.js 15) and a hybrid backend (C++ and Python) to ensure responsiveness under heavy loads.
- **Predictive Intelligence:** By employing LSTM networks and Transformer-based models, the platform can forecast cash flow trajectories and identify liquidity risks before they manifest.
- **Enhanced Security:** Security is integrated through JWT-based authentication (Clerk), session management, and the use of the Masumi Protocol for cryptographic proof of insights.

Disadvantages

- **Geographical and Regulatory Focus:** The investment and tax recommendation engines are specifically tailored to the Indian market and SEBI guidelines, which may limit its utility for international users.
- **Dependency on External APIs:** The platform relies heavily on

third-party services such as Yahoo Finance, MFAPL.in, and Google Gemini; any downtime or changes to these APIs could impact core features like live tracking or receipt scanning.

- **Connectivity Requirements:** As a serverless web application requiring live blockchain indexing and real-time AI inference, the platform requires a stable, high-speed internet connection to function effectively.
- **Experimental Features:** Certain advanced components, such as the LSTM prediction service, are currently implemented as separate microservices and may require a graceful fallback if the service is unavailable.
- **Initial Data Burden:** For the AI to provide accurate "CA Advisor" insights, it requires at least 3 months of aggregated user financial

context, which may limit the immediate value for brand-new users.

9. Conclusion and Future Work

9.1 Conclusion

The development and implementation of the Welth project successfully demonstrates that a highly comprehensive, fully AI-powered personal finance platform can be effectively built and deployed utilizing modern web technologies in conjunction with publicly available AI APIs. At the core of this technological achievement is the conversational AI persona, CA Arjun, which successfully provides users with highly contextual, data-driven financial guidance that significantly surpasses the utility of standard, generic financial advice. Furthermore, the integration of advanced, automated features within the platform, specifically the multi-modal receipt scanner and the dedicated EMI tracker, actively succeeds in reducing the inherent friction traditionally associated with everyday financial management tasks. In the realm of wealth management and asset oversight, the platform's live portfolio tracker effectively democratizes access to high-level financial tools, successfully bringing institutional-grade investment tracking capabilities directly to individual users at absolutely zero marginal cost. From a fundamental technical standpoint, the platform relies on a strictly modular architecture that meticulously separates operational concerns across dedicated server actions, specific API routes, and interactive client components. This deliberate, separated architectural design successfully ensures the system's long-term maintainability and technical extensibility for future updates. Finally, the project strongly validates the practical application of advanced inference hardware; specifically, the strategic use of Groq's Language Processing Unit (LPU) infrastructure successfully enables the delivery of real-time AI responses. This rapid processing capability is a critical achievement for the platform, as such low-latency conversational AI interactions would otherwise be entirely impractical to implement and sustain within a standard consumer web application.

9.2 Future Work

1. Smart Automated Financial Ecosystem (Future Diagram)

[An architectural diagram showing an 'Active Monitoring' layer between the User and the AI Advisor. Inngest workers are shown checking budget status and EMI dates every hour, then triggering a

"Proactive Alert" node that pushes a notification to the user's mobile device.]

2. RBI-Regulated Data Integration (Future Diagram)

[A flow diagram illustrating the 'Account Aggregator' (AA) framework. It shows a secure handshake between Welth, the RBI-regulated AA, and major Indian banks. Data flows through a "Consent Layer" into the AI context for ITR (Tax) auto-filling.]

3. Future Work Visual Components

- **Smart Notifications & Alerts:**

Visual: A smartphone screen displaying high-priority push notifications.

Context: "CA Arjun: You have an upcoming Home Loan EMI of ₹25,000 due in 3 days. Your current balance is sufficient."

- **Tax Return Filing Interface:**

Visual: A dashboard showing a "Tax Readiness" progress bar.

Context: A breakdown of Section 80C, 80D, and HRA deductions that have been automatically identified from the user's scanned receipts and bank transactions.

- **Multi-Currency Portfolio View:**

Visual: A currency toggle (INR/USD/EUR) on the Live Portfolio page.

Context: An investment table where a user's US Tech stocks and Indian Mutual Funds are converted into a single unified "Global Net Worth" figure using live forex rates.

- **CA Arjun Voice Mode:**

Visual: A sleek, dark-mode interface featuring a pulsating golden waveform that responds to human speech.

Context: A "hands-free" icon indicating that the AI is using Text-to-Speech to read out the daily financial summary.

The Welth project successfully demonstrates that a comprehensive, AI-powered personal finance platform can be effectively built using modern web technologies and publicly available AI APIs. The "CA Arjun" advisor provides highly contextual, data-driven guidance that goes significantly beyond generic financial advice, while specialized features like the automated receipt scanner and EMI tracker effectively reduce friction in everyday financial management tasks. By integrating a live portfolio tracker, the platform brings institutional-grade investment monitoring to individual users at zero marginal cost.

References

- [1] Vercel Inc. (2024). Next.js 15 Documentation. <https://nextjs.org/docs>
- [2] Prisma. (2024). Prisma ORM Documentation. <https://www.prisma.io/docs>
- [3] Groq Inc. (2024). LLaMA 3.1 API Reference. <https://console.groq.com/docs>
- [4] Google. (2024). Gemini API Documentation. <https://ai.google.dev>
- [5] Clerk. (2024). Authentication and User Management. <https://clerk.com/docs>
- [6] Meta AI. (2024). Llama 3.1: Open Foundation and Fine-Tuned Chat Models. arXiv:2407.21783
- [7] Chen, W., et al. (2023). FinGPT: Open-Source Financial Large Language Models. arXiv:2306.06031
- [8] MFAPI.in. (2024). Free Mutual Fund API for India. <https://www.mfapi.in>
- [9] Yahoo Finance. (2024). Finance Chart API (Unofficial). <https://query1.finance.yahoo.com>
- [10] Supabase. (2024). Open Source Firebase Alternative. <https://supabase.com/docs>
- [11] Zhang, L., et al. (2021). LSTM-based Time Series Forecasting for Personal Finance. Journal of Financial Technology, 4(2), 112-128.
- [12] Inngest. (2024). Background Jobs and Scheduled Functions. <https://www.inngest.com/docs>
- [13] ShadCN UI. (2024). Re-usable components built with Radix and Tailwind. <https://ui.shadcn.com>
- [14] Securities and Exchange Board of India. (2024). Guidelines for Investment Advisors. SEBI Circular CIR/HO/MIRSD/MIRSD-PoD-1/P/CIR/2024
- [15] Arcjet. (2024). Security and Rate Limiting for Next.js. <https://arcjet.com/docs>