A Review of Machine Learning and Deep Learning Approaches for Fraud Detection Across Financial and Supply Chain Domains

Óscar Martínez

omartinez@ull.es

Department of Computer Science and Artificial Intelligence, Universidad de La Laguna (ULL), Tenerife, Spain

Patricia Sánchez

Department of Computer Science and Artificial Intelligence, Universidad de La Laguna (ULL), Tenerife, Spain

Elena Alcaraz

Department of Computer Science and Artificial Intelligence, Universidad de La Laguna (ULL), Tenerife, Spain

---

Systematic Review

Keywords: Fraud Detection, Machine Learning, Deep Learning, Supply Chain, Financial Systems

Posted Date: October 21st, 2025

DOI: https://doi.org/10.21203/rs.3.rs-7893661/v1

Additional Declarations: The authors declare no competing interests.

A Review of Machine Learning and Deep Learning Approaches for Fraud Detection Across Financial and Supply Chain Domains

Óscar Martínez[1] · Patricia Sánchez[1] · Elena Alcaraz[1]

Abstract Fraud detection has emerged as a critical challenge in modern digital economies, affecting both financial institutions and supply chain networks. The rapid proliferation of online transactions, coupled with increasingly sophisticated fraud schemes, has necessitated the development of advanced detection systems capable of identifying anomalous patterns in real-time. This systematic review examines the state-of-the-art machine learning and deep learning approaches for fraud detection across financial and supply chain domains, covering research published between 2015 and 2025. We analyze over 80 studies, categorizing methodologies into traditional machine learning, deep learning, ensemble techniques, semi-supervised learning, and emerging paradigms. Our review identifies key challenges including extreme class imbalance, concept drift, interpretability requirements, and real-time processing constraints. We evaluate the effectiveness of various algorithms using standard performance metrics and discuss their practical deployment considerations. The findings reveal that ensemble methods and tree-based models consistently achieve superior performance in credit card fraud detection, while semi-supervised approaches show promise for supply chain applications where labeled data is scarce. This review provides researchers and practitioners with a comprehensive understanding of current methodologies, their comparative strengths, and future research directions in fraud detection systems.

✉ Ó. Martínez
oscar.martinez@ull.es

P. Sánchez
patricia.sanchez@ull.es

E. Alcaraz
elena.alcaraz@ull.es

[1] Department of Computer Science and Artificial Intelligence, Universidad de La Laguna (ULL), Tenerife, Spain

Keywords: Fraud Detection, Machine Learning, Deep Learning, Supply Chain, Financial Systems

1. Introduction

The exponential growth of digital financial transactions and the increasing complexity of global supply chain networks have created unprecedented opportunities for fraudulent activities. According to recent industry reports, financial fraud losses exceeded $32 billion globally in 2023, with credit card fraud accounting for a significant portion of these losses. Simultaneously, supply chain fraud has emerged as a critical concern, with estimates suggesting that procurement fraud alone costs organizations approximately 5% of their annual revenues. Traditional rule-based fraud detection systems, which rely on predefined patterns and expert knowledge, have proven inadequate in addressing the dynamic and evolving nature of modern fraud schemes. These limitations have driven the financial services industry and supply chain management sectors to adopt sophisticated machine learning (ML) and deep learning (DL) approaches that can automatically learn complex patterns from historical data and adapt to emerging fraud tactics. The application of artificial intelligence in fraud detection represents a paradigm shift from reactive, rules-based systems to proactive, adaptive solutions capable of processing vast amounts of transactional data in real-time.

Despite significant advances in ML and DL techniques, fraud detection systems face several persistent challenges that limit their effectiveness in real-world deployments. The most prominent challenge is extreme class imbalance, where fraudulent transactions typically represent less than 1% of total transactions, leading to biased models that favor the majority class (18). Concept drift, the phenomenon where fraud patterns evolve over time as fraudsters adapt their strategies to evade detection, poses another significant challenge requiring continuous model retraining and adaptation. The scarcity of labeled data, particularly in supply chain domains where fraud annotation requires extensive domain expertise, limits the applicability of supervised learning approaches (24). Additionally, regulatory requirements such as the General Data Protection Regulation (GDPR) and financial compliance standards mandate model interpretability and explainability, constraining the use of black-box deep learning models. Real-time processing requirements further complicate system design, as fraud detection models must make accurate predictions within milliseconds while handling millions of transactions daily. Privacy preservation has also emerged as a critical concern, necessitating the development of federated learning and privacy-preserving techniques that can detect fraud without exposing sensitive customer information.

The motivation for this systematic review stems from the rapid evolution of fraud detection methodologies and the fragmented nature of existing research across financial and supply chain domains. While previous systematic reviews have primarily focused on specific aspects of fraud detection, such as credit card fraud using machine learning techniques (18), there is a notable absence of comprehensive reviews that bridge financial and supply chain fraud detection while systematically comparing traditional ML, deep learning, ensemble, and emerging techniques. The increasing adoption of ensemble methods, which combine multiple models to achieve superior performance, has demonstrated remarkable success in handling class imbalance and improving detection accuracy on benchmark datasets such as IEEE-CIS (25; 26). Furthermore, recent developments in semi-supervised and unsupervised learning have opened new possibilities for fraud detection in data-scarce environments, as evidenced by novel two-phase frameworks combining unsupervised pre-filtering with semi-supervised refinement for supply chain applications (24). The emergence of quantum machine learning, graph neural networks, and federated learning paradigms presents promising but largely unexplored opportunities for next-generation fraud detection systems. This review aims to synthesize current knowledge across both financial and supply chain domains, identify research gaps, and provide actionable insights for both researchers developing novel algorithms and practitioners implementing fraud detection systems in production environments.

This systematic review examines fraud detection research through a multi-dimensional analytical lens, evaluating methodologies based on algorithmic approach, application domain, handling of class imbalance, model interpretability, computational efficiency, and real-world deployment feasibility. We conduct a comprehensive analysis of studies published between 2015 and 2025, with particular emphasis on research utilizing benchmark datasets such as the IEEE-CIS dataset for credit card fraud and the DataCo Smart Supply Chain Dataset for supply chain fraud. Our methodology involves systematic categorization of algorithms into five primary groups: traditional machine learning methods (including Support Vector Machines, Random Forests, and Decision Trees), deep learning approaches (encompassing Convolutional Neural Networks, Long Short-Term Memory networks, and Autoencoders), ensemble techniques (such as XGBoost, LightGBM, and stacking methods) (25; 26), semi-supervised and unsupervised learning paradigms (24), and emerging technologies (including quantum machine learning and graph neural networks). We evaluate each approach using standardized performance metrics including accuracy, precision, recall, F1-score, Area Under the ROC Curve (AUC-ROC), and Area Under the Precision-Recall Curve (AUC-PR), while also considering practical deployment factors such as training time, inference latency, memory requirements, and interpretability. The review critically examines data preprocessing techniques, feature engineering strategies, and imbalance handling methods including SMOTE, ADASYN, Borderline-SMOTE, and various ensemble-based approaches as demonstrated in recent empirical studies (25). Additionally, we analyze the comparative effectiveness of these methodologies across different fraud types and domains, identifying domain-specific requirements and cross-domain transferability of techniques.

The remainder of this paper is structured as follows: Section 2 reviews related work and existing surveys in fraud detection, positioning our contribution within the broader literature. Section 3 describes our systematic review methodology, including literature search strategy, inclusion and exclusion criteria, and the analytical framework used for categorizing and evaluating studies. Section 4 presents a comprehensive analysis of traditional machine learning approaches for fraud detection, examining algorithms such as Logistic Regression, Decision Trees, Random Forests, and Support Vector Machines. Section 5 explores deep learning methodologies, including Convolutional Neural Networks, Recurrent Neural Networks, Long Short-Term Memory networks, and Autoencoders, with emphasis on their ability to capture temporal and sequential patterns. Section 6 discusses ensemble learning techniques and their superior performance in handling imbalanced datasets, covering methods such as Random Forest, XGBoost, LightGBM, and advanced stacking approaches. Section 7 examines semi-supervised and unsupervised learning paradigms, particularly their application in scenarios with limited labeled data. Section 8 investigates emerging technologies including quantum machine learning, graph neural networks, and federated learning approaches. Section 9 provides an in-depth discussion of persistent challenges including class imbalance, concept drift, interpretability, privacy preservation, and real-time processing requirements. Section 10 reviews commonly used benchmark datasets and evaluation protocols. Section 11 presents a comparative analysis of methodologies across different performance dimensions and application domains. Section 12 identifies promising future research directions and emerging trends. Finally, Section 13 summarizes key findings and provides recommendations for researchers and practitioners.

2. Related Work

This section reviews existing literature on fraud detection, categorizing previous research into traditional machine learning approaches, deep learning methods, ensemble techniques, semi-supervised learning, and domain-specific applications. We position our contribution within the broader landscape of fraud detection research and highlight the gaps that motivate this comprehensive review.

2.1. Traditional Machine Learning Approaches

Traditional machine learning algorithms have been extensively applied to fraud detection problems over the past two decades. Support Vector Machines (SVM) have been widely investigated for their ability to handle high-dimensional data and non-linear decision boundaries (27). West and Bhattacharya demonstrated the effectiveness of SVMs in credit card fraud detection, achieving competitive performance compared to neural networks while offering better interpretability (28). Decision trees and their variants have been popular due to their inherent interpretability and ability to handle both numerical and categorical features (29). Random Forest, as an ensemble of decision trees, has consistently shown robust performance across various fraud detection tasks, with studies reporting accuracy rates exceeding 99% on balanced datasets (30).

Logistic Regression remains a baseline method in many fraud detection studies due to its simplicity, computational efficiency, and probabilistic output that facilitates threshold tuning for operational requirements (31). Naive Bayes classifiers have been employed in scenarios where feature independence assumptions approximately hold, particularly in text-based fraud detection and email spam filtering (32). K-Nearest Neighbors (KNN) algorithms have been applied to fraud detection with mixed results, often requiring careful feature scaling and suffering from computational inefficiency on large-scale datasets (33). One persistent limitation of traditional ML approaches is their dependence on manual feature engineering, which requires domain expertise and may not capture complex, latent patterns in transactional data (34).

2.2. Deep Learning Methods

Deep learning has emerged as a powerful paradigm for fraud detection, offering the ability to automatically learn hierarchical feature representations from raw data. Convolutional Neural Networks (CNNs), originally designed for image processing, have been adapted for fraud detection by treating transaction sequences as temporal images or applying 1D convolutions to feature vectors (35). Fu et al. proposed a CNN-based architecture for credit card fraud detection that achieved 99.6% accuracy on the European card transaction dataset (36). Recurrent Neural Networks (RNNs) and their variants, particularly Long Short-Term Memory (LSTM) networks, have demonstrated superior performance in capturing temporal dependencies in transaction sequences (37).

Autoencoders have been extensively used for unsupervised and semi-supervised fraud detection, learning compressed representations of normal transaction patterns and flagging anomalies based on reconstruction error (38). Variational Autoencoders (VAEs) extend this approach by learning probabilistic latent representations, enabling more robust anomaly detection (39). Generative Adversarial Networks (GANs) have been employed to address class imbalance by generating synthetic fraud samples, though their application remains limited due to training instability and mode collapse issues (40). Deep Belief Networks (DBNs) and Restricted Boltzmann Machines (RBMs) have been investigated for their ability to learn hierarchical representations, though their computational cost limits practical adoption (41).

Despite their powerful representation learning capabilities, deep learning methods face significant challenges in fraud detection applications. The requirement for large amounts of labeled training data conflicts with the inherent scarcity of fraud examples in real-world scenarios (42). The black-box nature of deep neural networks raises interpretability concerns, particularly in regulated financial environments where decisions must be explainable to customers and auditors (12). Computational requirements for training and inference can be prohibitive for real-time fraud detection systems processing millions of transactions per second (23). Transfer learning and domain adaptation techniques have been proposed to address data scarcity, but their effectiveness in fraud detection remains an active research area (43).

2.3. Ensemble Learning Techniques

Ensemble methods have gained prominence in fraud detection due to their ability to combine multiple models and achieve superior performance compared to individual classifiers. Bagging-based ensembles, particularly Random Forest, have been widely adopted for their robustness to overfitting and ability to handle high-dimensional feature spaces (44). Boosting algorithms, including AdaBoost, Gradient Boosting, XGBoost, and LightGBM, have demonstrated exceptional performance on imbalanced fraud datasets by iteratively focusing on misclassified examples (45).

XGBoost has emerged as a particularly effective algorithm for fraud detection, consistently ranking among top-performing methods in Kaggle competitions and academic benchmarks (46). Its implementation of regularization techniques, handling of missing values, and parallel processing capabilities make it well-suited for large-scale fraud detection applications (47). LightGBM, with its novel gradient-based one-side sampling and exclusive feature bundling techniques, offers faster training and lower memory consumption while maintaining competitive accuracy (48). CatBoost has been specifically designed to handle categorical features efficiently, which is particularly relevant for fraud detection where transaction categories, merchant types, and user demographics play crucial roles (49).

Stacking and blending approaches, which combine predictions from diverse base models through meta-learning, have shown promise in improving detection accuracy and reducing false positive rates (50). Hybrid ensembles that combine traditional ML algorithms with deep learning models have been proposed to leverage both interpretability and powerful representation learning (30). Heterogeneous ensemble methods that explicitly model diversity among base learners have demonstrated improved generalization on concept-drifting fraud patterns (51). Cost-sensitive ensemble learning, which assigns differential misclassification costs to fraud and legitimate transactions, has been effective in optimizing business-relevant metrics beyond accuracy (52).

2.4. Semi-Supervised and Unsupervised Learning

The scarcity of labeled fraud data has motivated extensive research in semi-supervised and unsupervised learning approaches. Isolation Forest, a tree-based anomaly detection algorithm, has been widely applied to fraud detection due to its efficiency and effectiveness in high-dimensional spaces (53). One-Class SVM and Support Vector Data Description (SVDD) learn a boundary around normal transactions, classifying deviations as potential fraud (54). Local Outlier Factor (LOF) and its variants detect anomalies based on local density deviations, proving effective when fraud patterns exhibit spatial clustering in feature space (55).

Clustering-based approaches, including K-Means, DBSCAN, and Gaussian Mixture Models, have been employed to identify unusual transaction patterns without requiring labeled examples (56). However, purely unsupervised methods often suffer from high false positive rates, necessitating human review or subsequent supervised refinement (57). Semi-supervised learning techniques that leverage small amounts of labeled data alongside abundant unlabeled transactions have shown promise in improving detection accuracy while reducing annotation costs (58). Self-training and co-training approaches iteratively label high-confidence unlabeled examples, gradually expanding the training set (59).

Graph-based semi-supervised learning methods exploit the relational structure in transaction networks, propagating labels through edges connecting similar transactions or linked accounts (60). Positive-Unlabeled (PU) learning treats unlabeled data as a mixture of fraudulent and legitimate transactions, using specialized algorithms to learn from only positive (fraud) examples and unlabeled data (61). Active learning strategies select the most informative unlabeled transactions for expert annotation, maximizing model improvement per labeled example (62). Transfer learning from related fraud detection tasks or domains has been investigated to reduce dependence on domain-specific labeled data (43).

2.5. Domain-Specific Applications

Credit card fraud detection has received the most extensive research attention, with numerous benchmark datasets and competitions driving algorithmic innovation. Studies have investigated transaction-level features such as amount, merchant category, time patterns, and geographic location, as well as aggregated features capturing customer spending behavior over various time windows (23). Device fingerprinting, IP address analysis, and behavioral biometrics have been integrated into fraud detection systems to enhance security (63). Real-time fraud detection architectures employing streaming data processing frameworks such as Apache Kafka and Apache Flink have been developed to meet the stringent latency requirements of payment authorization systems (51).

Insurance fraud detection presents unique challenges including diverse fraud types (claims fraud, application fraud, intermediary fraud) and long investigation cycles that delay ground truth labels (64). Network analysis techniques have been particularly effective in detecting organized fraud rings in insurance claims (65). Healthcare fraud detection systems must handle complex billing codes, diagnosis patterns, and provider-patient relationships while complying with strict privacy regulations such as HIPAA (66). Telecommunications fraud encompasses subscription fraud, call selling, and premium rate service fraud, requiring specialized approaches for call detail record analysis and network traffic monitoring (67).

E-commerce fraud detection systems address challenges including account takeover, payment fraud, return fraud, and promotional abuse (68). User behavior analytics incorporating clickstream data, session patterns, and device characteristics complement traditional transaction-based features (69). Supply chain fraud detection, including procurement fraud, invoice fraud, and counterfeit product detection, has received comparatively less attention in the ML literature despite significant financial impact (34). Blockchain-based approaches have been proposed to enhance supply chain transparency and traceability, though practical implementations remain limited (70). Mobile payment fraud presents emerging challenges due to the proliferation of digital wallets, QR code payments, and peer-to-peer transfer services (71).

2.6. Addressing Class Imbalance

Class imbalance is a fundamental challenge in fraud detection, with fraudulent transactions typically representing less than 1% of total volume. Data-level approaches including random undersampling, random oversampling, and synthetic minority oversampling (SMOTE) have been extensively studied (19). SMOTE variants such as Borderline-SMOTE, ADASYN, and SMOTE-ENN address limitations of the original algorithm by focusing on borderline examples or removing noise from oversampled data (20; 21).

Algorithm-level approaches modify learning algorithms to account for class imbalance, including cost-sensitive learning, threshold adjustment, and one-class classification (72). Ensemble-based methods specifically designed for imbalanced data, such as EasyEnsemble, BalanceCascade, and RUSBoost, combine sampling techniques with ensemble learning (73). Deep learning approaches to imbalance include focal loss, class-balanced loss, and attention mechanisms that emphasize minority class examples during training (22). Evaluation metrics appropriate for imbalanced classification, including precision, recall, F1-score, Matthews Correlation Coefficient, and Area Under the Precision-Recall Curve, have been advocated over accuracy (74).

2.7. Interpretability and Explainability

Model interpretability has become increasingly important in fraud detection due to regulatory requirements, customer trust considerations, and operational needs for fraud analysts to understand and act on model predictions. LIME (Local Interpretable Model-agnostic Explanations) and SHAP (SHapley Additive exPlanations) have been applied to explain individual fraud predictions from black-box models (9; 10). Attention mechanisms in deep learning models provide insight into which features or time steps contribute most to fraud predictions (75).

Rule extraction techniques convert complex models into interpretable rule sets, facilitating human understanding and regulatory compliance (76). Counterfactual explanations identify minimal changes to a transaction that would alter the fraud prediction, providing actionable insights for customers and fraud analysts (77). Prototype-based explanations present representative examples of fraud and legitimate transactions to aid human understanding (78). Inherently interpretable models such as decision trees, rule-based systems, and linear models with sparse feature selection remain relevant in regulated environments despite potentially lower accuracy (13).

2.8. Concept Drift and Model Adaptation

Concept drift, where the statistical properties of fraud patterns change over time, poses significant challenges for deployed fraud detection systems. Ensemble methods with dynamic member selection or weighting adapt to changing data distributions by emphasizing recently accurate models (79). Sliding window approaches retrain models on recent data, though determining optimal window size requires careful tuning (80). Online learning algorithms update model parameters incrementally as new data arrives, enabling continuous adaptation without full retraining (81).

Drift detection methods including DDM (Drift Detection Method), EDDM (Early Drift Detection Method), and ADWIN (Adaptive Windowing) trigger model retraining when performance degradation or distribution shift is detected

(82; 83). Hybrid approaches combine drift detection with ensemble learning, maintaining multiple models corresponding to different concepts and dynamically selecting appropriate models (84). Transfer learning and domain adaptation techniques leverage knowledge from previous time periods or related domains to accelerate adaptation to new fraud patterns (15).

2.9. Privacy-Preserving Fraud Detection

Privacy-preserving machine learning has gained prominence due to regulations such as GDPR and increasing consumer privacy concerns. Federated learning enables collaborative fraud detection across multiple institutions without sharing raw transaction data (4). Differential privacy techniques add calibrated noise to training data or model parameters to prevent inference of individual transaction details (85). Homomorphic encryption allows computation on encrypted data, enabling fraud detection without decrypting sensitive information, though at significant computational cost (86).

Secure multi-party computation protocols enable collaborative fraud detection while protecting each party's proprietary data and models (87). Synthetic data generation techniques create realistic but anonymized transaction datasets for model training and research (88). Privacy-utility trade-offs in fraud detection remain an active research area, balancing detection accuracy against privacy guarantees (5). Blockchain-based approaches provide transparency and auditability while preserving transaction privacy through cryptographic techniques (17).

2.10. Gaps in Existing Literature

Despite extensive research in fraud detection, several gaps remain in the literature. Most studies focus exclusively on either financial or supply chain fraud, lacking comparative analysis across domains. Benchmark datasets are often dated or synthetic, potentially limiting the generalizability of research findings to current fraud patterns. Evaluation protocols vary significantly across studies, making direct comparison of methods difficult. Real-world deployment considerations including latency, throughput, memory constraints, and operational integration receive insufficient attention. The interpretability-accuracy trade-off remains unresolved, particularly for deep learning methods. Comprehensive evaluation of emerging technologies such as quantum machine learning and federated learning in fraud detection contexts is lacking.

This review addresses these gaps by providing a comprehensive, cross-domain analysis of fraud detection methods, systematically comparing approaches across standardized evaluation criteria, discussing practical deployment considerations, and identifying promising directions for future research. By bridging financial and supply chain fraud detection and covering the full spectrum from traditional machine learning to emerging paradigms, this review offers a holistic perspective on the current state and future trajectory of fraud detection research.

3. Systematic Review Methodology

This section describes the systematic methodology employed to identify, select, and analyze relevant literature on machine learning and deep learning approaches for fraud detection. Our review follows the PRISMA (Preferred Reporting Items for Systematic Reviews and Meta-Analyses) guidelines to ensure transparency, reproducibility, and comprehensiveness (89).

3.1. Research Questions

This systematic review addresses the following research questions:

- **RQ1:** What are the primary machine learning and deep learning algorithms employed for fraud detection in financial and supply chain domains?
- **RQ2:** How do different algorithms perform in terms of accuracy, precision, recall, F1-score, and AUC metrics on benchmark datasets?
- **RQ3:** What techniques are most effective for addressing class imbalance in fraud detection datasets?
- **RQ4:** How do researchers handle concept drift and temporal dynamics in fraud patterns?
- **RQ5:** What are the trade-offs between model accuracy and interpretability in fraud detection systems?
- **RQ6:** What are the computational requirements and real-time performance characteristics of different approaches?
- **RQ7:** What are the emerging technologies and future research directions in fraud detection?

3.2. Search Strategy

We conducted a comprehensive literature search across multiple academic databases and digital libraries to identify relevant publications. The primary sources included:

- IEEE Xplore Digital Library
- ACM Digital Library
- ScienceDirect (Elsevier)
- Springer Link
- arXiv preprint repository
- Google Scholar

- Web of Science
- Scopus

The search was conducted between January 2024 and September 2025, covering publications from 2015 to 2025. This ten-year window captures the significant evolution in machine learning and deep learning techniques while focusing on contemporary approaches relevant to current fraud detection challenges.

3.3. Search Terms and Query Construction

We developed a comprehensive search query combining terms related to fraud detection, machine learning techniques, and application domains. The search string was adapted for each database's specific syntax but maintained the following core structure:

```
("fraud detection" OR "fraud prevention" OR "anomaly detection"
OR "outlier detection") AND ("machine learning" OR "deep learning"
OR "neural network" OR "ensemble learning" OR "random forest"
OR "XGBoost" OR "support vector machine" OR "SVM" OR "LSTM"
OR "autoencoder" OR "CNN") AND ("credit card" OR "financial"
OR "transaction" OR "payment" OR "supply chain" OR "procurement"
OR "e-commerce" OR "banking")
```

Additional searches were performed using specific algorithm names (e.g., ”LightGBM fraud detection”, ”semi-supervised fraud detection”, ”federated learning fraud”) to ensure comprehensive coverage of specialized techniques.

3.4. Inclusion and Exclusion Criteria

3.4.1. Inclusion Criteria

Studies were included if they met the following criteria:

1. Published in peer-reviewed journals, conferences, or reputable preprint repositories (arXiv, SSRN)
2. Published between 2015 and 2025
3. Focus on machine learning or deep learning approaches for fraud detection
4. Application to financial transactions, credit cards, banking, supply chain, or related domains
5. Provide empirical evaluation with quantitative performance metrics
6. Written in English
7. Full text available

3.4.2. Exclusion Criteria

Studies were excluded based on the following criteria:

1. Non-peer-reviewed sources (excluding reputable preprints)
2. Purely theoretical work without empirical validation
3. Focus solely on rule-based or traditional statistical methods without ML/DL components
4. Application domains unrelated to financial or supply chain fraud (e.g., academic plagiarism, social media fake accounts)
5. Duplicate publications (in such cases, the most comprehensive version was retained)
6. Short papers or abstracts without sufficient methodological detail
7. Non-English publications

3.5. Study Selection Process

The study selection process followed a multi-stage screening approach:

Stage 1: Initial Search and Deduplication The initial database searches yielded 1,847 potentially relevant publications. After removing duplicates using reference management software (Mendeley and Zotero), 1,312 unique publications remained.

Stage 2: Title and Abstract Screening Two reviewers independently screened titles and abstracts against the inclusion and exclusion criteria. Disagreements were resolved through discussion and, when necessary, consultation with a third reviewer. This stage resulted in 324 publications advancing to full-text review.

Stage 3: Full-Text Assessment Full-text articles were retrieved and assessed for eligibility. Studies were excluded if they did not meet the inclusion criteria upon detailed examination. This stage resulted in 156 publications meeting all criteria.

Stage 4: Quality Assessment and Final Selection A quality assessment was performed on the remaining 156 studies using a scoring rubric evaluating: (1) clarity of methodology, (2) appropriateness of evaluation metrics, (3) dataset transparency, (4) reproducibility, and (5) significance of contribution. Studies scoring below a threshold of 60% were excluded. The final corpus comprised 89 high-quality publications forming the basis of this systematic review.

Stage 5: Snowballing Forward and backward citation snowballing was performed on the selected studies to identify additional relevant publications not captured in the initial search. This resulted in 8 additional studies, bringing the total to 97 publications analyzed in this review.

3.6. Data Extraction

A standardized data extraction form was developed to systematically capture relevant information from each selected study. The extracted data included:

- **Publication metadata:** Authors, year, venue, citation count
- **Study characteristics:** Research objectives, methodology type, application domain
- **Dataset information:** Dataset name, size, class distribution, feature types, availability
- **Algorithms:** ML/DL techniques employed, hyperparameters, ensemble configurations
- **Preprocessing:** Feature engineering, normalization, sampling techniques
- **Imbalance handling:** SMOTE, ADASYN, undersampling, cost-sensitive learning
- **Evaluation metrics:** Accuracy, precision, recall, F1-score, AUC-ROC, AUC-PR
- **Performance results:** Reported metric values, comparison with baselines
- **Computational aspects:** Training time, inference latency, hardware requirements
- **Interpretability:** Explainability techniques, feature importance analysis
- **Limitations:** Acknowledged challenges, threats to validity

Data extraction was performed independently by two reviewers, with discrepancies resolved through discussion.

3.7. Quality Assessment

Each study was evaluated using a quality assessment checklist adapted from established guidelines for ML/DL research (90). The assessment criteria included:

1. **Methodological rigor:** Clear description of algorithms, appropriate experimental design
2. **Dataset quality:** Use of realistic datasets, proper train/test splitting, cross-validation
3. **Evaluation completeness:** Multiple metrics reported, statistical significance testing

4. **Comparison fairness:** Baseline comparisons, consistent evaluation protocols
5. **Reproducibility:** Sufficient implementation details, code/data availability
6. **Generalizability:** Evaluation on multiple datasets, discussion of limitations

Each criterion was scored on a 5-point Likert scale (1 = strongly disagree to 5 = strongly agree). Studies with an average score below 3.0 were excluded from the primary analysis but retained for contextual discussion.

3.8. Data Synthesis and Analysis

We employed both qualitative and quantitative synthesis methods to analyze the extracted data:

3.8.1. Qualitative Synthesis

A thematic analysis approach was used to identify recurring themes, challenges, and research trends across studies. Studies were categorized based on:

- Primary algorithm family (traditional ML, deep learning, ensemble, semi-supervised)
- Application domain (credit card, banking, supply chain, insurance, e-commerce)
- Methodological focus (imbalance handling, concept drift, interpretability, privacy)

3.8.2. Quantitative Synthesis

Where possible, we aggregated performance metrics across studies using similar datasets and evaluation protocols. For the widely-used IEEE-CIS Fraud Detection dataset, we compiled reported performance metrics from multiple studies to enable comparative analysis (25; 26). Similarly, for credit card fraud detection studies using the European cardholders dataset, we synthesized results to identify top-performing approaches (18).

Due to heterogeneity in datasets, preprocessing methods, and evaluation protocols, formal meta-analysis was not feasible. Instead, we provide narrative synthesis supplemented by performance comparison tables organized by dataset and metric type.

3.9. Taxonomy Development

Based on the analyzed literature, we developed a comprehensive taxonomy organizing fraud detection approaches along multiple dimensions:

1. **Algorithmic Approach:** Traditional ML, Deep Learning, Ensemble Methods, Hybrid Approaches
2. **Learning Paradigm:** Supervised, Semi-supervised, Unsupervised, Transfer Learning
3. **Application Domain:** Financial Transactions, Supply Chain, Insurance, E-commerce
4. **Temporal Handling:** Batch Learning, Online Learning, Streaming Processing
5. **Imbalance Strategy:** Sampling-based, Algorithm-based, Ensemble-based, Hybrid

This taxonomy guides the organization of subsequent sections and facilitates systematic comparison of approaches.

3.10. Limitations of the Review Methodology

Our systematic review methodology has several limitations:

1. **Publication bias:** Studies with negative or null results may be underrepresented in the literature
2. **Language bias:** Exclusion of non-English publications may omit relevant research
3. **Database coverage:** Despite comprehensive searching, some relevant publications may have been missed
4. **Preprint quality:** Inclusion of preprints provides recent findings but may include work not yet peer-reviewed
5. **Heterogeneity:** Variation in datasets, metrics, and protocols limits quantitative synthesis
6. **Proprietary systems:** Industry fraud detection systems are often proprietary, limiting available information

We mitigated these limitations through comprehensive multi-database searching, forward/backward citation tracking, and transparent reporting of inclusion/exclusion decisions.

4. Traditional Machine Learning Approaches

Traditional machine learning algorithms have formed the foundation of fraud detection systems for over two decades. This section examines the application of classical ML techniques including logistic regression, decision trees, random forests, support vector machines, naive Bayes, and k-nearest neighbors in fraud detection contexts. We analyze their strengths, limitations, and comparative performance across different datasets and fraud types.

4.1. Logistic Regression

Logistic regression remains one of the most widely deployed algorithms in fraud detection due to its simplicity, interpretability, and computational efficiency. The algorithm models the probability of fraud as a logistic function of input features, making it particularly suitable for scenarios requiring probabilistic outputs and regulatory compliance (91).

4.1.1. Applications and Performance

Logistic regression has been extensively applied to credit card fraud detection with varying degrees of success. Bhattacharyya et al. evaluated logistic regression on an Australian credit card dataset, achieving 85.3% accuracy with relatively low computational overhead (27). The algorithm's strength lies in its ability to provide coefficient interpretability, allowing fraud analysts to understand feature contributions to predictions. However, its linear decision boundary limits effectiveness on datasets with complex, non-linear fraud patterns.

In supply chain fraud detection, logistic regression serves as a reliable baseline method. West and Bhattacharya applied regularized logistic regression with L1 penalty to procurement fraud detection, achieving 78.2% accuracy while maintaining model sparsity through automatic feature selection (28). The regularization approach proved particularly valuable in high-dimensional scenarios with correlated features common in supply chain data.

4.1.2. Handling Class Imbalance

The primary challenge for logistic regression in fraud detection is extreme class imbalance. Standard implementations tend to predict the majority class exclusively when fraud rates fall below 1%. Several mitigation strategies have been investigated:

- **Threshold adjustment:** Moving the decision threshold below 0.5 increases sensitivity to fraud at the cost of higher false positive rates
- **Class weighting:** Assigning higher misclassification costs to fraud cases improves recall but requires careful tuning
- **Sampling techniques:** Combining logistic regression with SMOTE or random undersampling improves performance on minority class (19)

Dal Pozzolo et al. demonstrated that combining undersampling with probability calibration significantly improves logistic regression performance on imbalanced credit card fraud data, achieving AUC-ROC of 0.93 compared to 0.78 without calibration (92).

4.2. Decision Trees and CART

Decision trees provide inherently interpretable models through hierarchical rule structures, making them attractive for fraud detection applications requiring explainability. Classification and Regression Trees (CART) and their variants partition the feature space through recursive binary splitting based on information gain or Gini impurity criteria.

4.2.1. Advantages in Fraud Detection

Decision trees offer several advantages for fraud detection:

1. **Interpretability:** Rules can be directly translated into if-then statements understandable to non-technical stakeholders

2. **Mixed data types:** Native handling of both numerical and categorical features without encoding

3. **Non-linearity:** Capture complex decision boundaries through hierarchical splits

4. **Feature interactions:** Automatically detect interactions without explicit feature engineering

Phua et al. applied C4.5 decision trees to automobile insurance fraud detection, achieving 82.7% accuracy with a transparent rule set comprising 47 decision nodes (29). The extracted rules provided actionable insights for fraud investigators, identifying high-risk claim patterns such as specific combinations of claim amount, vehicle age, and claimant history.

4.2.2. Limitations and Overfitting

Single decision trees suffer from several critical limitations in fraud detection contexts:

- **High variance:** Small changes in training data can result in drastically different tree structures

- **Overfitting:** Unpruned trees memorize training data rather than learning generalizable patterns

- **Limited predictive power:** Single trees often underperform compared to ensemble methods

- **Greedy splitting:** Local optimization at each node may miss globally optimal splits

To address overfitting, pruning techniques including cost-complexity pruning and minimum leaf size constraints are commonly applied. Bhattacharyya et al. found that pruned decision trees achieved 5-7% higher accuracy on test data compared to fully grown trees in credit card fraud detection (27).

4.3. Random Forest

Random Forest, an ensemble of decision trees trained on bootstrap samples with random feature selection, has emerged as one of the most effective traditional ML algorithms for fraud detection. The algorithm combines predictions from multiple trees through majority voting (classification) or averaging (regression), significantly improving upon single decision tree performance.

4.3.1. Performance Across Datasets

Random Forest consistently ranks among top-performing algorithms in fraud detection benchmarks. Studies report the following representative performance metrics:

- **Credit card fraud:** 99.95% accuracy on European cardholders dataset (30)
- **IEEE-CIS dataset:** AUC-ROC of 0.925 with balanced precision-recall trade-offs (25)
- **Insurance fraud:** 94.3% accuracy with 89.7% recall on automobile insurance claims (64)

The algorithm's robustness stems from diversity among ensemble members combined with aggregation that reduces variance. Random Forest naturally handles class imbalance better than single trees through balanced bootstrap sampling options and weighted voting schemes.

4.3.2. Feature Importance Analysis

A key strength of Random Forest in fraud detection is built-in feature importance quantification. Two primary methods are employed:

1. **Mean Decrease Impurity (MDI):** Measures average impurity reduction contributed by each feature across all trees
2. **Permutation Importance:** Evaluates performance degradation when feature values are randomly permuted

Feature importance analysis has revealed critical fraud indicators across domains. In credit card fraud detection, transaction amount, time since last transaction, and merchant category consistently rank as top predictive features (23). For supply chain fraud, invoice amount deviations, vendor history, and approval bypass patterns emerge as key indicators (24).

4.3.3. Hyperparameter Optimization

Random Forest performance depends critically on several hyperparameters:

- **Number of trees (n_estimators):** Typical range 100-1000; diminishing returns beyond 500 trees
- **Maximum depth:** Controls tree complexity; deeper trees risk overfitting
- **Minimum samples split:** Prevents excessive fragmentation of decision nodes
- **Maximum features:** Number of features considered at each split; typically sqrt(n) or log2(n)

Zareapoor and Shamsolmoali demonstrated that grid search optimization of Random Forest hyperparameters improved F1-score from 0.82 to 0.91 on credit card fraud detection, highlighting the importance of systematic tuning (44).

4.4. Support Vector Machines

Support Vector Machines construct optimal hyperplanes that maximize margins between classes in high-dimensional feature spaces. Through kernel functions, SVMs can efficiently handle non-linear decision boundaries, making them theoretically well-suited for fraud detection.

4.4.1. Kernel Selection and Performance

Different kernel functions offer varying capabilities for capturing fraud patterns:

- **Linear kernel:** Computationally efficient but limited to linearly separable problems
- **Polynomial kernel:** Captures feature interactions up to specified degree
- **Radial Basis Function (RBF):** Most commonly used; maps data to infinite-dimensional space
- **Sigmoid kernel:** Mimics neural network activation functions

Bhattacharyya et al. compared SVM kernels on credit card fraud detection, finding RBF kernel achieved the highest accuracy (96.2%) followed by polynomial (94.8%) and linear (91.3%) kernels (27). However, RBF kernel training time was 3-5x longer than linear kernel, presenting trade-offs for large-scale applications.

4.4.2. Handling Imbalanced Data with SVM

Standard SVMs struggle with class imbalance due to bias toward the majority class in margin optimization. Several adaptations address this limitation:

1. **Class weight adjustment:** Penalizes misclassification of minority class more heavily through C parameter

2. **SMOTE-SVM combination:** Synthetic oversampling before SVM training improves minority class representation

3. **One-Class SVM:** Learns boundary around legitimate transactions; deviations flagged as fraud

Rtila and Enneya proposed SVM-based recursive feature elimination combined with hyperparameter optimization specifically for imbalanced fraud detection, achieving 98.7% accuracy on European cardholders dataset (30). Their approach reduced feature dimensionality from 30 to 18 features while improving computational efficiency by 40%.

4.4.3. Computational Challenges

SVM scalability limitations present significant challenges for fraud detection applications:

- **Training complexity:** $O(n^2)$ to $O(n^3)$ time complexity limits applicability to large datasets

- **Memory requirements:** Kernel matrix storage becomes prohibitive for millions of transactions

- **Hyperparameter sensitivity:** Grid search over C and gamma parameters computationally expensive

To address scalability, approximation methods including Nyström approximation and random Fourier features enable SVM application to large-scale fraud detection datasets (93). These techniques reduce computational complexity while maintaining competitive accuracy.

4.5. Naive Bayes Classifiers

Naive Bayes classifiers apply Bayes' theorem with strong independence assumptions between features. Despite the unrealistic assumption of feature independence in fraud detection contexts, Naive Bayes often performs surprisingly well due to its low variance and robustness to irrelevant features.

4.5.1. Variants and Applications

Three primary Naive Bayes variants are applied to fraud detection:

1. **Gaussian Naive Bayes:** Assumes continuous features follow Gaussian distributions
2. **Multinomial Naive Bayes:** Suitable for discrete count features
3. **Bernoulli Naive Bayes:** Appropriate for binary feature representations

Phua et al. evaluated Naive Bayes on automobile insurance fraud, achieving 79.4% accuracy with significantly faster training (5-10x) compared to decision trees and SVMs (29). The algorithm's computational efficiency makes it attractive for real-time fraud screening where rapid predictions are essential.

4.5.2. Limitations in Fraud Detection

Naive Bayes faces several limitations that restrict its effectiveness:

- **Feature independence assumption:** Violated in fraud detection where features exhibit strong correlations
- **Zero-frequency problem:** Unseen feature combinations in training data result in zero probability predictions
- **Poor probability estimates:** While classification accuracy may be acceptable, probability calibration is often poor
- **Limited expressiveness:** Cannot capture complex feature interactions critical for fraud detection

Despite limitations, Naive Bayes serves as a fast baseline method and performs competitively when combined with appropriate feature selection (32).

4.6. K-Nearest Neighbors

K-Nearest Neighbors (KNN) is a non-parametric, instance-based learning algorithm that classifies samples based on the majority class among k nearest neighbors in feature space. The algorithm requires no explicit training phase, instead storing the entire training dataset for prediction.

4.6.1. Distance Metrics and Performance

KNN performance depends critically on distance metric selection:

- **Euclidean distance:** Standard choice but sensitive to feature scaling
- **Manhattan distance:** More robust to outliers; suitable for high-dimensional spaces

- **Mahalanobis distance:** Accounts for feature correlations; computationally expensive
- **Cosine similarity:** Appropriate for directional data and text-based features

Jiang et al. applied KNN with Euclidean distance to credit card fraud detection, achieving 87.3% accuracy with k=5 (33). However, performance degraded significantly on high-dimensional datasets (¿50 features) due to the curse of dimensionality.

4.6.2. Challenges in Fraud Detection

KNN faces substantial challenges in fraud detection applications:

1. **Computational inefficiency:** O(n) prediction time scales poorly to millions of transactions
2. **Memory requirements:** Storing entire training dataset impractical for large-scale systems
3. **Curse of dimensionality:** Distance metrics lose discriminative power in high dimensions
4. **Class imbalance sensitivity:** Nearest neighbors predominantly belong to majority class
5. **Feature scaling sensitivity:** Requires careful normalization to prevent dominance by high-magnitude features

Approximate nearest neighbor methods including locality-sensitive hashing and k-d trees partially address computational challenges but introduce approximation errors (94). Despite optimizations, KNN is rarely competitive with tree-based and ensemble methods in fraud detection benchmarks.

4.7. Comparative Analysis of Traditional ML Methods

Table 1 summarizes the comparative strengths and weaknesses of traditional ML algorithms for fraud detection.

4.7.1. Performance on Benchmark Datasets

We synthesized performance metrics from multiple studies using common benchmark datasets:

European Credit Card Dataset: Random Forest achieved the highest average performance (AUC-ROC: 0.978), followed by SVM with RBF kernel (0.965) and Logistic Regression (0.912). Decision Trees underperformed (0.887) due to overfitting on the highly imbalanced dataset (92).

Table 1. Comparative Analysis of Traditional ML Algorithms for Fraud Detection

Algorithm	Accuracy	Interpretability	Scalability	Imbalance Handling	Training Time
Logistic Regression	Moderate	High	High	Poor	Fast
Decision Tree	Moderate	High	High	Moderate	Fast
Random Forest	High	Moderate	Moderate	Good	Moderate
SVM	High	Low	Low	Poor	Slow
Naive Bayes	Moderate	High	High	Poor	Very Fast
KNN	Moderate	Low	Low	Poor	Fast (no training

IEEE-CIS Dataset: Random Forest again demonstrated superior performance (AUC-ROC: 0.925, F1: 0.847), with careful handling of class imbalance and feature engineering (25; 26). Logistic Regression with regularization achieved competitive results (AUC-ROC: 0.901) while maintaining significantly lower computational requirements.

Supply Chain Datasets: In supply chain fraud detection with limited labeled data, simpler models often performed comparably to complex approaches. Logistic Regression and Decision Trees achieved 78-82% accuracy, while Random Forest improved performance to 85-88% (24).

4.7.2. Practical Deployment Considerations

Algorithm selection for production fraud detection systems must balance multiple factors:

- **Latency requirements:** Real-time authorization systems demand sub-100ms predictions, favoring computationally efficient methods
- **Interpretability mandates:** Regulatory requirements often necessitate explainable models like Logistic Regression or Decision Trees
- **Model maintenance:** Simple models require less frequent retraining and are easier to debug
- **Infrastructure constraints:** Memory and computational resources may limit viable algorithms

Random Forest emerges as the most balanced choice for many fraud detection applications, offering high accuracy, reasonable interpretability through feature importance, and acceptable computational requirements. However, domain-specific requirements may favor alternative approaches (23).

4.8. Limitations of Traditional ML Approaches

Despite widespread success, traditional ML methods face fundamental limitations in fraud detection:

1. **Manual feature engineering:** Requires domain expertise and may miss latent patterns
2. **Fixed representations:** Cannot adaptively learn hierarchical feature representations
3. **Temporal modeling:** Limited ability to capture sequential dependencies in transaction histories
4. **Concept drift:** Require frequent retraining as fraud patterns evolve
5. **Imbalance sensitivity:** Most algorithms struggle with extreme class imbalance without extensive preprocessing

These limitations motivated the development of deep learning approaches, which we examine in the following section.

5. Deep Learning Approaches for Fraud Detection

Deep learning has revolutionized fraud detection by enabling automatic feature learning from raw transactional data, capturing complex non-linear patterns, and modeling temporal dependencies in transaction sequences. This section examines the application of various deep neural network architectures including feedforward networks, convolutional neural networks, recurrent networks, autoencoders, and hybrid architectures in fraud detection contexts.

5.1. Feedforward Neural Networks

Feedforward Neural Networks (FNNs), also known as Multi-Layer Perceptrons (MLPs), consist of multiple layers of interconnected neurons with non-linear activation functions. These networks learn hierarchical representations through backpropagation, mapping input features to fraud predictions.

5.1.1. Architecture and Applications

Typical FNN architectures for fraud detection comprise:

- **Input layer:** Dimension matching the number of transaction features
- **Hidden layers:** 2-5 layers with 128-512 neurons per layer
- **Activation functions:** ReLU or Leaky ReLU for hidden layers, sigmoid for output
- **Regularization:** Dropout (0.3-0.5) and L2 regularization to prevent overfitting

Roy et al. applied deep feedforward networks to credit card fraud detection on the European cardholders dataset, achieving 99.2% accuracy with a 4-layer architecture (256-128-64-32 neurons) (95). The network automatically learned feature representations that outperformed manually engineered features, demonstrating the value of representation learning for fraud detection.

5.1.2. Handling Class Imbalance in Neural Networks

Deep neural networks are particularly susceptible to class imbalance, often converging to trivial solutions that predict only the majority class. Several strategies address this challenge:

1. **Weighted loss functions:** Assign higher weights to minority class samples during training
2. **Focal loss:** Down-weights easy examples and focuses learning on hard negatives (22)
3. **Cost-sensitive learning:** Incorporate business costs of false positives and false negatives
4. **Synthetic oversampling:** Generate artificial minority samples in feature space

Itoo et al. demonstrated that focal loss improved F1-score from 0.74 to 0.89 on highly imbalanced credit card fraud data compared to standard cross-entropy loss (96). The technique proved particularly effective when combined with data augmentation strategies.

5.1.3. Optimization Challenges

Training deep networks for fraud detection presents several challenges:

- **Vanishing gradients:** Deep architectures may suffer from gradient attenuation, addressed through batch normalization and skip connections
- **Hyperparameter sensitivity:** Learning rate, batch size, and architecture choices critically impact performance
- **Overfitting:** High model capacity risks memorizing training data rather than learning generalizable patterns
- **Computational cost:** Training time increases substantially with network depth and width

Bahnsen et al. applied Bayesian optimization for neural network hyperparameter tuning in fraud detection, reducing manual tuning effort while improving AUC-ROC from 0.89 to 0.94 (97).

5.2. Convolutional Neural Networks

Convolutional Neural Networks (CNNs), originally designed for image processing, have been adapted for fraud detection by treating transaction sequences or feature matrices as spatial data. CNNs apply learnable convolutional filters to extract local patterns and hierarchical features.

5.2.1. Adapting CNNs for Tabular Data

Several approaches enable CNN application to fraud detection:

1. **1D convolutions:** Apply temporal convolutions to transaction sequences
2. **2D convolutions:** Treat feature vectors as 2D images or construct feature correlation matrices
3. **Feature embedding:** Transform categorical features into learned embeddings suitable for convolution

Fu et al. proposed a CNN architecture for credit card fraud detection that treats transaction features as a 2D grid, applying multiple convolutional layers with max pooling (36). Their approach achieved 99.6% accuracy on a Chinese credit card dataset, demonstrating CNN's ability to capture local feature interactions without explicit feature engineering.

5.2.2. Temporal Pattern Detection

CNNs excel at detecting local temporal patterns in transaction sequences. Zhang et al. developed HOBA (Histogram of Behavior Attributes), a hybrid approach combining behavioral feature engineering with CNN-based pattern recognition (98). The system constructs histograms of user transaction behavior over time windows, which CNNs process to identify fraudulent patterns. This approach achieved AUC-ROC of 0.967 on real-world banking data, outperforming traditional methods by 8-12%.

5.2.3. Limitations for Fraud Detection

Despite promising results, CNNs face challenges in fraud detection:

- **Spatial assumption violation:** Fraud features lack the spatial locality present in images
- **Feature ordering sensitivity:** Performance depends on how features are arranged spatially
- **Interpretability loss:** Learned convolutional filters difficult to interpret in fraud context
- **Computational overhead:** Convolutions add computational cost compared to feedforward networks

The application of CNNs to tabular fraud data remains less common than to sequential or image-based domains, with recurrent architectures generally preferred for temporal modeling.

5.3. Recurrent Neural Networks and LSTM

Recurrent Neural Networks (RNNs) and their variants, particularly Long Short-Term Memory (LSTM) networks, are specifically designed to model sequential data and temporal dependencies. These architectures maintain hidden states that capture information from previous time steps, making them well-suited for analyzing transaction histories.

5.3.1. LSTM Architecture for Fraud Detection

LSTM networks address the vanishing gradient problem in standard RNNs through gating mechanisms (input, forget, and output gates) that regulate information flow. Typical LSTM architectures for fraud detection include:

- **Input:** Sequences of recent transactions for each account (typically 10-50 transactions)
- **LSTM layers:** 1-3 layers with 64-256 hidden units
- **Attention mechanism:** Optional attention layer to focus on critical transactions
- **Output:** Dense layer with sigmoid activation for binary classification

Jurgovsky et al. applied LSTM networks to credit card fraud detection using sequences of customer transactions (37). Their model achieved AUC-ROC of 0.89 on a real-world banking dataset, with performance improving as sequence length increased from 5 to 20 transactions. The LSTM learned to identify anomalous patterns in spending behavior, location sequences, and temporal rhythms without manual feature engineering.

5.3.2. Bidirectional LSTM and Attention Mechanisms

Bidirectional LSTM (Bi-LSTM) processes sequences in both forward and backward directions, capturing context from both past and future transactions. Wang et al. combined Bi-LSTM with attention mechanisms for e-commerce fraud detection, achieving 94.3% accuracy (68). The attention mechanism learned to focus on the most informative transactions in each sequence, providing interpretability through attention weight visualization.

The attention weights revealed that:

- High-value transactions received significantly higher attention
- Transactions immediately preceding fraud attempts were heavily weighted
- Geographic location changes triggered elevated attention scores

5.3.3. Temporal Feature Learning

A key advantage of LSTM networks is automatic learning of temporal patterns that would require extensive manual feature engineering in traditional approaches. Zheng et al. demonstrated that LSTM-learned features captured:

1. **Velocity features:** Rate of transaction submission and spending patterns
2. **Recency effects:** Time since last transaction and temporal clustering
3. **Sequential anomalies:** Deviations from typical transaction order and timing
4. **Behavioral changes:** Sudden shifts in purchasing patterns or merchant types

Their LSTM model achieved 15-20% higher fraud detection rate compared to Random Forest using manually engineered temporal features (99).

5.3.4. Computational and Practical Challenges

LSTM networks face several practical challenges in fraud detection deployment:

- **Training time:** Sequential processing prevents parallelization, increasing training duration
- **Inference latency:** Real-time systems require sub-100ms predictions; LSTM inference may exceed this threshold
- **Variable sequence lengths:** Transactions histories vary across accounts, requiring padding or masking
- **Cold start problem:** New accounts lack transaction history for sequence-based prediction
- **Memory requirements:** Storing hidden states for long sequences increases memory consumption

To address latency concerns, Jurgovsky et al. proposed a two-stage architecture: a fast rule-based filter for obvious cases, with LSTM applied only to ambiguous transactions requiring deeper analysis (37).

5.4. Gated Recurrent Units

Gated Recurrent Units (GRUs) offer a simplified alternative to LSTMs with fewer parameters and faster training. GRUs combine the forget and input gates into a single update gate, reducing computational complexity while maintaining comparable performance.

Ando et al. compared GRU and LSTM for credit card fraud detection, finding that GRU achieved similar AUC-ROC (0.91 vs 0.92) while reducing training time by 30-40% (100). For large-scale fraud detection systems processing millions of transactions, GRU's efficiency advantage makes it an attractive alternative to LSTM.

5.5. Autoencoders for Anomaly Detection

Autoencoders are unsupervised neural networks that learn compressed representations by encoding input data into lower-dimensional latent space and reconstructing the original input. In fraud detection, autoencoders trained on legitimate transactions learn normal behavior patterns; fraudulent transactions yield high reconstruction errors, enabling anomaly detection.

5.5.1. Autoencoder Architectures

Several autoencoder variants are applied to fraud detection:

1. **Standard autoencoders:** Symmetric encoder-decoder with bottleneck layer
2. **Variational autoencoders (VAE):** Learn probabilistic latent representations with regularization
3. **Denoising autoencoders:** Train on corrupted inputs to learn robust representations
4. **Sparse autoencoders:** Enforce sparsity in latent representations through L1 regularization

Pumsirirat and Yan applied autoencoders to credit card fraud detection on the European cardholders dataset, achieving 94.2% accuracy using reconstruction error as anomaly score (38). The autoencoder learned latent representations that captured legitimate transaction patterns, with fraudulent transactions producing reconstruction errors 3-5x higher than normal transactions.

5.5.2. Advantages for Imbalanced and Unlabeled Data

Autoencoders offer unique advantages for fraud detection:

- **Unsupervised learning:** Training requires only legitimate transactions, abundant in most datasets
- **Imbalance robustness:** No need for balanced training data or sampling techniques
- **Novel fraud detection:** Can detect previously unseen fraud patterns through deviation from normal behavior
- **Dimensionality reduction:** Learned latent representations compress high-dimensional feature spaces

Schreyer et al. applied deep autoencoders to accounting fraud detection, demonstrating effectiveness in scenarios with limited labeled fraud examples (39). The unsupervised approach achieved 89.7% precision at 80% recall, comparable to supervised methods despite using no labeled fraud data during training.

5.5.3. Hybrid Supervised-Unsupervised Approaches

Several studies combine autoencoder-based unsupervised learning with supervised fine-tuning:

1. **Pretraining:** Train autoencoder on legitimate transactions to learn representations
2. **Feature extraction:** Use encoder to transform transactions into latent features
3. **Supervised classification:** Train classifier on latent features with labeled data

Kazemi and Zarrabi demonstrated that this hybrid approach improved F1-score from 0.81 (purely supervised) to 0.88 (pretrained encoder + classifier) on credit card fraud detection (101). The unsupervised pretraining enabled effective learning from limited labeled fraud examples.

5.6. Generative Adversarial Networks

Generative Adversarial Networks (GANs) consist of two competing networks: a generator that creates synthetic samples and a discriminator that distinguishes real from generated samples. In fraud detection, GANs address class imbalance by generating synthetic fraud samples or detect fraud through discriminator scores.

5.6.1. GANs for Class Imbalance

Fiore et al. applied GANs to generate synthetic credit card fraud transactions, augmenting the minority class before training classifiers (40). The GAN-generated synthetic frauds improved Random Forest performance from F1-score 0.79 to 0.86 on highly imbalanced data. However, the approach required careful tuning to ensure synthetic samples reflected realistic fraud patterns rather than noise.

5.6.2. Challenges and Limitations

GAN application to fraud detection faces significant challenges:

- **Training instability:** GANs notoriously difficult to train, prone to mode collapse
- **Quality assessment:** Evaluating synthetic fraud quality is challenging
- **Computational cost:** Training GANs requires substantial computational resources
- **Limited adoption:** Few production fraud detection systems use GANs due to complexity

Despite theoretical promise, GANs remain a research topic with limited practical deployment in fraud detection systems. Simpler oversampling techniques like SMOTE often provide comparable benefits with greater reliability (25).

5.7. Graph Neural Networks

Graph Neural Networks (GNNs) operate on graph-structured data, making them suitable for fraud detection in scenarios with network relationships such as user-merchant interactions, social networks, or supply chain connections.

5.7.1. Applications to Fraud Detection

GNNs have been applied to several fraud detection contexts:

1. **Transaction networks:** Nodes represent users/merchants, edges represent transactions

2. **Social fraud rings:** Detect coordinated fraud through social network analysis

3. **Supply chain fraud:** Model relationships between suppliers, distributors, and products

Wang et al. proposed a GNN-based approach for e-commerce fraud detection that models user-product-merchant interaction graphs (8). Their Graph Convolutional Network (GCN) achieved 96.8% AUC-ROC by aggregating information from neighboring nodes, capturing collaborative fraud patterns invisible to transaction-level models.

5.7.2. Advantages for Network-Based Fraud

GNNs offer unique capabilities for detecting fraud involving multiple entities:

- **Relationship modeling:** Capture connections between fraudsters and compromised accounts

- **Information propagation:** Aggregate features from multi-hop neighborhoods

- **Fraud ring detection:** Identify coordinated attacks through network structure

- **Transfer learning:** Fraud patterns learned in one part of network generalize to others

Despite advantages, GNN application requires graph-structured data and faces scalability challenges on networks with millions of nodes and billions of edges.

5.8. Transformer Models

Transformer architectures, based on self-attention mechanisms, have revolutionized sequence modeling in natural language processing. Recent work explores their application to fraud detection through transaction sequence modeling.

Cheng et al. adapted BERT (Bidirectional Encoder Representations from Transformers) for credit card fraud detection, treating transaction sequences as "sentences" and individual transactions as "words" (102). The self-attention mechanism learned to identify relevant transactions in long sequences more effectively than LSTM, achieving AUC-ROC of 0.945 compared to 0.921 for LSTM on real-world banking data.

However, transformers face challenges in fraud detection:

- **Computational cost:** Quadratic complexity in sequence length limits scalability
- **Data requirements:** Transformers require large training datasets for effective learning
- **Interpretability:** Attention weights provide some interpretability but are less transparent than decision trees

5.9. Hybrid and Ensemble Deep Learning Architectures

Hybrid architectures combine multiple deep learning components to leverage complementary strengths. Common combinations include:

1. **CNN-LSTM:** CNN extracts spatial features, LSTM models temporal dependencies
2. **Autoencoder-Classifier:** Autoencoder learns representations, classifier makes predictions
3. **Multi-modal networks:** Process different feature types (numerical, categorical, sequential) through specialized sub-networks

Randhawa et al. proposed a hybrid CNN-LSTM architecture for credit card fraud detection that achieved 99.7% accuracy on the European cardholders dataset (103). The CNN component extracted local patterns from engineered features while LSTM captured temporal dependencies, outperforming both architectures individually.

5.10. Comparative Performance Analysis

Table 2 summarizes the comparative characteristics of deep learning approaches for fraud detection.

Table 2. Comparative Analysis of Deep Learning Approaches for Fraud Detection

Architecture	Temporal	Interpretability	Training Time	Inference Speed	Data Requiremen
Feedforward NN	No	Low	Moderate	Fast	Moderate
CNN	Limited	Low	Moderate	Fast	Moderate
LSTM/GRU	Excellent	Low	Slow	Moderate	High
Autoencoder	No	Low	Moderate	Fast	Low (unlabeled)
GAN	No	Very Low	Very Slow	Fast	High
GNN	No	Low	Slow	Moderate	High (+ graph)
Transformer	Excellent	Moderate	Very Slow	Moderate	Very High

5.10.1. Performance on Benchmark Datasets

Synthesizing results across studies using common datasets:

European Credit Card Dataset:

- CNN-LSTM hybrid: 99.7% accuracy (103)
- Standard LSTM: 99.2% accuracy (37)
- Feedforward NN: 99.0% accuracy (95)
- Autoencoder: 94.2% accuracy (38)

IEEE-CIS Dataset: Deep learning approaches have been less extensively evaluated on IEEE-CIS compared to ensemble methods, with reported performance generally comparable to or slightly below optimized Random Forest and XGBoost models (25; 26).

5.11. Interpretability and Explainability in Deep Learning

The black-box nature of deep neural networks poses challenges for fraud detection deployment in regulated environments. Several approaches enhance deep learning interpretability:

1. **Attention visualization:** Attention weights in LSTM/Transformer models highlight influential transactions
2. **SHAP values:** Model-agnostic explanations quantify feature contributions to predictions
3. **Layer-wise Relevance Propagation (LRP):** Traces prediction back through network layers to input features
4. **Activation maximization:** Identifies input patterns that maximally activate specific neurons

Despite these techniques, deep learning models remain less interpretable than traditional ML methods like decision trees or logistic regression, limiting adoption in contexts requiring regulatory compliance or customer-facing explanations.

5.12. Practical Deployment Challenges

Deploying deep learning models in production fraud detection systems faces several challenges:

- **Latency requirements:** Payment authorization systems demand sub-100ms predictions; deep networks may exceed this threshold
- **Model serving infrastructure:** GPU acceleration may be required for real-time inference at scale
- **Model updates:** Retraining deep networks to address concept drift is computationally expensive
- **A/B testing complexity:** Evaluating new models in production requires careful experimental design
- **Monitoring and debugging:** Deep network failures are harder to diagnose than traditional ML issues

These practical considerations explain why many production fraud detection systems continue to rely on ensemble tree methods despite deep learning's superior performance in academic benchmarks (23).

5.13. Future Directions in Deep Learning for Fraud Detection

Several promising research directions are emerging:

1. **Self-supervised learning:** Leverage abundant unlabeled transaction data through contrastive learning and other self-supervised objectives
2. **Few-shot learning:** Detect novel fraud types from limited examples through meta-learning
3. **Continual learning:** Adapt models to concept drift without catastrophic forgetting
4. **Neural architecture search:** Automatically discover optimal architectures for specific fraud detection tasks
5. **Federated learning:** Collaborative fraud detection across institutions without sharing sensitive data

These directions aim to address current limitations while maintaining the representation learning advantages that make deep learning attractive for fraud detection.

6. Ensemble Learning Techniques

Ensemble learning combines multiple base models to achieve superior predictive performance compared to individual classifiers. This section examines ensemble methods including bagging, boosting, stacking, and hybrid approaches, with particular emphasis on their effectiveness in addressing class imbalance and achieving state-of-the-art performance in fraud detection.

6.1. Fundamentals of Ensemble Learning

Ensemble methods leverage the wisdom of crowds principle: aggregating predictions from diverse models often produces more accurate and robust results than any single model. The effectiveness of ensembles depends on two key factors:

1. **Accuracy:** Individual models should perform better than random guessing
2. **Diversity:** Models should make different types of errors, enabling complementary predictions

Ensemble methods are categorized into three primary paradigms:

- **Bagging (Bootstrap Aggregating):** Train models on bootstrap samples and aggregate predictions through voting or averaging
- **Boosting:** Sequentially train models, with each focusing on examples misclassified by predecessors
- **Stacking:** Train a meta-model to combine predictions from diverse base models

6.2. Random Forest

Random Forest, discussed briefly in Section 4, represents one of the most successful ensemble methods for fraud detection. Its combination of bagging with random feature selection creates diverse decision trees that collectively achieve robust performance.

6.2.1. Advanced Random Forest Techniques

Several advanced Random Forest variants have been specifically developed for fraud detection:

1. **Balanced Random Forest:** Uses balanced bootstrap sampling to address class imbalance by ensuring equal representation of fraud and legitimate transactions in each tree's training set
2. **Weighted Random Forest:** Assigns higher misclassification costs to minority class, steering trees toward better fraud detection

3. **Cost-Sensitive Random Forest:** Incorporates business costs of false positives and false negatives directly into splitting criteria

Dal Pozzolo et al. evaluated Balanced Random Forest on the European cardholders dataset, achieving AUC-PR of 0.85 compared to 0.72 for standard Random Forest (92). The balanced sampling approach significantly improved fraud recall from 78% to 89% while maintaining acceptable precision.

6.2.2. Performance on Fraud Detection Benchmarks

Random Forest consistently achieves top-tier performance across fraud detection datasets:

European Credit Card Dataset:

- Accuracy: 99.95% (30)
- AUC-ROC: 0.978 (18)
- AUC-PR: 0.85 (Balanced RF) (92)

IEEE-CIS Dataset:

- Accuracy: 96.8% (25)
- AUC-ROC: 0.925 (26)
- F1-Score: 0.847 (25)

The algorithm's robust performance, interpretability through feature importance, and computational efficiency make it a preferred choice for many production fraud detection systems.

6.3. Gradient Boosting Machines

Gradient Boosting builds an ensemble of weak learners (typically shallow decision trees) sequentially, with each tree correcting errors made by the ensemble of previous trees. The algorithm minimizes a loss function through gradient descent in function space, enabling optimization of various objectives including accuracy, precision-recall trade-offs, and custom business metrics.

6.3.1. Core Algorithm

The gradient boosting process for fraud detection follows these steps:

1. Initialize model with a constant prediction (typically log-odds of fraud rate)
2. For each boosting iteration:
    - Compute pseudo-residuals (gradients of loss function)
    - Fit a decision tree to predict pseudo-residuals

  - Update ensemble by adding the new tree with a learning rate multiplier

3. Final prediction combines all trees: $F(x) = \sum_{m=1}^{M} \gamma_m h_m(x)$

Zareapoor and Shamsolmoali applied Gradient Boosting to credit card fraud detection, achieving 97.3% accuracy with 500 trees of depth 4 (44). The sequential error correction enabled the ensemble to focus learning capacity on difficult-to-classify borderline cases.

6.3.2. Handling Class Imbalance in Gradient Boosting

Several strategies address class imbalance in gradient boosting:

- **Sample weighting:** Assign higher weights to minority class samples in loss function
- **Focal loss:** Down-weight easy examples and focus on hard negatives
- **Scale_pos_weight parameter:** Adjusts the balance between positive and negative classes
- **Custom evaluation metrics:** Optimize for F1-score or AUC-PR rather than accuracy

6.4. XGBoost

XGBoost (eXtreme Gradient Boosting) extends gradient boosting with algorithmic and systems optimizations that dramatically improve performance and scalability. XGBoost has become the dominant algorithm in Kaggle competitions and is widely deployed in production fraud detection systems.

6.4.1. Key Innovations

XGBoost introduces several innovations over standard gradient boosting:

1. **Regularization:** L1 and L2 penalties on leaf weights prevent overfitting
2. **Sparsity awareness:** Efficient handling of missing values and sparse features
3. **Weighted quantile sketch:** Optimal split finding for weighted datasets
4. **System optimizations:** Cache-aware access, out-of-core computation, parallel tree construction

Chen and Guestrin's original XGBoost paper demonstrated 10x faster training compared to standard gradient boosting implementations while achieving superior accuracy (46).

6.4.2. Performance in Fraud Detection

XGBoost consistently ranks among top-performing algorithms in fraud detection benchmarks:

IEEE-CIS Fraud Detection Competition: Multiple teams achieved top-10 finishes using XGBoost as their primary algorithm, with winning solutions typically employing XGBoost in ensemble stacks. Reported performance metrics include:

- AUC-ROC: 0.942–0.958 (depending on feature engineering)
- Training time: 2–5 hours on full dataset with GPU acceleration
- Inference latency: ¡10ms per transaction

The algorithm's combination of accuracy, speed, and built-in handling of missing values and imbalanced data makes it particularly well-suited for fraud detection (25; 26).

6.4.3. Hyperparameter Tuning

XGBoost's performance depends critically on hyperparameter selection:

- **n_estimators:** Number of boosting rounds; typically 500–3000
- **max_depth:** Tree depth; 3–10 for fraud detection (deeper for more complex patterns)
- **learning_rate (eta):** Step size shrinkage; 0.01–0.1 (lower for better generalization)
- **subsample:** Fraction of samples used per tree; 0.6–1.0
- **colsample_bytree:** Fraction of features used per tree; 0.6–1.0
- **scale_pos_weight:** Balance between positive and negative classes; set to ratio of negative to positive samples
- **gamma:** Minimum loss reduction for split; 0–10 (higher values increase regularization)
- **min_child_weight:** Minimum sum of instance weights in a child; 1–10

Van Vlasselaer et al. demonstrated that Bayesian optimization of XGBoost hyperparameters improved AUC-ROC from 0.89 to 0.94 on credit card fraud detection, highlighting the importance of systematic tuning (108).

6.5. LightGBM

LightGBM (Light Gradient Boosting Machine) implements novel techniques that significantly improve training speed and memory efficiency compared to XGBoost while maintaining comparable or superior accuracy.

6.5.1. Novel Techniques

LightGBM introduces two key innovations:

1. **Gradient-based One-Side Sampling (GOSS):** Retains all instances with large gradients (difficult to classify) but randomly samples from instances with small gradients, reducing computational cost while maintaining accuracy

2. **Exclusive Feature Bundling (EFB):** Bundles mutually exclusive features (features that rarely take non-zero values simultaneously) to reduce dimensionality, particularly beneficial for sparse fraud detection features

Ke et al. demonstrated that LightGBM achieved accuracy comparable to XGBoost while training 20x faster on large datasets (48).

6.5.2. Application to Fraud Detection

LightGBM has been successfully applied to various fraud detection domains:

Credit Card Fraud: Rtayli and Enneya applied LightGBM to the European cardholders dataset, achieving 99.6% accuracy with 40% shorter training time compared to XGBoost (30). The algorithm's efficiency enables more frequent model retraining to adapt to concept drift.

E-commerce Fraud: LightGBM's handling of categorical features through native categorical splitting (rather than one-hot encoding) proved particularly advantageous for e-commerce fraud detection with features like product category, user agent, and geographic location.

6.5.3. Categorical Feature Handling

Unlike XGBoost which requires one-hot encoding of categorical features, LightGBM can directly split on categorical features by finding optimal partitions. This capability:

- Reduces memory consumption by avoiding one-hot encoding expansion
- Improves training speed by considering fewer split candidates
- Enhances interpretability by keeping categorical features intact
- Better captures information in high-cardinality categorical features common in fraud detection

6.6. CatBoost

CatBoost (Categorical Boosting) is specifically designed to handle categorical features efficiently while addressing prediction shift and target leakage issues in gradient boosting.

6.6.1. Key Features

CatBoost introduces several innovations:

1. **Ordered boosting:** Prevents target leakage by using different subsets for gradient calculation and base model training

2. **Categorical feature encoding:** Novel encoding scheme based on target statistics with appropriate smoothing

3. **Oblivious decision trees:** Symmetric trees with same splitting criterion at each level, improving prediction speed

Prokhorenkova et al. demonstrated that CatBoost's handling of categorical features yielded superior performance on datasets with many categorical variables (49).

6.6.2. Fraud Detection Applications

CatBoost's strengths align well with fraud detection requirements:

- Merchant category, user agent, country code, and other categorical features are prevalent

- Target encoding provides powerful representations of categorical variables

- Ordered boosting reduces overfitting on small fraud datasets

- Fast prediction through symmetric trees benefits real-time systems

However, CatBoost has seen less adoption in fraud detection compared to XGBoost and LightGBM, possibly due to longer training times and more complex parameterization.

6.7. AdaBoost

AdaBoost (Adaptive Boosting) was one of the first successful boosting algorithms, adaptively re-weighting training samples based on classification errors. While largely superseded by gradient boosting methods, AdaBoost remains relevant for certain fraud detection scenarios.

Randhawa et al. combined AdaBoost with majority voting for credit card fraud detection, achieving 99.2% accuracy on the European cardholders dataset (103). The approach proved particularly effective when combined with cost-sensitive learning to address class imbalance.

6.8. Stacking and Blending

Stacking (stacked generalization) trains a meta-model to combine predictions from multiple diverse base models. Unlike simple voting or averaging, stacking learns optimal weighting and non-linear combinations of base model predictions.

6.8.1. Stacking Architecture

A typical stacking ensemble for fraud detection comprises:

1. **Level 0 (Base models):** Diverse algorithms trained on original features
    - XGBoost with different hyperparameters
    - LightGBM
    - Random Forest
    - Neural Network
    - Logistic Regression
2. **Level 1 (Meta-model):** Classifier trained on base model predictions
    - Common choices: Logistic Regression, XGBoost, Neural Network
    - Input: Out-of-fold predictions from base models
    - Output: Final fraud probability

6.8.2. Performance in Fraud Detection

Stacking ensembles have achieved state-of-the-art results in fraud detection competitions and benchmarks:

IEEE-CIS Dataset: The stacking approach proposed by Moradi et al. achieved superior performance compared to individual models (25; 26):

- AUC-ROC: 0.918 (vs. 0.905 for best single model)
- AUC-PR: 0.891 (vs. 0.874 for best single model)
- F1-Score: 0.847 (vs. 0.829 for best single model)

The ensemble combined XGBoost, LightGBM, and Random Forest as base models with Logistic Regression as meta-learner, demonstrating that stacking can effectively leverage complementary strengths of different algorithms.

6.8.3. Avoiding Overfitting in Stacking

Stacking ensembles risk overfitting, particularly on small fraud datasets. Several strategies mitigate this risk:

- **Out-of-fold predictions:** Use k-fold cross-validation to generate base model predictions, ensuring meta-model trains on unseen predictions
- **Regularized meta-model:** Apply L1/L2 regularization to meta-model to prevent complex weighting schemes
- **Simple meta-models:** Logistic Regression or linear models often outperform complex meta-learners
- **Feature selection:** Include only most diverse base models to maximize complementarity

6.8.4. Blending vs. Stacking

Blending is a simpler alternative to stacking that splits training data into:

1. Training set for base models
2. Holdout set for meta-model training

While blending is computationally simpler, stacking generally achieves better performance by utilizing all training data through cross-validation. In fraud detection where labeled data is scarce, stacking's more efficient data usage typically outweighs blending's simplicity advantage.

6.9. Voting Ensembles

Voting ensembles combine predictions from multiple models through:

- **Hard voting:** Majority vote determines final prediction
- **Soft voting:** Average predicted probabilities, then threshold
- **Weighted voting:** Assign weights to models based on validation performance

Randhawa et al. demonstrated that majority voting across AdaBoost, Random Forest, and XGBoost achieved 99.4% accuracy on credit card fraud detection, outperforming individual models (103). The simplicity of voting makes it attractive for production systems requiring transparent ensemble logic.

6.10. Diversity in Ensemble Construction

Ensemble effectiveness depends critically on diversity among base models. Several approaches ensure diversity:

1. **Algorithm diversity:** Combine tree-based, linear, and neural models
2. **Hyperparameter diversity:** Train same algorithm with different configurations
3. **Feature diversity:** Provide different feature subsets to base models
4. **Sample diversity:** Train on different bootstrap samples or cross-validation folds
5. **Objective diversity:** Optimize different loss functions (accuracy, precision, recall)

Dastile et al. analyzed ensemble diversity in credit scoring and fraud detection, finding that combining models from different algorithm families (e.g., tree-based + linear + neural) yielded greater improvement than combining multiple tree-based models (50).

6.11. Imbalance-Aware Ensemble Methods

Several ensemble methods specifically target class imbalance:

6.11.1. EasyEnsemble

EasyEnsemble creates multiple balanced subsets by undersampling the majority class, trains a model on each subset, and combines predictions through voting. Liu et al. demonstrated that EasyEnsemble achieved 12–15% higher fraud recall compared to standard ensemble methods while maintaining acceptable precision (**?**).

6.11.2. BalanceCascade

BalanceCascade iteratively trains classifiers on balanced subsets and removes correctly classified majority class instances from subsequent training sets. This cascade approach focuses later classifiers on difficult-to-classify cases.

6.11.3. RUSBoost

RUSBoost combines Random Under-Sampling with AdaBoost, creating balanced training sets at each boosting iteration. Seiffert et al. showed that RUSBoost outperformed both standard AdaBoost and SMOTE-based approaches on highly imbalanced fraud datasets (104).

6.11.4. SMOTEBoost

SMOTEBoost applies SMOTE oversampling at each boosting iteration to balance class distribution. While computationally expensive, this approach achieved strong performance on fraud detection benchmarks with extreme imbalance (fraud rate ¡0.1%).

6.12. Ensemble Learning for Different Fraud Types

6.12.1. Credit Card Fraud

Ensemble methods dominate credit card fraud detection, with Random Forest, XGBoost, and stacking ensembles achieving near-perfect performance on benchmark datasets. The IEEE-CIS Fraud Detection competition demonstrated that all top-performing solutions employed ensemble approaches (25; 26).

6.12.2. Supply Chain Fraud

For supply chain fraud with limited labeled data, ensemble methods combining supervised and unsupervised approaches show promise. Moradi et al. proposed a two-phase ensemble framework using Isolation Forest for pre-filtering followed by semi-supervised SVM, achieving F1-score of 0.817 with ¡3% false positive rate (24).

6.12.3. Insurance Fraud

Insurance fraud detection benefits from ensembles that combine traditional fraud indicators with network analysis. Subelj et al. used Random Forest to combine behavioral features with social network metrics derived from claim relationships, achieving 94.3% accuracy on automobile insurance fraud (64).

6.13. Computational Considerations

6.13.1. Training Efficiency

Ensemble training time depends on:

- Number of base models
- Individual model complexity
- Parallelization capabilities

Bagging methods (Random Forest) parallelize easily across trees, while boosting methods (XGBoost, LightGBM) parallelize across features and samples within each tree. Stacking requires sequential training: base models first, then meta-model.

6.13.2. Inference Latency

Production fraud detection systems require real-time predictions (¡100ms). Ensemble inference latency considerations:

- **Random Forest:** Parallel tree evaluation; 1000 trees typically ¡50ms
- **XGBoost/LightGBM:** Sequential tree evaluation; optimized implementations achieve ¡20ms for 1000 trees
- **Stacking:** Sequential evaluation of base models then meta-model; latency sum of components

Model compression techniques including tree pruning, quantization, and knowledge distillation can reduce ensemble size and improve inference speed with minimal accuracy loss.

6.14. Comparative Analysis

Table 3 summarizes the comparative characteristics of ensemble methods for fraud detection.

Table 3. Comparative Analysis of Ensemble Learning Methods for Fraud Detection

Method	Accuracy	Training Speed	Inference Speed	Interpretability	Imbalance Han
Random Forest	High	Moderate	Fast	Moderate	Good
Gradient Boosting	High	Slow	Fast	Moderate	Moderate
XGBoost	Very High	Moderate	Fast	Moderate	Excellent
LightGBM	Very High	Fast	Very Fast	Moderate	Excellent
CatBoost	Very High	Slow	Very Fast	Moderate	Excellent
AdaBoost	Moderate	Moderate	Fast	Moderate	Poor
Stacking	Very High	Slow	Moderate	Low	Excellent
Voting	High	Moderate	Moderate	Moderate	Good

6.15. Best Practices for Ensemble Learning in Fraud Detection

Based on extensive research and practical deployments, we recommend:

1. **Start with XGBoost or LightGBM:** These algorithms provide excellent performance with reasonable training time and built-in imbalance handling
2. **Systematic hyperparameter tuning:** Use Bayesian optimization or random search for critical hyperparameters
3. **Feature engineering remains critical:** Ensemble methods benefit from well-engineered features despite automatic feature learning capabilities
4. **Address class imbalance explicitly:** Use scale_pos_weight, custom loss functions, or sampling techniques
5. **Validate on multiple metrics:** Accuracy alone is insufficient; evaluate AUC-ROC, AUC-PR, F1-score, and business metrics
6. **Consider stacking for maximum performance:** When computational budget permits, stacking diverse models achieves best results
7. **Monitor inference latency:** Ensure ensemble complexity doesn't violate real-time requirements
8. **Implement A/B testing:** Gradually deploy new ensemble models with careful monitoring of production performance

6.16. Limitations and Challenges

Despite superior performance, ensemble methods face limitations:

- **Interpretability:** Combining hundreds or thousands of trees reduces transparency compared to single decision trees

- **Overfitting risk:** Highly flexible ensembles can overfit small fraud datasets
- **Computational resources:** Training large ensembles requires significant computation and memory
- **Hyperparameter complexity:** Many hyperparameters require tuning, increasing development time
- **Model maintenance:** Updating ensemble models in production is more complex than updating single models

These challenges explain why some production systems balance ensemble complexity against operational requirements, sometimes favoring simpler models with slightly lower accuracy but greater maintainability.

7. Semi-Supervised and Unsupervised Learning Approaches

The scarcity of labeled fraud data presents a fundamental challenge in fraud detection. Labeling requires expert investigation, is time-consuming and expensive, and often lags behind fraud occurrence due to detection delays. Semi-supervised and unsupervised learning approaches address this challenge by leveraging abundant unlabeled transaction data alongside limited labeled examples, or by identifying anomalies without requiring labels at all. This section examines various paradigms including anomaly detection, clustering, semi-supervised classification, and hybrid approaches.

7.1. Motivation and Problem Formulation

Traditional supervised learning for fraud detection assumes availability of large labeled datasets with balanced or manageable class distributions. In practice, several factors limit labeled data availability:

1. **Labeling delay:** Fraud often discovered days or weeks after occurrence
2. **Investigation cost:** Expert review of suspicious transactions is expensive
3. **Label uncertainty:** Some cases remain ambiguous even after investigation
4. **Novel fraud patterns:** New fraud types lack historical labeled examples
5. **Privacy constraints:** Data sharing restrictions limit access to labeled fraud data

Semi-supervised learning addresses these challenges by learning from both labeled and unlabeled data, while unsupervised approaches identify fraud through deviation from normal behavior patterns without requiring any labels.

7.2. Anomaly Detection Fundamentals

Anomaly detection identifies rare observations that deviate significantly from the majority of data. In fraud detection, the assumption is that fraudulent transactions are anomalous compared to legitimate transactions. The effectiveness of anomaly detection depends on:

- **Anomaly definition:** Clear characterization of what constitutes abnormal behavior
- **Feature representation:** Appropriate features that distinguish fraud from legitimate transactions
- **Threshold selection:** Balancing true positive rate against false positive rate
- **Concept drift handling:** Adapting to evolving definitions of normal behavior

7.3. Isolation Forest

Isolation Forest is a tree-based anomaly detection algorithm that identifies anomalies based on the principle that anomalies are easier to isolate than normal points. The algorithm constructs multiple isolation trees through random feature selection and random split value selection.

7.3.1. Algorithm Principles

Isolation Forest operates on the insight that:

1. Anomalies are few and different, thus easier to isolate
2. Anomalies require fewer random partitions to be isolated
3. Path length in isolation trees inversely correlates with anomaly score

For each transaction, the anomaly score is computed as:

$$s(x, n) = 2^{-\frac{E(h(x))}{c(n)}} \tag{1}$$

where $h(x)$ is the path length, $E(h(x))$ is the expected path length over all trees, and $c(n)$ is the average path length for a binary search tree.

7.3.2. Application to Fraud Detection

Liu et al. introduced Isolation Forest and demonstrated its effectiveness for anomaly detection on various datasets (53). The algorithm's key advantages for fraud detection include:

- **No labeled data required:** Purely unsupervised approach

- **Computational efficiency:** Linear time complexity and low memory footprint
- **Handling high dimensions:** Effective in high-dimensional feature spaces
- **No distance metric needed:** Avoids issues with distance metrics in high dimensions

Moradi et al. applied Isolation Forest as the first phase of a two-stage fraud detection system for supply chain fraud (24). The unsupervised pre-filtering stage:

- Identified 23.7% of transactions as potential anomalies
- Captured 94.2% of actual fraud cases in the anomalous subset
- Reduced data volume for subsequent supervised learning by 76.3%

This approach achieved F1-score of 0.817 with false positive rate below 3.0%, demonstrating the effectiveness of Isolation Forest for fraud pre-screening in data-scarce environments.

7.3.3. Hyperparameters and Tuning

Key Isolation Forest hyperparameters include:

- **n_estimators:** Number of trees; typically 100-200 sufficient
- **max_samples:** Subsample size; smaller values (256) improve efficiency
- **contamination:** Expected proportion of anomalies; guides threshold selection
- **max_features:** Number of features per tree; affects tree diversity

The contamination parameter is critical: setting it too high increases false positives, while too low misses fraud. In practice, contamination is often set based on business tolerance for false alarms or tuned on a small validation set with known fraud labels.

7.4. One-Class SVM

One-Class Support Vector Machine (OC-SVM) learns a decision boundary encompassing normal data points, classifying points outside this boundary as anomalies. The algorithm solves an optimization problem to find the smallest hypersphere (or hyperplane in modified formulations) containing most training data.

7.4.1. Mathematical Formulation

OC-SVM with RBF kernel maps data to a high-dimensional feature space and finds a hyperplane with maximum margin from the origin. The optimization objective is:

$$\min_{w,\rho,\xi} \frac{1}{2}||w||^2 - \rho + \frac{1}{\nu n}\sum_{i=1}^{n}\xi_i \tag{2}$$

subject to $(w \cdot \phi(x_i)) \geq \rho - \xi_i$ and $\xi_i \geq 0$, where ν controls the fraction of outliers and support vectors.

7.4.2. Applications and Performance

Tax and Duin introduced Support Vector Data Description (SVDD), a variant of OC-SVM specifically designed for anomaly detection (54). In fraud detection applications:

Credit Card Fraud: Bhattacharyya et al. applied OC-SVM to credit card fraud detection, achieving 87.3% precision at 75% recall (27). The model trained on legitimate transactions successfully identified fraudulent patterns without requiring labeled fraud examples.

Insurance Fraud: OC-SVM has been applied to detect anomalous insurance claims by learning normal claim patterns and flagging deviations for investigation.

7.4.3. Challenges and Limitations

OC-SVM faces several challenges in fraud detection:

- **Kernel selection:** Performance heavily depends on kernel choice and parameters
- **Scalability:** Computational complexity limits application to very large datasets
- **Sensitivity to parameter ν:** Requires careful tuning or domain knowledge
- **Assumption of single normal cluster:** Struggles when legitimate behavior is multimodal

Despite limitations, OC-SVM remains valuable for small to medium-sized fraud detection problems where labeled fraud examples are unavailable.

7.5. Local Outlier Factor

Local Outlier Factor (LOF) detects anomalies based on local density deviation. Unlike global methods that identify points far from all other points, LOF identifies points whose local density is significantly lower than their neighbors' density.

7.5.1. LOF Computation

For each point, LOF computes:

1. **k-distance:** Distance to k-th nearest neighbor
2. **Reachability distance:** Maximum of k-distance of neighbor and distance to that neighbor
3. **Local reachability density:** Inverse of average reachability distance
4. **LOF score:** Ratio of neighbor densities to point's own density

Points with LOF $\gg 1$ are anomalies (lower density than neighbors), while LOF ≈ 1 indicates normal points.

7.5.2. Fraud Detection Applications

Breunig et al. introduced LOF for density-based outlier detection (55). In fraud detection:

- LOF effectively identifies fraud in spatially clustered transaction data
- Local approach handles multiple normal behavior clusters better than global methods
- Particularly useful for geographic fraud detection where normal patterns vary by region

Kou et al. surveyed fraud detection techniques and highlighted LOF's effectiveness for detecting fraud clusters that deviate from regional norms (56).

7.5.3. Computational Complexity

LOF's main limitation is computational cost:

- $O(n^2)$ complexity for brute-force k-nearest neighbor search
- Mitigated through approximate nearest neighbor algorithms (LSH, k-d trees)
- Ball tree and KD-tree acceleration enable application to datasets with millions of transactions

7.6. Clustering-Based Anomaly Detection

Clustering algorithms group similar transactions together, enabling anomaly detection through:

1. **Distance to cluster center:** Points far from nearest cluster center are anomalies
2. **Small cluster identification:** Clusters with very few members may represent fraud
3. **Cluster membership confidence:** Weak cluster membership indicates potential anomalies

7.6.1. K-Means for Fraud Detection

K-Means partitions data into k clusters by minimizing within-cluster variance. Fraud detection applications:

- Cluster legitimate transactions into k groups representing normal behavior patterns
- Compute distance of new transactions to nearest cluster center
- Flag transactions exceeding distance threshold as suspicious

Ahmed et al. applied K-Means clustering to credit card fraud detection, achieving 94.2% accuracy by combining clustering with outlier score computation (105). However, K-Means assumes spherical clusters and struggles with varying density patterns common in fraud data.

7.6.2. DBSCAN

Density-Based Spatial Clustering of Applications with Noise (DBSCAN) identifies clusters of arbitrary shape and explicitly labels low-density points as outliers. DBSCAN's advantages for fraud detection include:

- No need to specify number of clusters
- Automatic outlier detection through noise points
- Handles clusters of varying density and shape

Ester et al. introduced DBSCAN for spatial database applications (106). In fraud detection, DBSCAN successfully identifies anomalous transaction patterns that deviate from dense legitimate transaction clusters.

7.6.3. Gaussian Mixture Models

Gaussian Mixture Models (GMM) represent data as a mixture of Gaussian distributions, providing probabilistic cluster assignments. Fraud detection via GMM:

1. Fit GMM to legitimate transactions
2. Compute likelihood of new transactions under fitted model
3. Flag low-likelihood transactions as potential fraud

GMM's probabilistic framework enables natural threshold setting based on likelihood ratios and provides uncertainty quantification for borderline cases.

7.7. Autoencoder-Based Anomaly Detection

Autoencoders, discussed in Section 5.5, are particularly effective for unsupervised fraud detection. The approach trains autoencoders exclusively on legitimate transactions to learn compressed representations of normal behavior.

7.7.1. Reconstruction Error as Anomaly Score

The reconstruction error quantifies how well the autoencoder can reconstruct input transactions:

$$\text{Anomaly Score} = ||x - \hat{x}||^2 \tag{3}$$

where x is the original transaction and $\hat{x}$ is the reconstruction. High reconstruction error indicates anomalous transactions that deviate from learned normal patterns.

7.7.2. Performance in Fraud Detection

Pumsirirat and Yan demonstrated that autoencoder reconstruction error effectively identifies credit card fraud, achieving 94.2% accuracy without labeled fraud data during training (38). The approach:

- Trained on 99% legitimate transactions
- Set anomaly threshold at 95th percentile of training reconstruction errors
- Achieved precision of 89.7% at 85% recall on test data

Schreyer et al. applied deep autoencoders to accounting fraud detection, demonstrating effectiveness in scenarios with extremely limited labeled examples (39).

7.7.3. Variational Autoencoders

Variational Autoencoders (VAEs) extend standard autoencoders by learning probabilistic latent representations. VAEs provide:

- More robust anomaly detection through probabilistic framework
- Better handling of data uncertainty
- Improved generalization to novel fraud patterns

An et al. proposed VAE-based anomaly detection using reconstruction probability rather than reconstruction error, achieving superior performance on several fraud detection benchmarks (107).

7.8. Semi-Supervised Learning

Semi-supervised learning leverages both labeled and unlabeled data to improve model performance beyond what is achievable with labeled data alone. This paradigm is particularly relevant for fraud detection where unlabeled transactions are abundant but labeled examples are scarce.

7.8.1. Self-Training

Self-training iteratively expands the labeled dataset:

1. Train classifier on initial labeled data
2. Predict labels for unlabeled data
3. Add high-confidence predictions to labeled set
4. Retrain classifier and repeat

Moradi et al. applied self-training Support Vector Machine (SVM) for supply chain fraud detection following Isolation Forest pre-filtering (24). The approach:

- Started with 2,847 labeled transactions (1.2% of dataset)
- Used Isolation Forest to reduce unlabeled data volume by 76.3%
- Applied self-training SVM to iteratively label high-confidence anomalies
- Achieved F1-score of 0.817 with false positive rate under 3.0%

The two-phase framework demonstrated that combining unsupervised pre-filtering with semi-supervised refinement effectively addresses fraud detection in severely data-constrained environments.

7.8.2. Co-Training

Co-training trains multiple classifiers on different feature views and uses their agreement to label unlabeled data. Requirements:

- Multiple feature views (e.g., transaction attributes vs. user behavior features)
- Views should be conditionally independent given the class
- Each view should be sufficient for learning

Li et al. applied co-training to fraud detection by splitting features into transaction-level and account-level views, achieving 7-10% improvement over single-view supervised learning (59).

7.8.3. Graph-Based Semi-Supervised Learning

Graph-based methods propagate labels through similarity graphs:

1. Construct graph where nodes are transactions and edges connect similar transactions
2. Initialize labeled nodes with known fraud/legitimate labels
3. Propagate labels through graph based on edge weights
4. Assign labels to unlabeled nodes based on neighbor labels

Van Vlasselaer et al. applied graph-based semi-supervised learning to social security fraud detection, leveraging network connections between claimants to improve detection accuracy (108).

7.9. Positive-Unlabeled Learning

Positive-Unlabeled (PU) learning addresses scenarios where only positive examples (fraud) are labeled, while unlabeled data contains a mixture of positive and negative examples. This setting is common in fraud detection when only confirmed fraud cases are labeled, but many unlabeled transactions remain uninvestigated.

7.9.1. PU Learning Approaches

Two primary PU learning strategies:

1. **Two-step approach:**
    - Identify reliable negative examples from unlabeled data
    - Train classifier on positive examples and reliable negatives
2. **Biased learning approach:**
    - Treat unlabeled examples as negative with appropriate weighting
    - Account for label noise in unlabeled set

7.9.2. Applications in Fraud Detection

Jain et al. developed PU learning methods for scenarios with only labeled fraud examples, recovering near-optimal classifier performance despite absence of labeled legitimate transactions (61). The approach:

- Estimated class prior (fraud rate) from unlabeled data
- Identified likely negative examples through density estimation
- Trained classifier with appropriate sample weighting

PU learning achieved within 2-3% of fully supervised performance while requiring labels for only fraud cases, significantly reducing labeling cost.

7.10. Active Learning

Active learning strategically selects the most informative unlabeled examples for expert labeling, maximizing model improvement per labeled example. This approach addresses the high cost of fraud investigation by focusing human effort on the most valuable cases.

7.10.1. Query Strategies

Common active learning query strategies for fraud detection:

1. **Uncertainty sampling:** Select examples where model is least confident
2. **Query-by-committee:** Select examples where ensemble models disagree most
3. **Expected model change:** Select examples that would most change the model if labeled
4. **Expected error reduction:** Select examples that would most reduce expected error

7.10.2. Fraud Detection Applications

Das et al. applied active learning to anomaly detection, demonstrating that strategic labeling of 5% of data achieved performance comparable to fully supervised learning with 40% labeled data (62). In fraud detection contexts:

- Active learning reduces investigation costs by 60-80%
- Most effective when combined with expert feedback loops
- Particularly valuable for adapting to novel fraud patterns

7.10.3. Practical Considerations

Deploying active learning in production fraud detection systems requires:

- Infrastructure for human-in-the-loop labeling
- Appropriate batch sizes balancing labeling efficiency and model updates
- Handling labeling delays (fraud confirmation may take days/weeks)
- Quality control for expert labels

7.11. Transfer Learning

Transfer learning leverages knowledge from related domains or tasks to improve fraud detection in target domains with limited labeled data. This approach is particularly relevant for:

- New financial institutions with limited fraud history
- Emerging fraud types lacking labeled examples
- Cross-domain fraud detection (e.g., credit card → e-commerce)

7.11.1. Domain Adaptation Techniques

Lebichot et al. applied deep learning domain adaptation to credit card fraud detection, transferring knowledge from one bank's fraud patterns to another bank (45). Techniques include:

1. **Feature-based adaptation:** Learn domain-invariant representations
2. **Instance-based adaptation:** Re-weight source domain samples to match target distribution
3. **Model-based adaptation:** Fine-tune pre-trained models on target domain

The study demonstrated that transfer learning improved AUC-ROC from 0.78 (training on target domain only) to 0.89 (transferring from source domain) when target labeled data was limited to 1% of total transactions.

7.11.2. Cross-Domain Transfer

Zheng et al. developed TrAdaBoost for transfer learning in fraud detection, adapting models trained on one fraud type to detect different fraud types (99). The approach:

- Automatically identifies transferable patterns between domains
- Down-weights non-transferable source domain examples
- Achieves 12-18% improvement over target-only training when target labels are scarce

7.12. Hybrid Supervised-Unsupervised Approaches

Hybrid approaches combine strengths of supervised and unsupervised learning to address fraud detection challenges comprehensively.

7.12.1. Two-Phase Frameworks

The two-phase approach proposed by Moradi et al. exemplifies effective hybrid design (24):

Phase 1 - Unsupervised Pre-filtering:

- Isolation Forest identifies potential anomalies
- Reduces data volume by 76.3%
- Captures 94.2% of fraud in filtered subset

Phase 2 - Semi-supervised Refinement:

- Self-training SVM refines predictions on filtered data
- Leverages limited labeled examples more effectively
- Achieves F1-score of 0.817 with ¡3% false positive rate

This framework demonstrates superior performance compared to purely supervised or unsupervised approaches, particularly in data-scarce supply chain fraud detection scenarios.

7.12.2. Ensemble of Supervised and Unsupervised Models

Combining supervised classifiers with unsupervised anomaly detectors through ensemble voting or stacking has shown promise:

- Supervised models capture known fraud patterns from labeled examples
- Unsupervised models detect novel fraud types
- Ensemble combines complementary strengths

7.13. Comparative Analysis

Table 4 summarizes the comparative characteristics of semi-supervised and unsupervised approaches.

7.14. Evaluation Challenges

Evaluating semi-supervised and unsupervised methods presents unique challenges:

1. **Lack of ground truth:** Unsupervised methods cannot be evaluated on true labels
2. **Threshold selection:** Anomaly scores require threshold tuning, often requiring labeled validation data

Table 4. Comparative Analysis of Semi-Supervised and Unsupervised Methods

Method	Label Requirements	Novel Fraud Detection	Computational Cost	False Positive Rate	Scalability
Isolation Forest	None	Excellent	Low	High	Excellent
One-Class SVM	None	Good	Moderate	High	Moderate
LOF	None	Excellent	High	High	Poor
Clustering	None	Good	Moderate	High	Good
Autoencoder	None	Excellent	Moderate	Moderate	Good
Self-Training	Few	Good	Moderate	Moderate	Good
Co-Training	Few	Good	Moderate	Moderate	Good
PU Learning	Positive only	Moderate	Moderate	Moderate	Good
Active Learning	Few (strategic)	Good	Low	Low	Excellent
Transfer Learning	Source + Few target	Good	High	Moderate	Moderate

3. **Novel fraud evaluation:** Models detecting previously unseen fraud cannot be evaluated on historical test sets

4. **False positive costs:** High false positive rates incur investigation costs

Practical evaluation often employs:

- Small labeled validation sets for threshold tuning
- Expert review of top-ranked anomalies
- A/B testing in production with careful monitoring
- Precision at fixed recall levels (e.g., precision at 80% recall)

7.15. Best Practices and Recommendations

Based on research findings and practical deployments:

1. **Start with unsupervised pre-filtering:** Isolation Forest or autoencoders reduce false positives in subsequent supervised stages

2. **Combine multiple approaches:** Hybrid methods outperform pure supervised or unsupervised approaches

3. **Leverage domain knowledge:** Feature engineering remains critical even for unsupervised methods

4. **Use semi-supervised learning when labels are scarce:** Self-training and PU learning effectively leverage limited labels

5. **Implement active learning:** Strategic labeling reduces investigation costs by 60-80%

6. **Consider transfer learning for new deployments:** Leverage fraud patterns from related domains

7. **Carefully tune anomaly thresholds:** Balance fraud detection rate against investigation capacity

8. **Monitor and adapt:** Concept drift affects unsupervised methods; regular retraining necessary

7.16. Future Directions

Promising research directions include:

- **Self-supervised learning:** Leverage contrastive learning and other self-supervised objectives for fraud representation learning

- **Few-shot learning:** Detect novel fraud types from 1-5 examples through meta-learning

- **Weakly supervised learning:** Learn from noisy or imprecise labels (e.g., "suspicious" vs. confirmed fraud)

- **Federated unsupervised learning:** Collaborative anomaly detection across institutions without sharing data

8. Emerging Technologies and Future Paradigms

As fraud detection continues to evolve, several emerging technologies and novel paradigms show promise for addressing current limitations and enabling next-generation fraud detection systems. This section examines quantum machine learning, federated learning, graph neural networks, explainable AI, and other cutting-edge approaches that represent the future of fraud detection research and deployment.

8.1. Quantum Machine Learning

Quantum machine learning (QML) leverages quantum computing principles including superposition, entanglement, and quantum interference to potentially achieve exponential speedups over classical algorithms for certain machine learning tasks. While practical quantum computers remain in early stages, QML research explores theoretical foundations and near-term applications to fraud detection.

8.1.1. Quantum Computing Fundamentals

Quantum computers operate on qubits that can exist in superposition of states, enabling parallel exploration of exponentially large solution spaces. Key quantum computing concepts relevant to machine learning include:

- **Superposition:** Qubits simultaneously represent multiple states, enabling massive parallelism

- **Entanglement:** Correlated qubit states capture complex feature interactions
- **Quantum interference:** Amplifies correct solutions while suppressing incorrect ones
- **Quantum gates:** Unitary transformations implement quantum algorithms

8.1.2. Quantum Algorithms for Fraud Detection

Several quantum algorithms have theoretical potential for fraud detection applications:

1. **Quantum Support Vector Machine (QSVM):** Rebentrost et al. proposed quantum SVM with exponential speedup in dimensionality and polynomial speedup in training samples (1). For high-dimensional fraud detection features, QSVM could theoretically achieve dramatic acceleration over classical SVM.
2. **Quantum Neural Networks:** Quantum circuits implement parameterized quantum gates analogous to neural network weights. Farhi and Neven introduced quantum neural networks for supervised learning with potential advantages in optimization landscape (2).
3. **Quantum k-Means:** Lloyd et al. developed quantum clustering algorithms with exponential speedup over classical k-means (3), potentially enabling real-time clustering of massive transaction datasets.
4. **Quantum Sampling:** Quantum algorithms for sampling from probability distributions could accelerate Bayesian inference and generative modeling for fraud detection.

8.1.3. Current Limitations and Challenges

Despite theoretical promise, quantum machine learning faces significant practical challenges:

- **Hardware limitations:** Current quantum computers have 50-1000 noisy qubits, insufficient for practical fraud detection
- **Error rates:** Quantum decoherence and gate errors limit computation depth
- **Input/output bottleneck:** Loading classical data into quantum states and reading results can eliminate theoretical speedups
- **Algorithm design:** Few quantum algorithms demonstrate proven advantages for machine learning

- **Specialized problems:** Quantum advantages may only manifest for specific problem structures

Moradi et al. noted in their systematic review that quantum machine learning remains constrained by scalability and implementation complexity, with practical fraud detection applications likely years away (18).

8.1.4. Near-Term Prospects

Variational quantum algorithms represent the most promising near-term approach, combining quantum circuits with classical optimization:

- **Quantum variational classifiers:** Train parameterized quantum circuits for classification
- **Quantum kernel methods:** Use quantum computers to evaluate kernel functions for classical ML
- **Hybrid quantum-classical algorithms:** Partition computation between quantum and classical processors

Research focuses on identifying fraud detection problem structures where quantum approaches offer measurable advantages over classical methods, even on near-term noisy quantum devices.

8.2. Federated Learning

Federated learning enables collaborative machine learning across multiple institutions without sharing raw transaction data, addressing privacy concerns while leveraging collective knowledge to improve fraud detection.

8.2.1. Federated Learning Architecture

The federated learning paradigm for fraud detection involves:

1. **Local training:** Each institution trains model on its private transaction data
2. **Model aggregation:** Central server aggregates model updates (not raw data)
3. **Global model distribution:** Updated global model distributed back to institutions
4. **Iterative refinement:** Process repeats until convergence

The most common aggregation strategy is Federated Averaging (FedAvg), which computes weighted average of local model parameters based on training set sizes.

8.2.2. Applications to Fraud Detection

Federated learning addresses several fraud detection challenges:

- **Data privacy:** Financial institutions share model improvements without exposing customer transactions
- **Regulatory compliance:** Satisfies GDPR and other privacy regulations
- **Cross-institutional learning:** Small banks benefit from fraud patterns learned by larger institutions
- **Novel fraud detection:** Fraud observed at one institution quickly incorporated into global model

Yang et al. surveyed federated learning applications in finance, highlighting fraud detection as a key use case (4). Multiple financial institutions collaboratively trained fraud detection models, achieving performance improvements of 12-18% compared to institution-specific models while maintaining data privacy.

8.2.3. Privacy and Security Considerations

Despite privacy benefits, federated learning faces security challenges:

1. **Model inversion attacks:** Adversaries reconstruct training data from model updates
2. **Membership inference:** Determine whether specific transaction was in training set
3. **Poisoning attacks:** Malicious participants corrupt global model through adversarial updates
4. **Free-riding:** Institutions benefit from global model without contributing quality updates

Differential privacy techniques add calibrated noise to model updates, providing formal privacy guarantees at the cost of reduced accuracy. Wei et al. analyzed privacy-accuracy trade-offs in federated learning, demonstrating that carefully tuned differential privacy maintains fraud detection performance while preventing privacy breaches (5).

8.2.4. Federated Learning Variants

Several federated learning variants address specific fraud detection requirements:

- **Vertical federated learning:** Institutions hold different features for same customers (e.g., bank transactions + credit bureau data)

- **Federated transfer learning:** Transfer fraud patterns across institutions with different feature spaces
- **Personalized federated learning:** Global model adapted to each institution's specific fraud patterns
- **Asynchronous federated learning:** Institutions contribute updates asynchronously, improving practical feasibility

8.2.5. Challenges and Future Directions

Key challenges for federated fraud detection include:

- **Non-IID data:** Fraud patterns vary across institutions, violating independent and identically distributed assumptions
- **Communication costs:** Frequent model synchronization generates significant network overhead
- **Heterogeneous systems:** Institutions have different computational resources and data volumes
- **Incentive mechanisms:** Encouraging participation and high-quality contributions

Research directions include adaptive aggregation strategies robust to non-IID data, compressed communication protocols, and blockchain-based incentive mechanisms for federated fraud detection consortia.

8.3. Graph Neural Networks for Network-Based Fraud

Graph Neural Networks (GNNs), introduced briefly in Section 5.7, represent a powerful emerging paradigm for fraud detection in scenarios with explicit network structure.

8.3.1. Advanced GNN Architectures

Recent GNN architectures offer enhanced capabilities for fraud detection:

1. **Graph Attention Networks (GAT):** Learn attention weights over neighbors, focusing on most relevant connections (6)
2. **GraphSAINT:** Scalable GNN training through graph sampling, enabling application to transaction networks with billions of edges (7)
3. **Heterogeneous GNNs:** Handle multiple node and edge types (users, merchants, transactions, devices)
4. **Temporal GNNs:** Model evolving transaction networks where edges and node features change over time

8.3.2. Fraud Ring Detection

GNNs excel at detecting organized fraud rings where multiple fraudsters coordinate attacks. Wang et al. developed a semi-supervised GNN for financial fraud detection that achieved 96.8% AUC-ROC by modeling user-merchant-device interaction graphs (8). The approach identified fraud rings invisible to transaction-level models by analyzing:

- Shared devices across multiple accounts
- Coordinated transaction timing patterns
- Suspicious merchant networks
- Cascading fraud propagation through referral networks

8.3.3. Relational Features

GNNs automatically learn relational features that capture fraud indicators:

- **Node centrality:** Accounts central to fraud networks exhibit suspicious connectivity patterns
- **Community structure:** Fraudsters often form densely connected communities
- **Velocity patterns:** Rapid propagation of transactions through network subgraphs
- **Behavioral similarity:** Aggregated features from neighbors reveal account types

8.3.4. Scalability Challenges

Applying GNNs to large-scale fraud detection faces computational challenges:

- **Graph size:** Transaction networks contain millions of nodes and billions of edges
- **Neighborhood explosion:** K-hop neighborhoods grow exponentially
- **Dynamic graphs:** Transaction networks evolve continuously, requiring incremental updates
- **Memory constraints:** Storing and processing large graphs exceeds GPU memory

Sampling strategies including GraphSAGE, FastGCN, and GraphSAINT enable scalable GNN training on large transaction networks, though with potential information loss from sampling.

8.4. Explainable AI for Fraud Detection

Explainable AI (XAI) techniques address the black-box nature of complex machine learning models, providing interpretability required for regulatory compliance, customer communication, and fraud analyst decision support.

8.4.1. Model-Agnostic Explanation Methods

Several model-agnostic techniques explain predictions from any fraud detection model:

1. **SHAP (SHapley Additive exPlanations):** Lundberg and Lee developed SHAP based on game-theoretic Shapley values, providing consistent feature importance quantification (9). SHAP values indicate each feature's contribution to pushing prediction from base value to final output.

2. **LIME (Local Interpretable Model-agnostic Explanations):** Ribeiro et al. proposed LIME, which explains individual predictions by fitting interpretable models locally (10). For fraud detection, LIME identifies which transaction features most influenced the fraud prediction.

3. **Anchors:** High-precision rules that "anchor" predictions, providing if-then explanations for fraud decisions (11).

4. **Counterfactual Explanations:** Identify minimal changes to transaction that would alter prediction, answering "what would need to change for this transaction to be classified as legitimate?"

8.4.2. Application to Fraud Detection

Hasan et al. applied explainable AI techniques to credit card fraud detection, demonstrating that SHAP-based explanations improved fraud analyst efficiency by 35% (12). Explanations enabled analysts to:

- Quickly verify whether model predictions aligned with domain knowledge
- Identify spurious correlations learned by models
- Communicate fraud reasons to customers during transaction declines
- Debug model failures and improve feature engineering

8.4.3. Regulatory Requirements

Financial regulations increasingly mandate explainability for automated decisions:

- **GDPR Article 22:** Right to explanation for automated decisions affecting individuals

- **Fair Credit Reporting Act:** Requires adverse action reasons for credit decisions
- **Model Risk Management:** Regulators require understanding of model decision logic

XAI techniques enable compliance while maintaining sophisticated fraud detection models. However, tension remains between detailed explanations and protecting intellectual property in fraud detection algorithms.

8.4.4. Interpretable-by-Design Models

Alternative to post-hoc explanations, interpretable-by-design models build explainability into architecture:

- **Attention mechanisms:** Attention weights in neural networks provide interpretability
- **Prototype-based models:** Explain predictions through similarity to prototypical examples
- **Neural additive models:** Combine neural network flexibility with additive structure interpretability
- **Rule-based neural networks:** Extract symbolic rules from trained neural networks

Rudin argued for inherently interpretable models over complex models with post-hoc explanations, particularly for high-stakes decisions like fraud detection (13).

8.5. Continual Learning and Concept Drift Adaptation

Fraud patterns evolve continuously as fraudsters adapt tactics, causing concept drift that degrades model performance over time. Continual learning (also called lifelong learning) addresses this challenge by enabling models to learn new patterns while retaining knowledge of previous patterns.

8.5.1. Continual Learning Approaches

Three primary continual learning strategies:

1. **Regularization-based:** Penalize changes to important parameters learned from previous data (14)
2. **Rehearsal-based:** Store representative samples from previous data and replay during training on new data
3. **Architecture-based:** Allocate different network components to different time periods or concepts

8.5.2. Application to Fraud Detection

Continual learning addresses fraud detection challenges:

- **Catastrophic forgetting:** Standard retraining on new data causes models to forget previous fraud patterns
- **Incremental adaptation:** Models continuously adapt to emerging fraud without full retraining
- **Computational efficiency:** Incremental updates more efficient than periodic batch retraining
- **Rapid response:** Quick adaptation to novel fraud tactics

Lu et al. surveyed learning under concept drift, highlighting fraud detection as a key application domain requiring robust drift adaptation (15). They identified several drift types in fraud detection:

- **Sudden drift:** Abrupt change in fraud patterns (e.g., new attack vector)
- **Gradual drift:** Slow evolution of fraud tactics over time
- **Recurring concepts:** Seasonal or cyclical fraud patterns
- **Incremental drift:** Continuous small changes in fraud distribution

8.5.3. Drift Detection Methods

Detecting concept drift triggers model updates:

1. **Statistical tests:** Monitor prediction error distribution changes (e.g., Page-Hinkley test)
2. **Performance monitoring:** Track metrics like AUC-ROC on recent data
3. **Data distribution monitoring:** Detect feature distribution shifts
4. **Ensemble diversity:** Increased disagreement among ensemble members signals drift

8.5.4. Challenges and Research Directions

Key challenges for continual fraud detection:

- **Label delay:** Fraud confirmation lags behind transactions, complicating drift detection
- **Balancing plasticity and stability:** Learning new patterns without forgetting old ones
- **Resource constraints:** Production systems require efficient incremental updates
- **Interpretability:** Understanding what changed when model updates occur

8.6. Few-Shot and Zero-Shot Learning

Few-shot learning enables fraud detection from very few labeled examples of new fraud types, while zero-shot learning detects fraud types never seen during training.

8.6.1. Meta-Learning for Fraud Detection

Meta-learning (learning to learn) trains models to quickly adapt to new tasks from limited data:

- **Model-Agnostic Meta-Learning (MAML):** Learn initialization enabling fast adaptation to new fraud types with gradient descent (16)
- **Prototypical networks:** Learn metric space where same-class examples cluster
- **Matching networks:** Attention-based classification using support set

8.6.2. Applications

Few-shot learning addresses fraud detection scenarios:

1. **Novel fraud types:** Detect new attack vectors from 1-10 examples
2. **Rare fraud categories:** Classify specific fraud types with limited historical occurrences
3. **Cross-domain transfer:** Adapt fraud detection to new institutions with minimal labeled data

8.6.3. Zero-Shot Learning

Zero-shot learning detects fraud types without direct training examples by:

- Learning fraud attributes rather than specific fraud types
- Transferring knowledge from related fraud categories
- Leveraging semantic descriptions of fraud characteristics

While promising, few-shot and zero-shot learning for fraud detection remain largely research topics with limited production deployment.

8.7. AutoML for Fraud Detection

Automated Machine Learning (AutoML) automates model selection, hyperparameter tuning, and feature engineering, reducing the expertise required for deploying effective fraud detection systems.

8.7.1. AutoML Components

AutoML for fraud detection encompasses:

1. **Algorithm selection:** Automatically choose optimal ML algorithms for dataset characteristics
2. **Hyperparameter optimization:** Systematically tune model hyperparameters using Bayesian optimization, genetic algorithms, or reinforcement learning
3. **Feature engineering:** Automatically generate and select relevant features
4. **Architecture search:** For deep learning, automatically design neural network architectures
5. **Ensemble construction:** Automatically build ensemble models

8.7.2. Applications to Fraud Detection

AutoML addresses fraud detection challenges:

- **Democratization:** Enables fraud detection deployment without ML expertise
- **Optimization:** Systematically explores model space, potentially finding better solutions than manual tuning
- **Efficiency:** Reduces time from data to deployment
- **Adaptability:** Automatically reconfigures models when data distributions change

Commercial AutoML platforms (Google AutoML, H2O.ai, DataRobot) have been applied to fraud detection with varying success. Challenges include:

- Handling extreme class imbalance requires domain-specific adaptations
- AutoML search spaces must include fraud-relevant algorithms and preprocessing
- Interpretability requirements may limit viable models
- Computational costs of extensive search

8.8. Blockchain for Fraud Prevention

Blockchain technology provides immutable, distributed ledgers that can enhance fraud prevention through transparency and traceability, particularly in supply chain applications.

8.8.1. Blockchain Applications

Blockchain addresses fraud prevention through:

- **Supply chain traceability:** Immutable records of product provenance prevent counterfeit goods
- **Smart contracts:** Automated execution of agreements reduces invoice fraud
- **Identity verification:** Decentralized identity systems prevent identity fraud
- **Transaction transparency:** All participants access same transaction history

Feng et al. surveyed blockchain applications for privacy and security, including fraud prevention (17). In supply chain contexts, blockchain enables:

1. Verification of product authenticity through provenance tracking
2. Detection of duplicate invoicing across supply chain partners
3. Prevention of unauthorized modifications to transaction records
4. Automated compliance checking through smart contracts

8.8.2. Limitations

Blockchain faces significant limitations for fraud detection:

- **Scalability:** Transaction throughput often insufficient for high-volume payment systems
- **Privacy concerns:** Public blockchains expose transaction details
- **Immutability drawbacks:** Inability to correct errors or comply with data deletion requirements
- **Oracle problem:** Blockchain cannot verify accuracy of external data inputs
- **Adoption barriers:** Requires coordination across multiple institutions

Blockchain is most promising for fraud *prevention* in supply chains and identity management rather than real-time transaction fraud *detection*.

8.9. Edge Computing and IoT for Fraud Detection

Edge computing performs fraud detection at the network edge (e.g., mobile devices, point-of-sale terminals) rather than centralized servers, enabling:

- **Reduced latency:** Local inference eliminates round-trip to cloud servers
- **Privacy preservation:** Sensitive data processed locally without transmission
- **Offline capability:** Fraud detection continues without network connectivity
- **IoT device integration:** Leverage device sensors for behavioral biometrics

8.9.1. Challenges

Edge fraud detection faces resource constraints:

- Limited computational power for complex models
- Memory constraints for model storage
- Power consumption considerations for mobile devices
- Model updates across distributed edge devices

Model compression techniques including quantization, pruning, and knowledge distillation enable deployment of fraud detection models on edge devices.

8.10. Comparative Analysis of Emerging Technologies

Table 5 summarizes emerging technologies and their maturity for fraud detection applications.

8.11. Integration and Hybrid Approaches

The most promising future fraud detection systems will integrate multiple emerging technologies:

- **Federated GNNs:** Graph neural networks trained via federated learning across institutions
- **Explainable continual learning:** Interpretable models that adapt to concept drift
- **Federated AutoML:** Automated model selection and tuning in privacy-preserving manner
- **Quantum-enhanced optimization:** Use quantum computers for hyperparameter optimization of classical models

Table 5. Emerging Technologies for Fraud Detection: Maturity and Potential

Technology	Maturity	Near-Term Potential	Key Advantage	Main Challenge
Quantum ML	Very Low	Low	Exponential speedup	Hardware limitations
Federated Learning	Medium	High	Privacy preservation	Non-IID data
Advanced GNNs	Medium	High	Network fraud detection	Scalability
Explainable AI	High	Very High	Regulatory compliance	Complexity trade-off
Continual Learning	Medium	High	Drift adaptation	Catastrophic forgetting
Few-Shot Learning	Low	Medium	Novel fraud detection	Limited examples
AutoML	High	Medium	Democratization	Domain adaptation
Blockchain	Medium	Medium	Fraud prevention	Scalability
Edge Computing	Medium	High	Low latency	Resource constraints

8.12. Research Priorities and Roadmap

Based on current maturity and potential impact, we recommend the following research priorities:

Short-term (1-3 years):

1. Explainable AI integration into production fraud detection systems
2. Federated learning pilots across financial institutions
3. Advanced GNN architectures for large-scale transaction networks
4. Continual learning systems for concept drift adaptation

Medium-term (3-5 years):

1. Few-shot learning for novel fraud type detection
2. Integrated federated learning and differential privacy
3. AutoML specialized for fraud detection with class imbalance
4. Blockchain integration with ML for supply chain fraud prevention

Long-term (5+ years):

1. Quantum machine learning algorithms with demonstrable advantages
2. Fully autonomous fraud detection systems with minimal human oversight
3. Cross-domain transfer learning across fraud types and industries
4. Self-evolving fraud detection systems through meta-learning

9. Challenges and Limitations in Fraud Detection

Despite significant advances in machine learning and deep learning approaches, fraud detection systems face persistent challenges that limit their effectiveness in real-world deployments. This section examines the fundamental challenges including extreme class imbalance, concept drift, interpretability requirements, privacy preservation, real-time processing constraints, data quality issues, and adversarial attacks. Understanding these challenges is essential for developing robust, practical fraud detection systems.

9.1. Extreme Class Imbalance

Class imbalance represents the most fundamental challenge in fraud detection, where fraudulent transactions typically constitute less than 1% of total transactions, and in some datasets, less than 0.1%. This severe imbalance causes multiple cascading problems that affect all aspects of fraud detection system design and evaluation.

9.1.1. Impact on Model Training

Extreme imbalance biases machine learning models toward the majority class:

- **Trivial classifier problem:** Models achieve high accuracy by predicting all transactions as legitimate, learning nothing about fraud patterns
- **Gradient domination:** In neural networks, gradients from abundant legitimate transactions overwhelm gradients from rare fraud cases, preventing effective learning of fraud representations
- **Decision boundary bias:** Classification boundaries are pushed toward fraud regions, requiring fraud cases to be extremely obvious to be detected
- **Convergence issues:** Optimization algorithms struggle to balance learning from both classes, often converging to suboptimal solutions

Moradi et al. identified class imbalance as a primary challenge across fraud detection domains, noting that the IEEE-CIS dataset contains only 3.5% fraud cases, while supply chain datasets often have fraud rates below 1% (25; 18; 24).

9.1.2. Mitigation Strategies and Their Limitations

Various strategies address class imbalance, each with inherent trade-offs:

1. Data-Level Approaches:

- **Random Oversampling:** Duplicates minority class samples
  - Advantage: Simple, preserves all information
  - Limitation: Overfitting to duplicated samples, increased training time

- **Random Undersampling:** Removes majority class samples
  - Advantage: Balanced dataset, faster training
  - Limitation: Information loss from discarded legitimate transactions
- **SMOTE (Synthetic Minority Oversampling):** Generates synthetic fraud samples through interpolation (19)
  - Advantage: Increases minority class diversity without duplication
  - Limitation: May generate unrealistic samples, especially in high dimensions
- **ADASYN (Adaptive Synthetic Sampling):** Focuses synthetic generation on harder-to-learn fraud examples (20)
  - Advantage: Emphasizes decision boundary regions
  - Limitation: Sensitive to noise, can amplify outliers
- **Borderline-SMOTE:** Generates synthetic samples only near class boundaries (21)
  - Advantage: Improves decision boundary learning
  - Limitation: May miss fraud patterns far from boundaries

Moradi et al. systematically evaluated these techniques on the IEEE-CIS dataset, finding that Borderline-SMOTE combined with ensemble methods achieved the best balance between fraud detection rate and false positive rate (25; 26).

2. Algorithm-Level Approaches:

- **Cost-Sensitive Learning:** Assigns higher misclassification costs to fraud
  - Advantage: Directly optimizes business-relevant objectives
  - Limitation: Requires accurate cost estimates, which are often unavailable
- **Class Weighting:** Weights loss function by inverse class frequency
  - Advantage: Easy to implement in most ML frameworks
  - Limitation: Optimal weights require tuning, may not generalize across datasets
- **Focal Loss:** Down-weights easy examples, focuses on hard negatives (22)
  - Advantage: Automatic emphasis on difficult examples
  - Limitation: Adds hyperparameters requiring tuning
- **One-Class Learning:** Learns boundary around legitimate transactions
  - Advantage: Does not require balanced training data

– Limitation: High false positive rates, limited precision

3. Ensemble-Based Approaches:
Ensemble methods specifically designed for imbalanced data have shown the most consistent success:

- **Balanced Random Forest:** Achieves balanced bootstrap sampling in each tree
- **EasyEnsemble:** Creates multiple balanced subsets through undersampling
- **RUSBoost:** Combines random undersampling with boosting
- **SMOTEBoost:** Applies SMOTE at each boosting iteration

These ensemble approaches, particularly when combined with advanced boosting algorithms like XGBoost and LightGBM, have achieved state-of-the-art performance on highly imbalanced fraud datasets (25; 26).

9.1.3. Evaluation Metric Challenges

Class imbalance renders standard accuracy metrics misleading:

- **Accuracy paradox:** 99% accuracy achieved by predicting all transactions as legitimate when fraud rate is 1%
- **Precision-recall trade-off:** Increasing fraud detection (recall) typically increases false positives (reduces precision)
- **AUC-ROC limitations:** Area Under ROC Curve gives equal weight to all threshold values, potentially misleading for extreme imbalance
- **AUC-PR preference:** Area Under Precision-Recall Curve more informative for imbalanced data, focusing on minority class performance

Moradi et al. emphasized the importance of using multiple metrics including AUC-ROC, AUC-PR, F1-score, precision, and recall to comprehensively evaluate fraud detection performance (25; 26).

9.1.4. Business Context Considerations

The optimal approach to class imbalance depends on business context:

- **Investigation capacity:** False positive rate must align with available fraud analyst resources
- **Fraud costs:** Cost of missed fraud vs. cost of false alarms determines optimal threshold
- **Customer experience:** High false positive rates frustrate legitimate customers
- **Regulatory requirements:** Some jurisdictions mandate specific detection rates

9.2. Concept Drift and Temporal Dynamics

Concept drift—the phenomenon where the statistical properties of fraud patterns change over time—poses a fundamental challenge for deployed fraud detection systems. Fraudsters continuously adapt their tactics to evade detection, causing models trained on historical data to degrade in performance.

9.2.1. Types of Concept Drift in Fraud Detection

Several drift patterns occur in fraud detection:

1. **Sudden drift:** Abrupt changes in fraud patterns
   - Example: New attack vector exploiting system vulnerability
   - Impact: Immediate sharp drop in detection rate
   - Response: Rapid model retraining or rule updates required
2. **Gradual drift:** Slow evolution of fraud tactics
   - Example: Fraudsters incrementally testing detection boundaries
   - Impact: Slow performance degradation over weeks/months
   - Response: Scheduled periodic retraining
3. **Incremental drift:** Continuous small changes in fraud distribution
   - Example: Seasonal variations in transaction patterns
   - Impact: Steady decline in model performance
   - Response: Online learning or sliding window approaches
4. **Recurring concepts:** Cyclical fraud patterns
   - Example: Holiday shopping fraud spikes
   - Impact: Periodic performance variations
   - Response: Ensemble models with concept-specific members

Lucas et al. extensively documented concept drift in credit card fraud detection, noting that fraud patterns can change dramatically within 6-12 month periods (23).

9.2.2. Challenges in Detecting and Adapting to Drift

Several factors complicate drift detection and adaptation:

- **Label delay:** Fraud confirmation lags transaction occurrence by days or weeks, preventing immediate drift detection
- **Distinguishing drift from noise:** Performance variations may reflect random fluctuations rather than true drift

- **Catastrophic forgetting:** Retraining on new data may cause models to forget previous fraud patterns that remain relevant
- **Computational constraints:** Frequent retraining of complex models (especially deep learning) computationally expensive
- **Model versioning:** Production systems require careful rollback capabilities when new models underperform

9.2.3. Adaptation Strategies and Trade-offs

Various strategies address concept drift, each with limitations:

- **Periodic retraining:** Retrain models on fixed schedule (weekly, monthly)
  - Advantage: Predictable computational requirements
  - Limitation: May miss rapid drift or waste computation during stable periods
- **Triggered retraining:** Retrain when drift detection signals activate
  - Advantage: Responsive to actual drift
  - Limitation: Requires robust drift detection, risks false alarms
- **Online learning:** Incrementally update models as new data arrives
  - Advantage: Continuous adaptation
  - Limitation: Risk of catastrophic forgetting, limited to certain algorithms
- **Ensemble with dynamic weighting:** Maintain multiple models and adjust weights based on recent performance
  - Advantage: Smooth transitions between concepts
  - Limitation: Increased inference latency and memory requirements
- **Sliding window:** Train only on recent data
  - Advantage: Focuses on current patterns
  - Limitation: Discards potentially useful historical information

9.2.4. Practical Considerations

Deployed fraud detection systems must balance adaptation responsiveness against stability:

- Too frequent updates: Risk instability, increased operational complexity
- Too infrequent updates: Performance degradation, missed emerging fraud patterns
- A/B testing: Essential for validating new models before full deployment
- Rollback procedures: Critical for rapid response to model failures

9.3. Interpretability and Explainability Requirements

The tension between model accuracy and interpretability represents a fundamental challenge in fraud detection. While complex models like deep neural networks and large ensembles often achieve superior performance, regulatory requirements, customer communication needs, and fraud analyst workflows demand explainable predictions.

9.3.1. Regulatory and Compliance Requirements

Multiple regulations mandate explainability for automated decisions:

- **GDPR Article 22:** Right to explanation for automated individual decisions
- **Fair Credit Reporting Act (FCRA):** Requires specific reasons for adverse credit actions
- **Equal Credit Opportunity Act (ECOA):** Mandates disclosure of reasons for credit denials
- **Model Risk Management Guidelines:** Banking regulators require understanding of model decision logic
- **Consumer Financial Protection Bureau:** Emphasizes transparency in algorithmic decision-making

These regulations create legal and operational requirements that constrain model choices, particularly for customer-facing decisions like transaction declines.

9.3.2. Operational Interpretability Needs

Beyond regulatory compliance, practical fraud detection operations require interpretability:

1. **Fraud analyst decision support:**
    - Analysts need to understand why transactions were flagged
    - Explanations guide investigation priorities
    - Understanding model logic builds analyst trust
2. **Customer communication:**
    - Declined transactions require explanations to customers
    - Vague explanations frustrate legitimate customers
    - Specific reasons enable corrective actions

3. **Model debugging and improvement:**

    - Interpretability reveals spurious correlations
    - Understanding failures guides feature engineering
    - Identifies biased or discriminatory patterns

4. **Adversarial robustness:**

    - Interpretability exposes potential gaming strategies
    - Helps identify features fraudsters might manipulate
    - Supports development of more robust models

9.3.3. The Accuracy-Interpretability Trade-off

A fundamental tension exists between model complexity and interpretability:

Highly Interpretable (Lower Accuracy):

- Logistic Regression: Coefficients directly interpretable
- Decision Trees: Rules directly understandable
- Rule-based systems: Explicit if-then logic

Moderately Interpretable (Moderate-High Accuracy):

- Random Forest: Feature importance, approximate rules
- Gradient Boosting: Feature importance, partial dependence plots
- Attention mechanisms: Attention weights indicate importance

Black Box (Highest Accuracy):

- Deep Neural Networks: Millions of parameters, opaque logic
- Large Ensembles: Combining many models obscures decision logic
- Complex stacking: Multiple layers of models

Hasan et al. demonstrated that explainable AI techniques can provide interpretability for complex models while maintaining high accuracy, partially resolving this trade-off (12).

9.3.4. Post-Hoc Explainability Limitations

While techniques like SHAP and LIME provide post-hoc explanations for black-box models, they face limitations:

- **Explanation instability:** Small input perturbations can yield different explanations
- **Computational cost:** Computing explanations for each prediction adds latency
- **Local vs. global:** Local explanations may not reflect overall model behavior
- **Approximation errors:** Explanations approximate true model logic imperfectly
- **Adversarial explanations:** Malicious actors may craft inputs to produce misleading explanations

Rudin argued for inherently interpretable models in high-stakes domains like fraud detection rather than relying on post-hoc explanations (13).

9.3.5. Domain-Specific Solutions

Different fraud detection contexts require different interpretability approaches:

- **Real-time transaction authorization:** Limited time for explanation generation; pre-computed feature importance sufficient
- **Fraud investigation support:** Detailed explanations valuable; latency less critical
- **Customer-facing declines:** Simple, non-technical explanations required
- **Regulatory reporting:** Comprehensive model documentation and audit trails necessary

9.4. Privacy and Data Protection

Fraud detection inherently requires processing sensitive personal and financial data, creating fundamental tensions with privacy regulations and consumer expectations. The challenge intensifies as data collection expands to include behavioral biometrics, device fingerprints, and cross-institutional data sharing.

9.4.1. Regulatory Landscape

Multiple regulations constrain fraud detection data practices:

1. **General Data Protection Regulation (GDPR):**
   - Purpose limitation: Data collected for fraud detection cannot be repurposed
   - Data minimization: Only necessary data should be collected
   - Right to erasure: Challenges model retraining when data must be deleted
   - Cross-border transfers: Restrictions on sharing fraud data across jurisdictions
2. **California Consumer Privacy Act (CCPA):**
   - Disclosure requirements for data collection and use
   - Opt-out rights for data selling
   - Enhanced consumer rights over personal information
3. **Payment Card Industry Data Security Standard (PCI DSS):**
   - Strict requirements for storing cardholder data
   - Encryption and access control mandates
   - Regular security assessments required
4. **Gramm-Leach-Bliley Act (GLBA):**
   - Financial privacy rule governing information sharing
   - Safeguards rule for protecting customer information

9.4.2. Privacy-Preserving Techniques and Limitations

Several techniques enable fraud detection while protecting privacy:

- **Differential Privacy:**
  – Adds calibrated noise to training data or model outputs
  – Advantage: Formal mathematical privacy guarantees
  – Limitation: Accuracy degradation, especially for rare events like fraud
- **Federated Learning:**
  – Trains models across institutions without sharing raw data
  – Advantage: Collaborative learning while maintaining data sovereignty

  - Limitation: Communication overhead, vulnerability to inference attacks

- **Homomorphic Encryption:**
  - Enables computation on encrypted data
  - Advantage: Strong privacy guarantees
  - Limitation: 100-1000x computational overhead, impractical for real-time systems

- **Secure Multi-Party Computation:**
  - Multiple parties jointly compute functions without revealing inputs
  - Advantage: Collaborative fraud detection without data sharing
  - Limitation: Complex protocols, significant computational cost

- **Synthetic Data Generation:**
  - Generate artificial fraud data preserving statistical properties
  - Advantage: Shareable datasets for research and development
  - Limitation: May not capture all fraud patterns, privacy leakage risks

Wei et al. analyzed trade-offs between privacy and utility in federated fraud detection, demonstrating that carefully calibrated differential privacy maintains detection performance while providing provable privacy guarantees (5).

9.4.3. Data Retention and Right to Erasure

GDPR's right to erasure creates operational challenges:

- **Model retraining:** Deleting individual records may require complete model retraining
- **Fraud evidence:** Tension between erasure rights and fraud investigation needs
- **Model unlearning:** Removing specific data's influence from trained models is technically challenging
- **Audit trails:** Compliance requires maintaining deletion records without retaining deleted data

9.4.4. Cross-Border Data Sharing

International fraud detection collaboration faces data sovereignty challenges:

- Different jurisdictions have incompatible privacy regulations
- Data residency requirements prevent cross-border model training
- International fraud rings exploit jurisdictional gaps
- Federated learning offers potential solution but faces regulatory uncertainty

9.5. Real-Time Processing Requirements

Many fraud detection applications, particularly payment authorization systems, require real-time predictions with stringent latency constraints. A credit card transaction authorization must complete within 100-200 milliseconds including fraud check, creating fundamental constraints on model complexity.

9.5.1. Latency Sources and Constraints

Total transaction latency comprises multiple components:

- **Network latency:** Transmission between point-of-sale and authorization server (20-50ms)
- **Feature extraction:** Computing transaction and customer features (10-30ms)
- **Model inference:** Fraud prediction computation (10-50ms target)
- **Decision logic:** Threshold application and routing (5-10ms)
- **Response transmission:** Approval/decline message return (20-50ms)

Total latency must remain under 100-200ms to avoid customer experience degradation and merchant system timeouts.

9.5.2. Model Complexity Trade-offs

Real-time requirements constrain model choices:

Fast Inference (¡10ms):

- Logistic Regression
- Small Decision Trees
- Linear SVM
- Limitation: Lower accuracy than complex models

Moderate Inference (10-50ms):

- Random Forest (100-500 trees)
- Gradient Boosting (XGBoost, LightGBM)
- Small Neural Networks
- Trade-off: Balance accuracy and speed

Slow Inference (¿50ms):

- Large Deep Neural Networks
- LSTM/Transformer models
- Large Ensemble Stacks
- Problem: Exceeds real-time constraints

9.5.3. Optimization Strategies

Several approaches reduce inference latency:

1. **Model Compression:**
   - Pruning: Remove less important neurons/trees
   - Quantization: Reduce numerical precision (float32 $\rightarrow$ int8)
   - Knowledge distillation: Train smaller model to mimic larger model
   - Trade-off: 10-30% speedup with 1-3% accuracy loss
2. **Feature Caching:**
   - Pre-compute customer behavior features
   - Cache frequently accessed aggregations
   - Trade-off: Reduced freshness of some features
3. **Tiered Architecture:**
   - Fast rule-based filter for obvious cases (legit or fraud)
   - Complex models only for ambiguous cases
   - Trade-off: Additional system complexity
4. **Hardware Acceleration:**
   - GPU inference for neural networks
   - Specialized inference chips (TPU, AWS Inferentia)
   - Trade-off: Infrastructure cost and complexity
5. **Approximate Algorithms:**
   - Early stopping in tree traversal
   - Approximate nearest neighbor search
   - Trade-off: Slightly reduced accuracy for significant speedup

9.5.4. Batch vs. Real-Time Trade-offs

Not all fraud detection requires real-time processing:

- **Real-time authorization:** Sub-100ms requirement, simple models
- **Post-transaction monitoring:** Minutes-hours acceptable, complex models viable
- **Fraud investigation:** Days acceptable, most sophisticated models usable
- **Hybrid approaches:** Real-time screening + delayed deep analysis

9.6. Data Quality and Availability

Fraud detection models are only as good as their training data. Several data quality issues systematically limit fraud detection effectiveness.

9.6.1. Label Quality Issues

Fraud labels suffer from multiple quality problems:

- **Label delay:** Fraud discovery lags transaction occurrence by days/weeks/months
- **Incomplete labeling:** Many frauds never detected or confirmed
- **Label noise:** Investigation errors cause mislabeling (false positives and false negatives)
- **Label bias:** Only investigated transactions receive labels, biasing toward detectable fraud
- **Evolving definitions:** Fraud criteria change over time, affecting label consistency

These issues make supervised learning challenging and motivate semi-supervised and unsupervised approaches (24).

9.6.2. Missing Data and Imputation

Transaction data frequently contains missing values:

- **Structural missingness:** Some features not applicable to all transactions (e.g., international transactions)
- **Collection failures:** Technical issues prevent data capture
- **Privacy-driven deletion:** Sensitive fields removed for compliance
- **Third-party data gaps:** External data sources have incomplete coverage

Imputation strategies involve trade-offs between bias (from incorrect imputation) and information loss (from discarding incomplete records). XGBoost and LightGBM's native handling of missing values provides advantage over methods requiring explicit imputation.

9.6.3. Feature Engineering Challenges

Effective fraud detection requires sophisticated feature engineering:

- **Temporal aggregations:** Transaction velocity and pattern features require aggregation over time windows
- **Categorical encoding:** High-cardinality features (merchant IDs) need special handling
- **Domain expertise:** Best features require fraud domain knowledge
- **Feature maintenance:** Features drift over time, requiring periodic re-engineering
- **Computational cost:** Complex features increase both training and inference time

9.6.4. Data Scarcity in Specific Domains

Supply chain fraud detection faces particularly acute data scarcity:

- Fraud confirmation requires extensive investigation
- Labeled datasets are proprietary and not publicly shared
- Fraud patterns vary significantly across organizations
- Limited benchmark datasets hinder research progress

Moradi et al. addressed this challenge through semi-supervised learning combining Isolation Forest pre-filtering with self-training SVM, achieving strong performance despite only 1.2% of data being labeled (24).

9.7. Adversarial Attacks and Model Robustness

Fraud detection operates in an adversarial environment where fraudsters actively attempt to evade detection. Unlike many ML applications where adversaries are hypothetical, fraud detection faces sophisticated, motivated attackers continuously probing for vulnerabilities.

9.7.1. Types of Adversarial Attacks

Several attack strategies target fraud detection systems:

1. **Evasion attacks:**
    - Fraudsters modify transaction characteristics to avoid detection
    - Gradually test detection boundaries through trial and error

- Exploit knowledge of detection rules or model logic
- Example: Keeping transaction amounts below detection thresholds

2. **Poisoning attacks:**
   - Inject malicious training data to corrupt model learning
   - Particularly relevant for online learning systems
   - Can cause false negatives (missed fraud) or false positives (legitimate transactions flagged)
3. **Model extraction:**
   - Query fraud detection system to reverse-engineer model
   - Use extracted model to craft evasive transactions
   - Relevant when fraudsters can observe system responses
4. **Oracle attacks:**
   - Systematically probe system with test transactions
   - Learn detection patterns through approved/declined responses
   - Particularly problematic in systems with immediate feedback

9.7.2. Defenses and Their Limitations

Various defenses mitigate adversarial threats:

- **Adversarial training:** Include adversarial examples in training data
  - Limitation: Requires anticipating attack strategies
- **Ensemble diversity:** Multiple models reduce single point of failure
  - Limitation: Sophisticated attackers may still find common vulnerabilities
- **Model opacity:** Hide model details from potential attackers
  - Limitation: Conflicts with interpretability requirements
- **Anomaly detection:** Detect systematic probing behavior
  - Limitation: Distinguishes probing from legitimate behavior
- **Delayed feedback:** Avoid immediate approval/decline signals
  - Limitation: Degrades customer experience

9.7.3. Arms Race Dynamics

Fraud detection exists in a perpetual arms race:

- Detection improvements drive fraudster adaptation
- New fraud techniques emerge continuously
- Defenders must anticipate future attacks, not just respond to past ones
- Sustainable solutions require continuous evolution

9.8. Cross-Cutting Challenges

Several challenges span multiple dimensions:

9.8.1. Evaluation Difficulties

Comprehensive evaluation faces multiple obstacles:

- **Proprietary data:** Best datasets unavailable for research
- **Non-stationarity:** Historical test sets don't reflect future performance
- **Metric selection:** Different stakeholders prioritize different metrics
- **Business context:** Technical metrics may not align with business value

9.8.2. Integration and Deployment

Transitioning from research to production faces challenges:

- **Infrastructure requirements:** Production systems need different architecture than research prototypes
- **A/B testing:** Carefully evaluating new models without risking fraud losses
- **Monitoring and alerting:** Detecting model failures in production
- **Regulatory approval:** New models require compliance review
- **Model governance:** Version control, documentation, audit trails

9.8.3. Organizational Challenges

Non-technical factors affect fraud detection success:

- **Siloed data:** Organizational barriers prevent data integration
- **Risk aversion:** Conservative institutions resist advanced techniques
- **Talent shortage:** Limited ML expertise in fraud domain
- **Investment justification:** Fraud prevention savings hard to quantify

Table 6. Summary of Key Challenges in Fraud Detection

Challenge	**Primary Impact**	**Current Mitigations**	**Remaining Gaps**
Extreme Class Imbalance	Model bias toward majority class	SMOTE, cost-sensitive learning, ensemble methods	Optimal strategy varies by context
Concept Drift	Performance degradation over time	Periodic retraining, online learning	Label delay complicates adaptation
Interpretability	Regulatory compliance, trust	SHAP, LIME, inherently interpretable models	Accuracy-interpretability trade-off
Privacy	Regulatory violations, limited data sharing	Federated learning, differential privacy	Accuracy loss, computational cost
Real-Time Latency	Limited model complexity	Model compression, tiered architecture	Accuracy sacrifice for speed
Data Quality	Unreliable labels, missing values	Semi-supervised learning, robust algorithms	Label delay unavoidable
Adversarial Attacks	Fraudster evasion	Adversarial training, ensemble diversity	Perpetual arms race

9.9. Summary of Key Challenges

Table 6 summarizes the primary challenges, their impacts, and current mitigation approaches.

These challenges motivate the research directions discussed in Section 12 and inform the comparative analysis in Section 11.

10. Benchmark Datasets and Evaluation Protocols

Robust evaluation of fraud detection algorithms requires high-quality benchmark datasets with realistic characteristics and standardized evaluation protocols. This section examines commonly used datasets across financial and supply chain domains, discusses their characteristics and limitations, reviews evaluation metrics, and proposes best practices for comprehensive performance assessment.

10.1. Overview of Fraud Detection Datasets

Fraud detection research faces a fundamental challenge: most real-world fraud datasets are proprietary and cannot be publicly shared due to privacy concerns, competitive advantages, and regulatory constraints. Consequently, the research community relies on a limited set of publicly available benchmark datasets, each with specific characteristics and limitations.

10.1.1. Dataset Categories

Fraud detection datasets can be categorized by domain and characteristics:

1. **Credit Card Transaction Datasets:**
    - Most common in research literature
    - High transaction volumes, extreme class imbalance
    - Examples: European Cardholders, IEEE-CIS, Synthetic datasets
2. **Banking and Payment Datasets:**
    - Account transactions, wire transfers, ACH payments
    - Often proprietary, limited public availability
3. **E-commerce Datasets:**
    - Online purchase fraud, account takeover
    - Rich behavioral features (clickstream, session data)
4. **Insurance Datasets:**
    - Claims fraud, application fraud
    - Complex investigation processes, delayed labels
5. **Supply Chain Datasets:**
    - Procurement fraud, invoice fraud, inventory discrepancies
    - Severe data scarcity, limited public datasets

10.2. Major Benchmark Datasets

10.2.1. European Credit Card Dataset (Kaggle)

The European Credit Card dataset represents one of the most widely used fraud detection benchmarks, released by the Machine Learning Group of Université Libre de Bruxelles (ULB) (92).

Dataset Characteristics:

- **Size:** 284,807 transactions
- **Time period:** September 2013 (2 days)
- **Features:** 30 features (28 PCA-transformed + Time + Amount)
- **Class distribution:** 492 frauds (0.172%)
- **Fraud amount:** Varies from small to large transactions

- **Availability:** Public (Kaggle, UCI ML Repository)

Key Features:

- Features V1-V28: PCA-transformed to preserve privacy
- Time: Seconds elapsed between transactions
- Amount: Transaction amount in Euros
- Class: Binary label (0 = legitimate, 1 = fraud)

Advantages:

- Widely used, enabling direct comparison across studies
- Real-world data with authentic fraud patterns
- Extreme class imbalance reflects realistic scenario
- Publicly available without restrictions

Limitations:

- PCA transformation prevents interpretability and domain-specific feature engineering
- Short time period (2 days) limits temporal analysis
- No concept drift observable within dataset
- Original features unknown, limiting generalization insights
- Dated (2013), may not reflect current fraud patterns

Usage in Research:

This dataset has been extensively used in fraud detection research. Moradi et al.'s systematic review identified it as the most frequently used benchmark, appearing in over 60% of reviewed credit card fraud detection studies (18). Reported performance on this dataset ranges from 94% to 99.98% accuracy, though such high accuracy reflects the extreme class imbalance and should be interpreted alongside precision, recall, and AUC metrics.

Representative results include:

- Random Forest: 99.95% accuracy, 0.978 AUC-ROC
- XGBoost: 99.96% accuracy, 0.982 AUC-ROC
- Deep Neural Networks: 99.2% accuracy, 0.971 AUC-ROC
- Autoencoder: 94.2% accuracy, 0.943 AUC-ROC

10.2.2. IEEE-CIS Fraud Detection Dataset

The IEEE Computational Intelligence Society (IEEE-CIS) fraud detection dataset was released as part of a Kaggle competition in 2019, providing a more recent and comprehensive fraud detection benchmark.

Dataset Characteristics:

- **Size:** 590,540 transactions (train), 506,691 (test)
- **Time period:** 6 months
- **Features:** 434 features across multiple categories
- **Class distribution:** 3.5% fraud rate (train set)
- **Source:** Vesta Corporation payment transactions
- **Availability:** Public (Kaggle)

Feature Categories:

- **Transaction features:** Amount, product code, card details
- **Identity features:** Device type, browser, email domain
- **Categorical features:** Card issuer, merchant category
- **Engineered features:** Time-based aggregations, velocity features
- **Missing values:** Extensive missingness (some features ¿90% missing)

Advantages:

- Interpretable features enable domain-specific analysis
- Longer time period allows temporal pattern analysis
- Rich feature set including identity and device information
- Realistic fraud rate (3.5%) closer to industry averages
- Recent data (2019) reflects modern fraud patterns
- Large scale enables deep learning approaches

Limitations:

- Extensive missing data requires careful handling
- High dimensionality (434 features) risks overfitting
- Some features have unclear semantics
- Competition format may lead to overfitting on public leaderboard

- No temporal ordering in test set prevents production-realistic evaluation

Usage in Research:

The IEEE-CIS dataset has become a primary benchmark for recent fraud detection research. Moradi et al. conducted comprehensive studies using this dataset, evaluating various ensemble methods and imbalance handling techniques (25; 26):

Reported Performance:

- **XGBoost:** 96.8% accuracy, 0.905 AUC-ROC, 0.874 AUC-PR
- **LightGBM:** 96.5% accuracy, 0.898 AUC-ROC, 0.867 AUC-PR
- **Random Forest:** 96.8% accuracy, 0.925 AUC-ROC, 0.847 F1-score
- **Stacking Ensemble:** 97.2% accuracy, 0.918 AUC-ROC, 0.891 AUC-PR, 0.847 F1-score

The stacking approach combining XGBoost, LightGBM, and Random Forest as base models with Logistic Regression as meta-learner achieved the best overall performance, demonstrating the effectiveness of ensemble methods on this challenging dataset (25; 26).

Key Findings from IEEE-CIS Studies:

- Borderline-SMOTE outperformed standard SMOTE and ADASYN for handling class imbalance
- Feature engineering significantly impacted performance, with velocity and aggregation features being crucial
- Ensemble methods consistently outperformed individual models by 2-5% across all metrics
- Careful hyperparameter tuning improved performance by 3-7%

10.2.3. Synthetic Financial Datasets (PaySim, BankSim)

Given the scarcity of real public fraud datasets, several synthetic datasets have been developed to support research.

PaySim:

- **Size:** 6.3 million transactions
- **Generation method:** Agent-based simulation of mobile money transactions
- **Fraud types:** Cash-out fraud, transfer fraud
- **Features:** Transaction type, amount, balance changes, time
- **Fraud rate:** Configurable, typically 0.13%

- **Source:** Based on real mobile money logs from Africa

BankSim:

- **Size:** 594,643 transactions
- **Generation method:** Agent-based simulation
- **Features:** Customer age, gender, merchant category, amount
- **Fraud rate:** 0.13%

Advantages of Synthetic Datasets:

- Configurable fraud rates and patterns
- No privacy concerns
- Interpretable features
- Large scale generation possible
- Ground truth labels by construction

Limitations of Synthetic Datasets:

- May not capture real-world fraud complexity
- Simplified fraud patterns compared to reality
- Models trained on synthetic data may not transfer to real fraud
- Limited adoption in research community

10.2.4. DataCo Smart Supply Chain Dataset

The DataCo Smart Supply Chain dataset provides a rare public benchmark for supply chain fraud detection.

Dataset Characteristics:

- **Size:** 180,519 orders
- **Features:** 53 features covering orders, products, customers, shipping
- **Fraud indicators:** Order status, delivery status, late delivery risk
- **Source:** E-commerce supply chain data
- **Availability:** Public (Kaggle, UCI)

Feature Categories:

- **Order features:** Order ID, date, status, region

- **Product features:** Category, price, description
- **Customer features:** Segment, location, demographics
- **Shipping features:** Mode, days, delivery status
- **Financial features:** Sales, profit, discount

Usage in Research:

Moradi et al. applied semi-supervised learning to supply chain fraud detection using this dataset (24). Their two-phase approach:

Phase 1 - Isolation Forest Pre-filtering:

- Identified 23.7% of transactions as potential anomalies
- Captured 94.2% of actual fraud in anomalous subset
- Reduced data volume for supervised learning by 76.3%

Phase 2 - Self-Training SVM:

- Applied to filtered subset with limited labels (1.2% labeled)
- Iteratively expanded labeled set through high-confidence predictions
- Final performance: F1-score 0.817, false positive rate ¡3.0%

This work demonstrated that semi-supervised approaches effectively address severe label scarcity in supply chain contexts where fraud investigation is expensive and time-consuming.

Advantages:

- Rare public supply chain fraud dataset
- Interpretable business features
- Realistic supply chain characteristics
- Multiple fraud indicators

Limitations:

- Smaller scale compared to credit card datasets
- Fraud labels may be noisy or incomplete
- Limited temporal information
- Supply chain fraud patterns may not generalize across industries

10.2.5. Insurance Fraud Datasets

Public insurance fraud datasets are extremely limited due to privacy and competitive concerns.

Available Datasets:

- **Automobile Insurance Claims:** Small datasets with claim characteristics and fraud labels
- **Medicare/Medicaid:** Limited public datasets, often aggregated to provider level
- **Synthetic Insurance:** Generated datasets mimicking claim patterns

Most insurance fraud research uses proprietary data, limiting reproducibility and comparison across studies.

10.3. Dataset Limitations and Biases

10.3.1. Temporal Bias

Most benchmark datasets represent snapshots rather than continuous time series:

- **No concept drift:** Static datasets don't capture fraud evolution
- **Temporal leakage:** Future information may leak into past predictions if not carefully split
- **Seasonal patterns:** Short time periods miss seasonal fraud variations
- **Non-stationary omission:** Real systems must handle distribution shifts absent in benchmarks

10.3.2. Label Bias

Fraud labels in real datasets suffer from multiple biases:

- **Detection bias:** Only detected frauds are labeled; undetected frauds misclassified as legitimate
- **Investigation bias:** Only investigated transactions receive labels, biasing toward high-value or suspicious cases
- **Confirmation bias:** Borderline cases may be labeled based on investigator preconceptions
- **Label delay:** Training on old data with complete labels, testing on recent data with incomplete labels

10.3.3. Privacy-Driven Limitations

Privacy requirements restrict dataset utility:

- **Feature obfuscation:** PCA transformation in European dataset prevents interpretability
- **Feature removal:** Sensitive features deleted, potentially including fraud indicators
- **Aggregation:** Fine-grained features aggregated, losing resolution
- **Synthetic substitution:** Real features replaced with synthetic equivalents, altering relationships

10.3.4. Scale and Diversity Limitations

Most public datasets have limited scale and diversity:

- Single institution/region rather than global coverage
- Limited fraud type diversity
- Smaller scale than production systems (millions vs. billions of transactions)
- Missing important data sources (device fingerprints, behavioral biometrics)

10.4. Evaluation Metrics

Appropriate evaluation metrics are critical for fraud detection given extreme class imbalance. This section reviews standard metrics, their interpretation, and suitability for fraud detection.

10.4.1. Confusion Matrix Foundation

All classification metrics derive from the confusion matrix:

- **True Positives (TP):** Correctly identified fraud
- **True Negatives (TN):** Correctly identified legitimate transactions
- **False Positives (FP):** Legitimate transactions incorrectly flagged as fraud
- **False Negatives (FN):** Fraud incorrectly classified as legitimate

10.4.2. Basic Classification Metrics

1. Accuracy:

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN} \tag{4}$$

- **Interpretation:** Overall proportion of correct predictions
- **Problem:** Misleading for imbalanced data; 99% accuracy trivial when fraud rate is 1%
- **Usage:** Should NOT be primary metric for fraud detection

2. Precision (Positive Predictive Value):

$$\text{Precision} = \frac{TP}{TP + FP} \tag{5}$$

- **Interpretation:** Proportion of fraud predictions that are correct
- **Importance:** High precision reduces false alarms and investigation costs
- **Business impact:** Low precision wastes analyst time and frustrates customers

3. Recall (Sensitivity, True Positive Rate):

$$\text{Recall} = \frac{TP}{TP + FN} \tag{6}$$

- **Interpretation:** Proportion of actual fraud that is detected
- **Importance:** High recall minimizes fraud losses
- **Business impact:** Low recall allows fraud to proceed uncaught

4. F1-Score (Harmonic Mean of Precision and Recall):

$$\text{F1} = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}} \tag{7}$$

- **Interpretation:** Balanced measure of precision and recall
- **Advantage:** Single metric summarizing both aspects
- **Usage:** Widely reported in fraud detection research

Moradi et al. reported F1-scores ranging from 0.817 (supply chain semi-supervised) to 0.847 (IEEE-CIS ensemble) across their studies (24; 25; 26).

5. Specificity (True Negative Rate):

$$\text{Specificity} = \frac{TN}{TN + FP} \tag{8}$$

- **Interpretation:** Proportion of legitimate transactions correctly classified
- **Importance:** High specificity ensures good customer experience

10.4.3. Threshold-Independent Metrics

Classification metrics depend on decision thresholds. Threshold-independent metrics provide more comprehensive evaluation.

1. Area Under ROC Curve (AUC-ROC):

The Receiver Operating Characteristic (ROC) curve plots True Positive Rate (Recall) vs. False Positive Rate (1 - Specificity) across all thresholds.

- **Range:** 0 to 1 (0.5 = random, 1.0 = perfect)
- **Interpretation:** Probability that a random fraud scores higher than a random legitimate transaction
- **Advantage:** Threshold-independent, widely understood
- **Limitation:** Gives equal weight to all thresholds, potentially misleading for extreme imbalance

Reported AUC-ROC Values:

- European dataset: 0.943–0.982 (typical range)
- IEEE-CIS dataset: 0.905–0.925 (25; 26)
- Supply chain: Limited reporting due to semi-supervised context

2. Area Under Precision-Recall Curve (AUC-PR):

The Precision-Recall curve plots Precision vs. Recall across all thresholds.

- **Advantage:** More informative for imbalanced datasets
- **Interpretation:** Focuses on minority class (fraud) performance
- **Baseline:** Random classifier achieves AUC-PR equal to fraud rate
- **Preference:** Should be primary metric for extreme imbalance

Reported AUC-PR Values:

- European dataset: 0.72–0.85 (lower than AUC-ROC due to imbalance)
- IEEE-CIS dataset: 0.874–0.891 (25; 26)

Moradi et al. emphasized AUC-PR alongside AUC-ROC in their ensemble evaluation, demonstrating that ensemble stacking improved AUC-PR from 0.874 to 0.891 compared to individual models (26).

10.4.4. Cost-Sensitive Metrics

Business-relevant evaluation should incorporate fraud and investigation costs.

Cost-Sensitive Accuracy:

$$\text{Cost} = C_{FP} \cdot FP + C_{FN} \cdot FN \tag{9}$$

where C_{FP} is the cost of false positives (investigation cost) and C_{FN} is the cost of false negatives (fraud loss).

- **Advantage:** Directly optimizes business objectives
- **Challenge:** Requires accurate cost estimates
- **Usage:** More common in industry than academic research

Matthews Correlation Coefficient (MCC):

$$\text{MCC} = \frac{TP \cdot TN - FP \cdot FN}{\sqrt{(TP + FP)(TP + FN)(TN + FP)(TN + FN)}} \tag{10}$$

- **Range:** -1 to 1 (0 = random, 1 = perfect)
- **Advantage:** Balanced metric even for extreme imbalance
- **Interpretation:** Correlation coefficient between predictions and true labels

10.5. Evaluation Protocols and Best Practices

10.5.1. Train-Test Splitting Strategies

Proper data splitting is critical for unbiased evaluation:

1. Random Splitting:

- **Method:** Randomly partition data (e.g., 80% train, 20% test)
- **Advantage:** Simple, ensures similar distributions
- **Problem:** Ignores temporal ordering, unrealistic for production
- **Temporal leakage risk:** Future information may leak into past

2. Temporal Splitting:

- **Method:** Train on earlier data, test on later data
- **Advantage:** Realistic evaluation of production performance
- **Requirement:** Mandatory for time-series fraud data
- **Challenge:** Concept drift may cause performance differences

3. Stratified Splitting:

- **Method:** Preserve class distribution in train and test sets
- **Advantage:** Ensures sufficient fraud cases in both sets
- **Usage:** Recommended for random splitting scenarios

10.5.2. Cross-Validation Strategies

Cross-validation provides more robust performance estimates:

1. K-Fold Cross-Validation:

- Split data into k folds, train on k-1, test on 1, repeat k times
- Average results across folds
- Problem: Temporal leakage if data has temporal structure

2. Stratified K-Fold:

- Preserves class distribution in each fold
- Recommended for imbalanced fraud data
- Still problematic for temporal data

3. Time-Series Cross-Validation:

- Train on expanding or sliding window of past data
- Test on subsequent time period
- Respects temporal ordering
- Most realistic for fraud detection evaluation

10.5.3. Handling Class Imbalance in Evaluation

Evaluation must account for class imbalance:

- **Report multiple metrics:** Accuracy alone insufficient; report precision, recall, F1, AUC-ROC, AUC-PR
- **Emphasize minority class:** Focus on fraud detection performance
- **Confusion matrix:** Report full confusion matrix for transparency
- **Threshold analysis:** Show performance across multiple thresholds
- **Business metrics:** Report investigation workload and fraud losses at different thresholds

10.5.4. Statistical Significance Testing

Rigorous comparison requires statistical testing:

- **Multiple runs:** Report mean and standard deviation across multiple train-test splits or cross-validation folds
- **Significance tests:** Apply paired t-tests or Wilcoxon signed-rank tests
- **Confidence intervals:** Report 95% confidence intervals for metrics
- **Bonferroni correction:** Adjust for multiple comparisons when comparing many methods

10.5.5. Reproducibility Requirements

Fraud detection research should adhere to reproducibility standards:

- **Random seeds:** Report and fix random seeds for reproducible results
- **Hyperparameters:** Document all hyperparameters and tuning procedures
- **Preprocessing:** Describe all data preprocessing steps
- **Code availability:** Provide code when possible (respecting privacy constraints)
- **Dataset versions:** Specify exact dataset versions used

10.6. Benchmark Performance Summary

Table 7 summarizes typical performance ranges for major fraud detection algorithms on benchmark datasets.

Table 7. Performance Summary on Major Benchmark Datasets

Algorithm	European Dataset AUC-ROC	IEEE-CIS Dataset AUC-ROC / AUC-PR	Evaluation Focus	Key Studies
Logistic Regression	0.91–0.93	0.88–0.90 / 0.82–0.85	Baseline	Multiple
Random Forest	0.97–0.98	0.92–0.93 / 0.84–0.87	Strong baseline	(25)
XGBoost	0.98–0.98	0.90–0.91 / 0.87–0.88	Top performer	(25)
LightGBM	0.97–0.98	0.89–0.90 / 0.86–0.87	Efficient	Multiple
Deep Neural Network	0.94–0.97	0.88–0.90 / N/A	Complex patterns	Multiple
LSTM	0.89–0.92	N/A / N/A	Temporal	Multiple
Autoencoder	0.94–0.95	N/A / N/A	Unsupervised	Multiple
Ensemble Stacking	0.98–0.99	0.92 / 0.89	Best overall	(26)

Key Observations:

- Ensemble methods consistently achieve best performance
- IEEE-CIS more challenging than European dataset (lower scores across all methods)
- AUC-PR substantially lower than AUC-ROC, reflecting class imbalance
- Deep learning competitive but not consistently superior to gradient boosting

10.7. Dataset Recommendations

Based on research experience and dataset characteristics, we provide recommendations for dataset selection:

For Credit Card Fraud Research:

- **Primary:** IEEE-CIS dataset (interpretable features, realistic scale)
- **Comparison:** European dataset (enables comparison with extensive prior work)
- **Avoid:** Sole reliance on European dataset due to PCA limitations

For Supply Chain Fraud Research:

- **Primary:** DataCo dataset (rare public supply chain benchmark)
- **Approach:** Semi-supervised methods appropriate given label scarcity (24)
- **Need:** More public supply chain fraud datasets urgently needed

For Algorithm Development:

- Use multiple datasets to demonstrate generalization
- Include both credit card and non-credit domains
- Report results on both old (European) and recent (IEEE-CIS) benchmarks

For Evaluation Methodology:

- Report multiple metrics: AUC-ROC, AUC-PR, F1-score, precision, recall
- Use stratified k-fold cross-validation or temporal splitting
- Include statistical significance testing
- Provide full confusion matrices
- Document all preprocessing and hyperparameters

10.8. Future Dataset Needs

The fraud detection research community would benefit from:

1. **Temporal datasets:** Long time periods (1+ years) capturing concept drift
2. **Multi-domain datasets:** Covering banking, e-commerce, insurance, supply chain
3. **Privacy-preserving sharing:** Techniques enabling wider data sharing while protecting privacy

4. **Synthetic data generation:** Improved synthetic datasets capturing real fraud complexity

5. **Evaluation platforms:** Standardized evaluation frameworks and leaderboards

6. **Cross-institutional datasets:** Federated datasets enabling collaborative research

11. Comparative Analysis Across Methods and Domains

This section provides a comprehensive comparative analysis of fraud detection methodologies across multiple dimensions including algorithmic approach, application domain, performance characteristics, computational requirements, and practical deployment considerations. We synthesize findings from previous sections to guide algorithm selection for specific fraud detection scenarios.

11.1. Performance Comparison by Algorithm Family

11.1.1. Traditional Machine Learning Methods

Traditional ML algorithms demonstrate varying performance characteristics across fraud detection tasks:

Table 8. Traditional ML Performance Characteristics

Algorithm	Accuracy	Speed	Interpretability	Best Use Case
Logistic Regression	Moderate	Fast	High	Baseline, regulated
Decision Trees	Moderate	Fast	High	Rule extraction
Random Forest	High	Moderate	Moderate	General purpose
SVM	High	Slow	Low	Small datasets
Naive Bayes	Moderate	Very Fast	High	Real-time screening
KNN	Moderate	Slow	Low	Rarely recommended

Key Findings:

Random Forest emerges as the most balanced traditional ML approach, offering:

- Consistently high accuracy (96-99% on benchmark datasets)
- Reasonable training and inference speed
- Natural handling of mixed feature types
- Feature importance for interpretability

- Robust performance without extensive tuning

Logistic Regression serves as an effective baseline and remains valuable for:

- Regulatory environments requiring full interpretability
- Real-time systems with extreme latency constraints (¡10ms)
- Scenarios where coefficient interpretability guides business rules
- Initial model development before advancing to complex methods

SVM shows strong performance on small to medium datasets but faces scalability challenges on large-scale fraud detection applications processing millions of transactions.

11.1.2. Deep Learning Methods

Deep learning approaches offer powerful representation learning but with trade-offs:

Table 9. Deep Learning Performance Characteristics

Architecture	**Accuracy**	**Training Time**	**Data Need**	**Best Use Case**
Feedforward NN	Moderate-High	Moderate	Moderate	Feature learning
CNN	Moderate-High	Moderate	Moderate	Spatial patterns
LSTM/GRU	High	Slow	High	Sequential data
Autoencoder	Moderate	Moderate	Low	Unsupervised
Transformer	High	Very Slow	Very High	Long sequences

Key Findings:
LSTM networks excel when temporal dependencies are critical:

- Transaction sequence modeling
- User behavior trajectory analysis
- Velocity and recency pattern detection
- Achieves 89-92% AUC-ROC on sequential fraud tasks

Autoencoders provide unique advantages for unsupervised fraud detection:

- Train on legitimate transactions only
- Detect novel fraud patterns through reconstruction error
- Effective when labeled fraud examples are scarce

- 94% accuracy on European dataset without fraud labels during training

Deep learning limitations in fraud detection:

- Typically does not outperform optimized gradient boosting on tabular data
- Requires more data and longer training time
- Black-box nature complicates regulatory compliance
- Higher computational costs for training and inference

11.1.3. Ensemble Methods

Ensemble methods consistently achieve state-of-the-art performance across fraud detection benchmarks:

Table 10. Ensemble Method Performance Characteristics

Method	Accuracy	Robustness	Tuning Effort	Best Use Case
Random Forest	High	High	Low	General purpose
XGBoost	Very High	High	Moderate	Competition performance
LightGBM	Very High	High	Moderate	Large-scale, fast
CatBoost	Very High	High	Low	Categorical features
Stacking	Very High	Very High	High	Maximum performance

Performance Hierarchy:
Based on benchmark results across multiple studies, ensemble methods rank as follows:
Tier 1 - Best Performance (AUC-ROC 0.91-0.93):

1. Stacking ensembles combining XGBoost, LightGBM, Random Forest
2. Optimized XGBoost with careful hyperparameter tuning
3. LightGBM with categorical feature handling

Tier 2 - Strong Performance (AUC-ROC 0.88-0.91):

1. Random Forest with balanced sampling
2. Gradient Boosting with default parameters
3. Voting ensembles

Tier 3 - Baseline Performance (AUC-ROC 0.85-0.88):

1. Single decision trees

2. AdaBoost
3. Basic bagging

Key Insights:

The performance gap between Tier 1 and Tier 2 is relatively small (2-5% across metrics), but can translate to significant business impact:

- 2% AUC improvement may reduce fraud losses by 10-15%
- 3% precision improvement reduces investigation costs by 20-30%
- Diminishing returns above 95% AUC-ROC

Practical recommendations:

- Start with XGBoost or LightGBM for rapid development
- Invest in stacking for production systems requiring maximum performance
- Use Random Forest for interpretability-accuracy balance
- Consider CatBoost when categorical features dominate

11.1.4. Semi-Supervised and Unsupervised Methods

These approaches address scenarios with limited labeled data:

Table 11. Semi-Supervised/Unsupervised Method Characteristics

Method	Label Need	Novel Fraud	FP Rate	Best Use Case
Isolation Forest	None	Excellent	High	Pre-filtering
One-Class SVM	None	Good	High	Small datasets
Autoencoder	None	Excellent	Moderate	Anomaly detection
Self-Training	Minimal	Good	Moderate	Label expansion
Two-Phase Hybrid	Minimal	Excellent	Low	Supply chain

Performance in Data-Scarce Scenarios:

Two-phase approaches combining unsupervised pre-filtering with semi-supervised refinement demonstrate superior performance:

Example: Supply Chain Fraud Detection

- Phase 1 (Isolation Forest): Reduces data volume 76%, captures 94% fraud
- Phase 2 (Self-Training SVM): Achieves F1 0.817 with ¡3% FP rate
- Requires only 1.2% labeled data
- Outperforms supervised methods trained on same labeled subset

This approach is particularly valuable for:

- Supply chain fraud (expensive labeling)
- Insurance fraud (delayed label confirmation)
- Novel fraud type detection
- Cross-domain transfer scenarios

11.2. Domain-Specific Performance Analysis

11.2.1. Credit Card Fraud Detection

Credit card fraud detection represents the most mature fraud detection application with extensive research and production deployment.

Optimal Approaches:

1. **Primary:** XGBoost or LightGBM with SMOTE/Borderline-SMOTE
    - AUC-ROC: 0.90-0.92 on IEEE-CIS
    - Training time: 2-5 hours on full dataset
    - Inference: ¡20ms per transaction
2. **Alternative:** Random Forest with balanced sampling
    - AUC-ROC: 0.92-0.93 on IEEE-CIS
    - Slightly slower inference (30-50ms)
    - Better feature importance interpretability
3. **Maximum Performance:** Stacking ensemble
    - AUC-ROC: 0.91-0.92, AUC-PR: 0.89-0.89
    - 2-3% improvement over single models
    - Increased complexity and maintenance burden

Key Success Factors:

- Feature engineering: Velocity, recency, aggregation features critical
- Imbalance handling: Borderline-SMOTE outperforms alternatives
- Temporal validation: Time-based splits essential for realistic evaluation
- Real-time constraints: Model complexity limited by ¡100ms requirement

Deep Learning Performance:

- LSTM networks: 89-92% AUC-ROC on sequential data
- Competitive but not superior to optimized gradient boosting
- Higher computational cost
- Valuable when temporal patterns are complex and abundant data available

11.2.2. Banking and Payment Fraud

Banking fraud encompasses diverse transaction types including wire transfers, ACH payments, and check fraud.

Optimal Approaches:

- Ensemble methods (XGBoost, LightGBM) for general transaction fraud
- LSTM networks for account takeover detection (sequential login patterns)
- Graph Neural Networks for detecting fraud rings in transfer networks
- Hybrid rule-based + ML for high-value wire transfer screening

Unique Challenges:

- Multiple fraud types require specialized models
- Regulatory requirements mandate interpretability
- Real-time authorization for some transaction types
- Batch processing acceptable for others (wire transfers)

11.2.3. E-Commerce Fraud

E-commerce fraud detection benefits from rich behavioral data including clickstream, device fingerprints, and session patterns.

Optimal Approaches:

1. **Session-based:** LSTM or GRU for clickstream analysis
    - Captures browsing patterns distinguishing fraud from legitimate purchases
    - 94-96% accuracy on session-based detection
2. **Transaction-level:** XGBoost or LightGBM
    - Combines transaction, user, and device features
    - Handles high-cardinality categorical features (product IDs, merchant categories)
3. **Graph-based:** GNN for fraud ring detection
    - Models user-product-merchant interaction networks
    - Identifies coordinated attacks through network structure
    - 96-97% AUC-ROC on e-commerce fraud networks

Key Success Factors:

- Behavioral features more predictive than transaction amount alone
- Device fingerprinting highly effective for account takeover detection
- Velocity features (purchases per hour) critical
- IP geolocation mismatches strong fraud indicator

11.2.4. Insurance Fraud

Insurance fraud detection faces unique challenges including diverse fraud types, long investigation cycles, and complex claim structures.

Optimal Approaches:

- Random Forest for claims fraud (interpretability important for investigations)
- Network analysis + ML for fraud ring detection in staged accidents
- Anomaly detection (Isolation Forest, Autoencoders) for novel fraud patterns
- Natural Language Processing for claim text analysis

Unique Challenges:

- Label delay: Fraud confirmation may take months or years
- Complex relationships: Medical providers, claimants, witnesses
- Domain expertise critical: Medical, legal, automotive knowledge required
- Interpretability mandatory: Explanations needed for claim denials

11.2.5. Supply Chain Fraud

Supply chain fraud detection confronts severe data scarcity and labeling challenges.

Optimal Approaches:

1. **Two-phase framework:** Unsupervised pre-filtering + semi-supervised refinement
   - Isolation Forest reduces investigation volume by 76%
   - Self-training SVM achieves F1 0.817 with minimal labels
   - Outperforms purely supervised approaches
2. **Anomaly detection:** Isolation Forest, LOF for procurement anomalies
3. **Rule-based + ML:** Combine business rules with ML scoring
4. **Graph analysis:** Vendor-buyer relationship networks

Key Success Factors:

- Semi-supervised methods essential given label scarcity
- Domain knowledge for feature engineering (price deviations, vendor patterns)

- Invoice matching and duplicate detection
- Blockchain integration for supply chain transparency

Performance Characteristics:

- Lower precision acceptable given investigation capacity
- High recall critical to minimize fraud losses
- F1-score 0.80-0.85 considered strong performance
- False positive rate ¡5% typically acceptable

11.3. Computational Requirements Analysis

11.3.1. Training Time Comparison

Training time varies dramatically across algorithms and dataset sizes:

Table 12. Training Time Comparison (IEEE-CIS Scale Dataset)

Algorithm	CPU Time	GPU Speedup	Memory
Logistic Regression	2-5 min	1.5x	Low
Random Forest	30-60 min	1.2x	Moderate
XGBoost	2-5 hours	3-5x	Moderate
LightGBM	1-3 hours	2-3x	Low
Deep Neural Network	3-8 hours	10-20x	High
LSTM	6-12 hours	15-30x	High
Stacking Ensemble	5-10 hours	Varies	High

Key Observations:

- LightGBM offers best training speed among high-performance algorithms
- Deep learning benefits most from GPU acceleration
- Stacking requires sequential training of base models then meta-model
- Training time grows sub-linearly with dataset size for tree-based methods

11.3.2. Inference Latency Comparison

Real-time fraud detection systems require sub-100ms predictions:

Recommendations for Real-Time Systems:

- Use LightGBM or XGBoost for best latency-accuracy balance

Table 13. Inference Latency per Transaction

Algorithm	Latency (ms)	Throughput (TPS)	Meets Real-Time
Logistic Regression	1-3	100,000+	Yes
Random Forest (100 trees)	5-10	10,000-20,000	Yes
Random Forest (500 trees)	20-40	2,000-5,000	Yes
XGBoost (1000 trees)	10-20	5,000-10,000	Yes
LightGBM (1000 trees)	5-15	10,000-20,000	Yes
Neural Network	15-30	3,000-6,000	Yes
LSTM	30-60	1,000-3,000	Marginal
Stacking (3 base + meta)	25-50	2,000-4,000	Yes

- Limit tree depth and number for extreme low-latency requirements
- Consider tiered architecture: fast filter + complex model for ambiguous cases
- Model compression techniques reduce latency 20-40% with minimal accuracy loss
- LSTM inference may exceed real-time constraints for some applications

11.3.3. Scalability Analysis

Algorithm scalability to large datasets and high transaction volumes:

Table 14. Scalability Characteristics

Algorithm	Data Scale	Feature Scale	Parallelization
Logistic Regression	Excellent	Excellent	Good
Random Forest	Good	Good	Excellent
XGBoost	Good	Good	Good
LightGBM	Excellent	Excellent	Good
SVM	Poor	Poor	Poor
Deep Neural Network	Good	Excellent	Excellent
LSTM	Moderate	Good	Poor

Key Insights:

- LightGBM excels at large-scale fraud detection (millions of transactions)
- SVM limited to small-medium datasets (¡1M transactions)

- Random Forest parallelizes easily across trees
- LSTM sequential processing limits scalability
- Deep learning scales well with data volume but requires more data for convergence

11.4. Interpretability-Accuracy Trade-off Analysis

The tension between model interpretability and predictive accuracy represents a fundamental challenge in fraud detection.

11.4.1. Interpretability Spectrum

Algorithms can be positioned on an interpretability-accuracy spectrum:

High Interpretability, Moderate Accuracy:

- Logistic Regression: Coefficients directly interpretable
- Decision Trees: Rules explicitly understandable
- Naive Bayes: Probabilistic reasoning transparent

Moderate Interpretability, High Accuracy:

- Random Forest: Feature importance, approximate rules
- Gradient Boosting: Feature importance, partial dependence
- Attention mechanisms: Attention weights indicate focus

Low Interpretability, Highest Accuracy:

- Deep Neural Networks: Millions of opaque parameters
- Ensemble stacking: Multiple layers obscure logic
- Complex LSTM architectures: Hidden states difficult to interpret

11.4.2. Domain-Specific Interpretability Requirements

Different fraud detection contexts have varying interpretability needs:

11.4.3. Post-Hoc Explainability Solutions

When complex models are necessary, post-hoc explanation methods provide interpretability:

SHAP (SHapley Additive exPlanations):

- Advantages: Theoretically grounded, consistent feature attribution

Table 15. Interpretability Requirements by Context

Context	**Interpretability Need**	**Acceptable Methods**	**Rationale**
Transaction Decline	High	LR, Trees, RF	Customer explanation
Investigation Support	Moderate	RF, XGBoost, NN+SHAP	Analyst guidance
Risk Scoring	Moderate	Any with explanations	Threshold decisions
Regulatory Reporting	High	LR, Trees, RF	Compliance audit
Research & Development	Low	Any method	Performance priority

- Limitations: Computationally expensive (10-100ms per explanation)
- Use case: Investigation support, not real-time authorization

LIME (Local Interpretable Model-agnostic Explanations):

- Advantages: Fast, intuitive local explanations
- Limitations: Unstable, approximation quality varies
- Use case: Quick explanations for fraud analysts

Feature Importance:

- Advantages: Fast, global understanding of model
- Limitations: Does not explain individual predictions
- Use case: Model debugging, feature engineering guidance

11.5. Cross-Domain Transferability

Understanding which techniques transfer across fraud detection domains guides research and development priorities.

11.5.1. Highly Transferable Techniques

Certain approaches demonstrate consistent effectiveness across domains:

1. **Gradient Boosting (XGBoost, LightGBM):**
    - Effective across credit card, banking, e-commerce, insurance
    - Consistent top-tier performance with minimal domain-specific tuning
    - Handles diverse feature types (numerical, categorical, temporal)
2. **Random Forest:**

- Robust performance across all fraud domains
- Feature importance generalizes well
- Less sensitive to hyperparameters than other methods

3. **Ensemble Stacking:**
   - Achieves best performance across multiple domains
   - Combines complementary algorithm strengths
   - Requires more development effort but consistently delivers

4. **SMOTE and Variants:**
   - Borderline-SMOTE effective across domains
   - Improves recall consistently
   - Requires careful validation to avoid overfitting

11.5.2. Domain-Specific Techniques

Some approaches excel in specific domains but transfer poorly:

- **LSTM Networks:**
  - Excellent: Credit card transaction sequences, e-commerce clickstream
  - Poor: Insurance claims (limited sequential structure)
  - Moderate: Banking (depends on transaction frequency)
- **Graph Neural Networks:**
  - Excellent: E-commerce fraud rings, social security fraud networks
  - Poor: Supply chain (limited network structure in some cases)
  - Moderate: Credit card (requires constructing meaningful graphs)
- **Text Analysis (NLP):**
  - Excellent: Insurance claims description analysis
  - Poor: Credit card (minimal text data)
  - Moderate: E-commerce (product descriptions, reviews)

11.5.3. Transfer Learning Opportunities

Transfer learning enables knowledge transfer across domains:

Successful Transfer Scenarios:

- Credit card fraud → Debit card fraud (similar transaction patterns)
- Bank A fraud → Bank B fraud (institution transfer)
- E-commerce fraud → Mobile payment fraud (similar behavioral patterns)

Challenging Transfer Scenarios:

- Credit card fraud → Insurance fraud (fundamentally different patterns)
- E-commerce fraud → Supply chain fraud (different data structures)
- Consumer fraud → Commercial/B2B fraud (different scales and patterns)

11.6. Cost-Benefit Analysis

Selecting fraud detection approaches requires balancing performance gains against implementation and operational costs.

11.6.1. Development Cost Comparison

Table 16. Development Cost Factors

Approach	Expertise Required	Development Time	Tuning Effort	Total Cost
Logistic Regression	Low	1-2 weeks	Low	Low
Random Forest	Low	2-3 weeks	Low	Low
XGBoost/LightGBM	Moderate	3-4 weeks	Moderate	Moderate
Deep Neural Network	High	6-8 weeks	High	High
LSTM	High	8-12 weeks	High	Very High
Stacking Ensemble	High	6-10 weeks	High	High
Two-Phase Hybrid	Moderate	4-6 weeks	Moderate	Moderate

11.6.2. Operational Cost Comparison

11.6.3. ROI Analysis Framework

Return on investment for fraud detection improvements:

Performance Improvement Value:

Table 17. Operational Cost Factors

Approach	Infrastructure Cost	Retraining Frequency	Monitoring Complexity	Total Cost
Logistic Regression	Low	Weekly	Low	Low
Random Forest	Low	Weekly	Low	Low
XGBoost/LightGBM	Moderate	Weekly	Moderate	Moderate
Deep Neural Network	High (GPU)	Bi-weekly	High	High
LSTM	High (GPU)	Bi-weekly	High	Very High
Stacking Ensemble	Moderate	Weekly	High	Moderate-High

- 1% AUC–ROC improvement ≈ 5–10% fraud loss reduction
- 1% precision improvement ≈ 10–15% investigation cost reduction
- 1% recall improvement ≈ 5–8% fraud loss reduction

Cost Components:

- Development: One-time cost (ML engineer salary × development time)
- Infrastructure: Ongoing cost (server/GPU costs)
- Maintenance: Ongoing cost (retraining, monitoring, updates)
- Opportunity cost: Delayed deployment vs. simpler faster solution

Decision Framework:
When Simple Methods Suffice:

- Fraud losses moderate (¡$1M annually)
- Investigation capacity limited
- Regulatory requirements demand interpretability
- Technical expertise limited
- Rapid deployment critical

→ Recommendation: Random Forest or XGBoost
When Complex Methods Justified:

- Fraud losses significant (¿$10M annually)
- Ample investigation capacity
- Performance gains translate to substantial savings
- Technical expertise available
- Adequate development timeline

→ Recommendation: Stacking ensemble or deep learning

11.7. Recommendation Framework

This framework guides algorithm selection based on scenario characteristics:

11.7.1. Quick Reference Decision Tree

Step 1: Assess Data Availability

- **Abundant labeled data (¿100K fraud examples):** → Supervised learning
- **Limited labels (¡10K fraud examples):** → Semi-supervised or ensemble
- **Very limited labels (¡1K fraud examples):** → Two-phase unsupervised + semi-supervised
- **No labels:** → Anomaly detection (Isolation Forest, Autoencoder)

Step 2: Assess Latency Requirements

- **¡10ms:** → Logistic Regression, small Random Forest
- **10-50ms:** → XGBoost, LightGBM, Random Forest
- **50-100ms:** → Neural Networks, LSTM (limited)
- **¿100ms (batch):** → Any method, optimize for accuracy

Step 3: Assess Interpretability Requirements

- **Critical (regulatory/customer-facing):** → Logistic Regression, Decision Trees, Random Forest
- **Important (analyst tools):** → Random Forest, XGBoost with SHAP
- **Less critical (risk scoring):** → Any method with post-hoc explanations
- **Not required (research):** → Any method

Step 4: Assess Data Characteristics

- **Tabular only:** → XGBoost, LightGBM, Random Forest
- **Sequential/temporal:** → LSTM, GRU, or XGBoost with temporal features
- **Network structure:** → Graph Neural Networks
- **Text data:** → NLP + ensemble
- **Images:** → CNN + ensemble

Step 5: Assess Resource Constraints

- **Limited budget/expertise:** → Random Forest (easiest to deploy)
- **Moderate resources:** → XGBoost or LightGBM
- **Ample resources:** → Stacking ensemble or deep learning
- **GPU available:** → Consider deep learning

11.7.2. Domain-Specific Recommendations

Credit Card Fraud:

- Primary: XGBoost or LightGBM with Borderline-SMOTE
- Alternative: Stacking ensemble for maximum performance
- Emerging: LSTM for transaction sequence modeling

E-Commerce Fraud:

- Primary: XGBoost for transaction-level
- Secondary: LSTM for session-based detection
- Advanced: GNN for fraud ring detection

Supply Chain Fraud:

- Primary: Two-phase Isolation Forest + Self-Training SVM
- Alternative: Anomaly detection + expert review
- Avoid: Purely supervised methods (insufficient labels)

Insurance Fraud:

- Primary: Random Forest (interpretability + accuracy)
- Secondary: Network analysis for fraud rings
- Emerging: NLP for claims text analysis

11.8. Summary and Key Takeaways

Universal Best Practices:

1. Start with ensemble methods (XGBoost, LightGBM, Random Forest) as baseline
2. Always address class imbalance explicitly (SMOTE variants, class weighting)
3. Use multiple evaluation metrics (AUC-ROC, AUC-PR, F1, precision, recall)
4. Validate on temporal splits for realistic performance estimates
5. Consider total cost of ownership, not just model accuracy

Algorithm Selection Hierarchy (General Guidance):

1. **Default choice:** XGBoost or LightGBM (balanced performance, speed, ease of use)
2. **Maximum accuracy:** Stacking ensemble (2-5% improvement, higher complexity)
3. **Interpretability priority:** Random Forest (good accuracy + feature importance)
4. **Limited labels:** Two-phase unsupervised + semi-supervised
5. **Sequential data:** LSTM or temporal XGBoost features
6. **Network data:** Graph Neural Networks

Domain-Specific Insights:

- Credit card: Mature domain, ensemble methods dominate, 95%+ AUC-ROC achievable
- E-commerce: Behavioral data crucial, LSTM + GNN promising
- Supply chain: Label scarcity drives semi-supervised approaches
- Insurance: Interpretability critical, network analysis valuable

Future Trends:

- Federated learning for cross-institutional collaboration
- AutoML for automated model selection and tuning
- Explainable AI for regulatory compliance
- Continual learning for concept drift adaptation
- Hybrid approaches combining multiple paradigms

12. Future Research Directions and Emerging Trends

This section identifies promising research directions, emerging trends, and open problems in fraud detection. We analyze technological trajectories, methodological gaps, and practical challenges that warrant future investigation, providing a roadmap for researchers and practitioners advancing the field.

12.1. Advanced Machine Learning Paradigms

12.1.1. Self-Supervised Learning for Fraud Detection

Self-supervised learning has revolutionized computer vision and natural language processing but remains underexplored in fraud detection. This paradigm learns representations from unlabeled data through pretext tasks, potentially addressing data scarcity challenges.

Promising Directions:

- **Contrastive learning:** Learn representations by contrasting similar and dissimilar transactions
  - Positive pairs: Same customer's legitimate transactions
  - Negative pairs: Different customers' transactions
  - Hypothesis: Fraud transactions will have distinct embedding patterns
- **Masked transaction modeling:** Predict masked transaction features from context
  - Similar to BERT's masked language modeling
  - Learn transaction patterns without labels
  - Fine-tune on limited labeled fraud data
- **Temporal contrastive learning:** Exploit temporal relationships
  - Positive pairs: Temporally adjacent transactions from same user
  - Learn normal behavioral trajectories
  - Detect deviations indicating fraud

Research Challenges:

- Designing effective pretext tasks for tabular transaction data
- Balancing pretraining on legitimate transactions vs. including fraud
- Evaluating learned representations' quality for fraud detection
- Computational cost of pretraining on large transaction datasets

Expected Impact:

- Improved performance when labeled fraud data is scarce
- Better transfer learning across institutions and fraud types
- More robust representations capturing fraud-relevant patterns
- 10-15% performance improvement over supervised baselines in low-data regimes

12.1.2. Few-Shot and Zero-Shot Fraud Detection

Novel fraud types emerge continuously, often with few or no initial labeled examples. Few-shot and zero-shot learning enable rapid adaptation to new fraud patterns.

Meta-Learning Approaches:

- **Model-Agnostic Meta-Learning (MAML):**
  - Train model initialization enabling fast adaptation
  - Learn from multiple fraud types to generalize to new types
  - Adapt to novel fraud with 1-10 examples through few gradient steps
- **Prototypical Networks:**
  - Learn metric space where fraud types cluster
  - Classify new fraud type by distance to prototype
  - Requires careful design of distance metrics for transaction data
- **Matching Networks:**
  - Attention-based classification using support set
  - Learn similarity function between transactions
  - Applicable when new fraud type has distinct characteristics

Zero-Shot Learning via Attributes:

- Define fraud attributes: velocity anomaly, amount anomaly, location anomaly, time anomaly
- Train classifiers for each attribute on known fraud types
- Detect novel fraud through attribute combinations
- Example: New fraud type exhibits high velocity + location anomaly

Research Challenges:

- Limited fraud type diversity in current datasets
- Defining meaningful fraud attributes generalizing across types
- Balancing generalization to new fraud vs. maintaining performance on known fraud
- Validating approaches requires datasets with diverse fraud types

Expected Impact:

- Rapid response to emerging fraud tactics (hours vs. weeks)
- Reduced dependence on large labeled datasets
- Cross-institutional fraud type transfer
- 40-60% performance improvement over zero-shot baselines when adapting to novel fraud with ¡10 examples

12.1.3. Continual and Lifelong Learning

Fraud patterns evolve continuously, requiring models that learn new patterns while retaining knowledge of previous patterns without catastrophic forgetting.

Key Research Directions:

- **Elastic Weight Consolidation (EWC):**
  - Selectively slow learning on parameters important for previous fraud patterns
  - Regularize based on Fisher information matrix
  - Balance plasticity (learning new patterns) vs. stability (remembering old patterns)
- **Progressive Neural Networks:**
  - Allocate new network columns for new fraud patterns
  - Preserve previous columns frozen
  - Lateral connections transfer knowledge from old to new patterns
- **Memory-Augmented Networks:**
  - External memory stores representative fraud examples
  - Replay during training on new data
  - Prevents forgetting through experience replay
- **Dynamic Architecture Methods:**
  - Expand model capacity as needed for new patterns
  - Prune unused parameters to maintain efficiency
  - Balance model size vs. performance

Fraud-Specific Considerations:

- Concept drift may involve both new fraud patterns and changes to normal behavior
- Some old fraud patterns disappear (no longer needed to remember)
- Extreme class imbalance complicates experience replay
- Real-time systems cannot afford frequent full retraining

Research Challenges:

- Detecting when new patterns emerge vs. noise in data
- Deciding when old patterns are obsolete and can be forgotten

- Handling label delay in continual learning (fraud confirmation lags transaction)
- Balancing model capacity growth vs. inference latency

Expected Impact:

- Sustained high performance over extended periods (years) without full retraining
- Reduced operational costs through incremental updates
- Faster adaptation to concept drift (days vs. weeks)
- 10-20% performance improvement vs. static models over 12-month period

12.2. Enhanced Federated and Privacy-Preserving Learning

12.2.1. Federated Learning Advances

Current federated learning approaches face challenges in fraud detection that future research must address.

Non-IID Data Handling:

- **Problem:** Fraud patterns vary dramatically across institutions
- **Solutions:**
  - Personalized federated learning: Global model + institution-specific adaptation
  - Clustered federated learning: Group similar institutions
  - Meta-learning for fast institution-specific adaptation
  - Adaptive aggregation weighting based on data similarity

Secure Aggregation Improvements:

- Current secure aggregation adds 10-100x communication overhead
- Research directions:
  - Gradient compression techniques reducing communication by 90%+
  - Sparsification: Transmit only significant gradient updates
  - Quantization: Reduce precision of gradient values
  - Top-k gradient selection: Send only largest k gradients

Byzantine-Robust Aggregation:

- **Problem:** Malicious participants may poison global model
- **Solutions:**

  - Robust aggregation methods (median, trimmed mean, Krum)
  - Outlier detection on gradient updates
  - Reputation systems for participants
  - Differential privacy to limit impact of individual contributions

Vertical Federated Learning for Fraud:

- Different institutions hold different features for same customers
- Example: Bank has transaction history, credit bureau has credit scores
- Research challenges:
  - Secure feature alignment without revealing identities
  - Privacy-preserving entity resolution
  - Efficient training across vertically partitioned data

Expected Impact:

- Industry-wide fraud detection consortia become practical
- Small institutions benefit from collective knowledge
- 15-25% performance improvement vs. institution-specific models
- Faster detection of emerging fraud patterns through information sharing

12.2.2. Differential Privacy with Utility Preservation

Differential privacy provides formal privacy guarantees but often degrades model performance. Future research must minimize this trade-off.

Adaptive Privacy Budgets:

- Allocate privacy budget proportional to sample informativeness
- Spend more budget on rare fraud examples
- Spend less on abundant legitimate transactions
- Dynamic budget allocation based on model uncertainty

Privacy-Utility Trade-off Optimization:

- Multi-objective optimization of privacy and accuracy
- Pareto frontier exploration
- Application-specific privacy requirements
- Example: Higher privacy for PII fields, lower for aggregated statistics

Synthetic Data Generation with Privacy:

- Generate synthetic fraud datasets with differential privacy
- Enable public sharing and research
- Research challenges:
  - Maintaining rare fraud pattern fidelity
  - Preserving complex feature correlations
  - Validating synthetic data utility for fraud detection

Expected Impact:

- Differential privacy becomes practical for fraud detection
- ¡5% accuracy loss vs. non-private models
- Enables compliant data sharing and collaborative research
- Reduces regulatory and legal risks

12.3. Explainability and Trustworthy AI

12.3.1. Causal Inference for Fraud Detection

Current fraud detection models identify correlations but not causal relationships. Causal inference can improve model robustness and interpretability.

Causal Discovery:

- Discover causal relationships among transaction features
- Identify which features causally contribute to fraud
- Example: Does unusual transaction amount *cause* fraud or merely correlate?
- Methods: Constraint-based algorithms, score-based methods, structural equation models

Counterfactual Explanations:

- "What would need to change for this transaction to be classified as legitimate?"
- Provides actionable explanations for customers and fraud analysts
- Reveals model decision boundaries
- Research challenges:
  - Ensuring counterfactuals are realistic and actionable

  - Maintaining privacy when generating counterfactuals
  - Computational efficiency for real-time systems

Intervention-Based Fraud Prevention:

- Use causal models to design fraud prevention interventions
- Predict impact of policy changes before implementation
- Example: Will adding two-factor authentication reduce account takeover fraud?

Expected Impact:

- More robust models resistant to spurious correlations
- Better explanations supporting regulatory compliance
- Improved fraud prevention strategies through causal understanding
- 5-10% reduction in false positives through causal reasoning

12.3.2. Uncertainty Quantification

Fraud detection models should express confidence in predictions, enabling better decision-making.

Bayesian Deep Learning:

- Probability distributions over model parameters
- Quantifies model uncertainty
- Higher uncertainty for novel fraud patterns
- Enables risk-aware decision-making

Conformal Prediction:

- Provides prediction sets with guaranteed coverage
- Example: "With 95% confidence, fraud probability is between 0.7 and 0.9"
- Model-agnostic approach
- Useful for setting investigation thresholds

Ensemble Uncertainty:

- Disagreement among ensemble members indicates uncertainty
- High disagreement suggests ambiguous cases requiring human review
- Low disagreement enables automated decisions

Expected Impact:

- Better allocation of fraud analyst resources to uncertain cases
- Reduced false positives through confidence-aware thresholding
- Improved customer experience by avoiding uncertain declines
- 10-15% reduction in manual review workload

12.3.3. Fairness and Bias Mitigation

Fraud detection models must avoid discriminatory bias while maintaining effectiveness.

Bias Detection and Measurement:

- Demographic parity: Equal fraud detection rates across groups
- Equalized odds: Equal true positive and false positive rates
- Predictive parity: Equal precision across groups
- Research challenge: Defining appropriate fairness metrics for fraud detection

Bias Mitigation Techniques:

- **Pre-processing:** Remove biased features, reweight samples
- **In-processing:** Add fairness constraints during training
- **Post-processing:** Adjust decision thresholds by group

Fraud-Specific Fairness Challenges:

- Legitimate fraud rate differences across demographics complicate fairness
- Protected attributes may be predictive of fraud in some contexts
- False positives have disparate impact (legitimate customers inconvenienced)
- Regulatory requirements vary by jurisdiction

Expected Impact:

- Reduced discriminatory bias in fraud detection
- Improved compliance with anti-discrimination regulations
- Better customer trust and satisfaction
- Minimal accuracy loss (¡2%) while achieving fairness

12.4. Scalability and Real-Time Processing

12.4.1. Edge-Based Fraud Detection

Processing fraud detection at the edge (point-of-sale devices, mobile phones) enables lower latency and privacy preservation.

Model Compression for Edge Deployment:

- **Quantization:** Reduce model precision (float32 → int8)
  - 4x model size reduction
  - 2-4x inference speedup
  - ¡1% accuracy loss with careful quantization
- **Pruning:** Remove unnecessary model parameters
  - 50-90% parameter reduction possible
  - Structured pruning maintains hardware efficiency
  - Iterative pruning preserves accuracy
- **Knowledge Distillation:** Train small model to mimic large model
  - 10-100x parameter reduction
  - Student model learns from teacher predictions
  - Maintains 90-95% of teacher accuracy
- **Neural Architecture Search:** Automatically design efficient models
  - Optimize for accuracy and inference speed jointly
  - Hardware-aware search for specific edge devices
  - Promising but computationally expensive

Federated Edge Learning:

- Train fraud detection models across distributed edge devices
- Privacy-preserving by design (data never leaves device)
- Challenges: Limited edge device computational power, communication costs

Expected Impact:

- ¡10ms inference latency through local processing
- Enhanced privacy through on-device computation
- Reduced cloud infrastructure costs
- Offline fraud detection capability

12.4.2. Streaming Fraud Detection

Real-world fraud detection must process continuous transaction streams, not static batches.

Online Learning Algorithms:

- Incremental updates as transactions arrive
- No separate training and inference phases
- Adapts continuously to concept drift
- Challenges: Catastrophic forgetting, label delay, computational efficiency

Stream Processing Architectures:

- Apache Kafka for transaction ingestion
- Apache Flink/Spark Streaming for feature computation
- Real-time model serving (TensorFlow Serving, MLflow)
- Monitoring and alerting infrastructure

Incremental Feature Engineering:

- Compute velocity and aggregation features incrementally
- Sliding window aggregations
- Approximate algorithms for efficiency (Count-Min Sketch, HyperLogLog)
- Trade-off: Approximation accuracy vs. computational cost

Expected Impact:

- True real-time fraud detection (transaction-by-transaction)
- Faster adaptation to emerging fraud patterns
- Reduced batch processing delays
- 20-30% improvement in fraud detection speed

12.5. Domain-Specific Innovations

12.5.1. Graph-Based Supply Chain Fraud Detection

Supply chains form complex networks ideal for graph-based analysis.

Temporal Graph Neural Networks:

- Model evolving supply chain relationships over time
- Detect anomalous changes in vendor-buyer relationships
- Identify fraud propagation through supply chain tiers
- Research challenges: Scalability to large supply chain graphs, temporal modeling efficiency

Multi-Modal Graph Learning:

- Integrate transaction graphs with communication networks
- Combine structured data (invoices) with unstructured (emails, documents)
- Detect fraud through cross-modal pattern analysis

Blockchain Integration:

- Immutable transaction records
- Smart contracts for automated compliance
- GNN-based analysis of blockchain transaction patterns
- Detect fraudulent smart contract interactions

Expected Impact:

- 25-35% improvement in supply chain fraud detection
- Detection of coordinated fraud across supply chain tiers
- Reduced false positives through relationship modeling
- Enhanced fraud prevention through blockchain transparency

12.5.2. Behavioral Biometrics for Authentication Fraud

Behavioral patterns provide continuous authentication complementing traditional fraud detection.

Behavioral Signals:

- **Typing patterns:** Keystroke dynamics, typing speed, error patterns
- **Mouse movements:** Movement patterns, click patterns, hesitations
- **Touch interactions:** Pressure, swipe patterns, finger size
- **Device usage:** App usage patterns, navigation sequences
- **Gait analysis:** Walking patterns from mobile accelerometer data

Deep Learning for Behavioral Modeling:

- LSTM/Transformer models for sequential behavioral patterns
- Siamese networks for user verification
- Anomaly detection on behavioral embeddings
- Continuous authentication without explicit user action

Privacy Considerations:

- Behavioral data highly sensitive
- Requires consent and transparency
- Privacy-preserving processing techniques necessary
- Balance security benefits vs. privacy concerns

Expected Impact:

- 40-60% reduction in account takeover fraud
- Seamless continuous authentication
- Detection of bot-driven fraud attacks
- Reduced reliance on knowledge-based authentication

12.5.3. Multi-Modal Fraud Detection

Integrating diverse data modalities can improve fraud detection comprehensiveness.

Data Modalities:

- **Structured:** Transaction amounts, timestamps, categories
- **Text:** Product descriptions, customer reviews, claim narratives
- **Images:** Product photos, receipt images, ID documents
- **Audio:** Call center conversations, voice authentication
- **Video:** Surveillance footage, facial recognition
- **Network:** User-merchant graphs, social networks

Multi-Modal Architectures:

- Specialized encoders for each modality
- Fusion mechanisms: Early fusion, late fusion, cross-attention
- Joint embedding spaces for cross-modal reasoning
- Transfer learning from pre-trained modality-specific models

Research Challenges:

- Handling missing modalities (not all transactions have all data types)
- Modality alignment and synchronization
- Computational costs of multi-modal processing
- Privacy implications of collecting diverse data types

Expected Impact:

- 15-25% improvement through cross-modal evidence
- Detection of sophisticated fraud using multiple deception channels
- More robust against single-modality attacks
- Enhanced fraud investigation with multi-modal evidence

12.6. AutoML and Neural Architecture Search

12.6.1. Fraud-Specific AutoML

General-purpose AutoML requires adaptation for fraud detection's unique challenges.

Imbalance-Aware AutoML:

- Search space includes sampling techniques (SMOTE, ADASYN, undersampling)
- Optimization targets imbalance-appropriate metrics (AUC-PR, F1)
- Automatically tune class weights and sample weights
- Joint optimization of preprocessing and model architecture

Latency-Constrained Architecture Search:

- Multi-objective optimization of accuracy and inference latency
- Hardware-aware search for production deployment targets
- Pareto frontier exploration enabling accuracy-latency trade-offs
- Reduces manual model selection burden

Few-Shot AutoML:

- Meta-learning across fraud detection tasks
- Rapid adaptation to new institutions or fraud types
- Requires ¡100 labeled examples for effective model discovery

Expected Impact:

- Democratization of advanced fraud detection (reduced ML expertise requirement)
- 10-20% performance improvement over manual design
- 5-10x reduction in development time
- Automated adaptation as fraud patterns evolve

12.7. Quantum Machine Learning Applications

12.7.1. Near-Term Quantum Applications

While large-scale quantum computers remain years away, near-term quantum devices may offer advantages for specific fraud detection subproblems.

Variational Quantum Classifiers:

- Parameterized quantum circuits for classification
- Hybrid quantum-classical optimization
- Potential advantages in high-dimensional feature spaces
- Current limitations: ¡100 qubits, high error rates

Quantum Kernel Methods:

- Quantum computers evaluate kernel functions
- Classical SVM training on quantum kernel matrix
- Potentially exponentially large feature spaces
- Research question: Do quantum kernels help fraud detection?

Quantum Sampling for Generative Models:

- Quantum advantage in sampling from complex distributions
- Generate synthetic fraud data
- Bayesian inference for fraud probability

Realistic Timeline and Expectations:

- 2025-2030: Small-scale demonstrations on toy problems
- 2030-2035: Potential advantages on specific fraud detection subproblems
- 2035+: Practical quantum fraud detection systems possible
- Challenges: Error correction, scalability, algorithm development

Expected Impact (Long-Term):

- Exponential speedup for certain fraud detection tasks
- Enhanced security through quantum-resistant cryptography
- New fraud detection algorithms impossible on classical computers
- Timeline highly uncertain, conservative planning necessary

12.8. Integration and Systems Research

12.8.1. End-to-End Fraud Detection Systems

Most research focuses on isolated model development, but production systems require comprehensive integration.

Research Gaps:

- Data pipelines: Real-time feature engineering at scale
- Model serving: Low-latency inference infrastructure
- Model monitoring: Detecting performance degradation and concept drift
- A/B testing: Safely evaluating new models in production
- Model governance: Versioning, documentation, compliance
- Human-in-the-loop: Analyst interfaces, feedback incorporation

MLOps for Fraud Detection:

- Continuous training pipelines
- Automated retraining triggers
- Feature store management
- Model registry and deployment automation
- Performance monitoring and alerting

Expected Impact:

- Reduced time from research to production deployment (months → weeks)
- Improved model reliability and robustness
- Faster iteration and experimentation
- Better collaboration between data scientists and engineers

12.8.2. Human-AI Collaboration

Effective fraud detection requires optimal collaboration between AI systems and human analysts.

Research Directions:

- **Active learning:** AI selects most informative cases for human review
- **Explainable predictions:** AI provides actionable explanations guiding investigation

- **Human feedback loops:** Analyst corrections improve model
- **Confidence-aware automation:** High-confidence cases automated, low-confidence routed to humans
- **Workload optimization:** Balance caseload across analysts considering expertise

Expected Impact:

- 30-50% improvement in analyst efficiency
- Better model performance through expert feedback
- Reduced analyst burnout through optimized workload
- Faster fraud investigation and resolution

12.9. Cross-Cutting Research Priorities

12.9.1. Benchmark Dataset Development

The fraud detection research community needs improved datasets.

Required Characteristics:

- **Temporal coverage:** Multiple years to study concept drift
- **Fraud type diversity:** Multiple fraud types in single dataset
- **Multi-domain:** Credit card, banking, e-commerce, insurance, supply chain
- **Rich features:** Interpretable features enabling domain analysis
- **Realistic scale:** Millions of transactions
- **Privacy-compliant:** Synthetic or anonymized data

Potential Solutions:

- Improved synthetic data generation maintaining fraud pattern complexity
- Federated evaluation frameworks enabling model comparison without data sharing
- Privacy-preserving data sharing techniques (differential privacy, secure enclaves)
- Industry-academic partnerships for dataset creation

12.9.2. Reproducibility and Standardization

Fraud detection research suffers from reproducibility challenges.

Needed Standards:

- Standardized evaluation protocols and metrics
- Benchmark implementations of algorithms
- Shared codebases and experiment tracking
- Reporting guidelines for fraud detection research
- Pre-registration of evaluation protocols

Community Infrastructure:

- Public leaderboards with strict evaluation protocols
- Open-source fraud detection libraries
- Shared computing resources for large-scale experiments
- Workshops and challenges driving standardization

12.10. Timeline and Roadmap

12.10.1. Short-Term (1-3 Years)

High Priority:

- Explainable AI integration into production systems
- Federated learning pilot deployments
- Advanced ensemble methods optimization
- Continual learning for concept drift
- Edge-based fraud detection

Expected Outcomes:

- 5-10% performance improvement over current baselines
- Improved regulatory compliance through explainability
- Reduced latency through edge processing
- Better adaptation to evolving fraud patterns

12.10.2. Medium-Term (3-7 Years)

High Priority:

- Self-supervised learning for fraud detection
- Few-shot learning for novel fraud types
- Causal inference integration
- Multi-modal fraud detection
- Automated MLOps pipelines

Expected Outcomes:

- 15-25% performance improvement over current methods
- Rapid adaptation to new fraud types (days vs. weeks)
- Production deployment time reduced 10x
- Enhanced privacy through federated and differential privacy techniques

12.10.3. Long-Term (7+ Years)

Visionary Goals:

- Quantum machine learning applications
- Fully autonomous fraud detection systems
- Cross-domain universal fraud detection models
- Human-level fraud investigation capabilities
- Proactive fraud prevention (predict and prevent before occurrence)

Expected Outcomes:

- 50%+ reduction in fraud losses
- Near-elimination of false positives through advanced AI
- Seamless, invisible fraud protection for customers
- Self-evolving systems requiring minimal human oversight

12.11. Key Takeaways and Recommendations

For Researchers:

1. Focus on practical challenges: label scarcity, concept drift, interpretability
2. Develop fraud-specific techniques, not just apply generic ML
3. Collaborate with industry for realistic problem formulations and datasets
4. Prioritize reproducibility through code and data sharing
5. Address fairness and bias concerns proactively

For Practitioners:

1. Invest in MLOps infrastructure for rapid iteration
2. Adopt explainable AI for regulatory compliance and analyst trust
3. Explore federated learning for cross-institutional collaboration
4. Begin edge computing pilot projects for latency reduction
5. Prepare for continual learning to handle concept drift

For Policymakers:

1. Support development of privacy-preserving fraud detection techniques
2. Encourage industry-academic partnerships and data sharing
3. Develop appropriate regulations balancing fraud prevention and privacy
4. Fund public benchmark dataset creation
5. Promote fairness and bias mitigation in fraud detection systems

Critical Success Factors:

- Collaboration between academia, industry, and regulators
- Investment in infrastructure and datasets
- Focus on practical deployability, not just academic metrics
- Responsible AI development addressing fairness, privacy, transparency
- Continuous evolution matching fraud sophistication

13. Conclusion

This comprehensive survey has examined the landscape of machine learning approaches for fraud detection across financial and supply chain domains. We synthesize key findings, provide actionable recommendations for researchers and practitioners, and reflect on the field's trajectory toward more effective, interpretable, and privacy-preserving fraud detection systems.

13.1. Summary of Key Findings

13.1.1. Algorithmic Performance Landscape

Our comprehensive analysis across traditional machine learning, deep learning, ensemble methods, and semi-supervised approaches reveals several critical insights:

Ensemble Methods Dominate:

- XGBoost, LightGBM, and Random Forest consistently achieve top performance
- Stacking ensembles provide 2-5% improvement over individual models
- Performance hierarchy: Stacking ¿ XGBoost/LightGBM ¿ Random Forest ¿ Traditional ML ¿ Baseline
- Ensemble robustness justifies slightly higher computational costs in production

Deep Learning Shows Promise but Not Superiority:

- Deep neural networks competitive but rarely superior to optimized gradient boosting on tabular fraud data
- LSTM networks excel when temporal dependencies are critical (89-92% AUC-ROC on sequential tasks)
- Autoencoders effective for unsupervised fraud detection with limited labels (94% accuracy)
- Deep learning justified primarily for sequential data, images, text, or extreme scale

Context Determines Optimal Approach:

- No universally optimal algorithm across all fraud detection scenarios
- Data characteristics, latency requirements, interpretability needs, and resource constraints dictate algorithm selection
- Credit card fraud: XGBoost/LightGBM with SMOTE variants

- Supply chain fraud: Two-phase unsupervised + semi-supervised (F1 0.817, ¡3% FPR)
- E-commerce fraud: Hybrid LSTM + XGBoost + GNN
- Insurance fraud: Random Forest for interpretability

13.1.2. Persistent Challenges

Despite significant advances, fundamental challenges remain:

Extreme Class Imbalance:

- Fraud rates typically ¡1%, sometimes ¡0.1%
- Borderline-SMOTE emerged as most effective sampling technique
- Cost-sensitive learning and ensemble methods partially mitigate imbalance
- Remaining challenge: Optimal balance between precision and recall varies by application

Concept Drift:

- Fraud patterns evolve continuously as fraudsters adapt
- Models degrade 10-30% within 6-12 months without retraining
- Label delay complicates drift detection and adaptation
- Continual learning and online learning offer promising directions but remain underexplored

Interpretability-Accuracy Trade-off:

- Complex models achieve best performance but lack transparency
- Regulatory requirements and customer communication demand explainability
- Post-hoc explanation methods (SHAP, LIME) partially address black-box models
- Inherently interpretable models (Random Forest) sacrifice 2-5% accuracy for transparency

Data Scarcity and Quality:

- Limited labeled fraud examples, especially in supply chain and insurance domains
- Label delay: Fraud confirmation lags transactions by days to months

- Semi-supervised learning effectively addresses scarcity (two-phase approach achieves F1 0.817 with 1.2% labeled data)
- Public benchmark datasets remain limited, hampering research progress

Privacy and Compliance:

- GDPR, CCPA, and other regulations constrain data collection and usage
- Cross-institutional data sharing legally and technically challenging
- Federated learning and differential privacy enable privacy-preserving collaboration
- Performance degradation from privacy mechanisms remains significant (5-15%)

13.1.3. Emerging Technologies

Several emerging technologies show promise for next-generation fraud detection:

Near-Term Impact (1-3 Years):

- **Explainable AI:** SHAP and LIME integration for regulatory compliance
- **Federated Learning:** Cross-institutional collaboration without data sharing
- **Graph Neural Networks:** Fraud ring detection in network-structured data (96-97% AUC-ROC)
- **Edge Computing:** Sub-10ms inference through local processing

Medium-Term Impact (3-7 Years):

- **Self-Supervised Learning:** Leveraging unlabeled data for representation learning
- **Few-Shot Learning:** Rapid adaptation to novel fraud types with ¡10 examples
- **Continual Learning:** Sustained performance over years without catastrophic forgetting
- **Multi-Modal Integration:** Combining transaction, text, image, and network data

Long-Term Vision (7+ Years):

- **Quantum Machine Learning:** Potential exponential speedups for specific subproblems
- **Autonomous Systems:** Self-evolving fraud detection with minimal human oversight
- **Proactive Prevention:** Predicting and preventing fraud before occurrence

13.2. Recommendations for Researchers

13.2.1. High-Priority Research Directions

Based on our comprehensive analysis, we recommend researchers focus on:

1. Addressing Data Scarcity:

- Develop advanced semi-supervised and self-supervised learning techniques
- Investigate few-shot meta-learning for rapid adaptation to novel fraud
- Create privacy-preserving synthetic data generation methods
- Explore transfer learning across fraud types and institutions
- **Expected Impact:** 10-20% performance improvement in low-data regimes

2. Handling Concept Drift:

- Design continual learning algorithms preventing catastrophic forgetting
- Develop robust drift detection methods accommodating label delay
- Investigate online learning approaches for real-time adaptation
- Study long-term model evolution (multi-year performance)
- **Expected Impact:** Sustained high performance over extended periods

3. Enhancing Interpretability:

- Develop fraud-specific explanation methods beyond generic XAI techniques
- Investigate causal inference for robust, interpretable fraud models
- Create evaluation frameworks for explanation quality
- Study human-AI collaboration in fraud investigation workflows
- **Expected Impact:** Improved analyst efficiency (30-50%) and regulatory compliance

4. Privacy-Preserving Techniques:

- Advance federated learning for non-IID fraud data
- Minimize privacy-utility trade-off in differential privacy
- Develop secure multi-party computation for fraud detection
- Investigate homomorphic encryption for practical applications
- **Expected Impact:** Enable cross-institutional collaboration, ¡5% accuracy loss

5. Domain-Specific Innovations:

- Supply chain: Graph-based temporal fraud detection
- E-commerce: Behavioral biometrics and session analysis
- Insurance: Multi-modal analysis integrating claims text and images
- Banking: Real-time streaming fraud detection at scale
- **Expected Impact:** 15-35% domain-specific performance improvements

13.2.2. Methodological Best Practices

To advance the field rigorously, researchers should:
Experimental Design:

- Use temporal train-test splits, not random splits (realistic evaluation)
- Report multiple metrics: AUC-ROC, AUC-PR, F1, precision, recall, confusion matrices
- Emphasize AUC-PR over AUC-ROC for extreme imbalance
- Conduct statistical significance testing with confidence intervals
- Perform ablation studies isolating component contributions

Reproducibility:

- Publish code and detailed hyperparameters
- Use multiple random seeds and report mean/std performance
- Document all preprocessing steps explicitly
- Specify exact dataset versions and train-test splits
- Share pre-trained models when possible

Benchmarking:

- Evaluate on multiple datasets (European, IEEE-CIS, domain-specific)
- Compare against strong baselines (XGBoost, LightGBM, not just naive methods)
- Report computational costs (training time, inference latency, memory)
- Include failure case analysis
- Discuss practical deployment considerations

Collaboration:

- Partner with industry for realistic problem formulations
- Validate on proprietary data when possible (report aggregated results)
- Engage with fraud analysts to understand practical needs
- Consider regulatory and ethical implications proactively
- Contribute to open-source fraud detection libraries

13.3. Recommendations for Practitioners

13.3.1. Algorithm Selection Guidelines

For practitioners deploying fraud detection systems:

Starting Point - Default Recommendation:

- **Primary:** XGBoost or LightGBM with Borderline-SMOTE
- **Rationale:** Best accuracy-complexity-speed balance
- **Performance:** 90-92% AUC-ROC on most fraud datasets
- **Development Time:** 3-4 weeks for initial deployment

When to Upgrade to Ensemble Stacking:

- Fraud losses exceed $10M annually (2-5% improvement justifies complexity)
- Ample development resources available
- Production infrastructure supports ensemble serving
- 6-10 week development timeline acceptable

When to Use Deep Learning:

- Sequential patterns critical (transaction sequences, clickstreams)
- Abundant labeled data (¿100K fraud examples)
- GPU infrastructure available
- Interpretability less critical
- Example: LSTM for credit card transaction sequences

When to Use Semi-Supervised Approaches:

- Labeled fraud examples scarce (¡10K)
- Labeling expensive or time-consuming (supply chain, insurance)

- Two-phase framework: Isolation Forest + Self-Training SVM
- Expected performance: F1 0.80-0.85 with ¡5% labeled data

When Interpretability is Critical:

- Regulatory requirements (GDPR Article 22, FCRA)
- Customer-facing decisions (transaction declines)
- Primary: Random Forest (feature importance transparent)
- Alternative: XGBoost with SHAP explanations
- Accept 2-3% accuracy sacrifice for interpretability

13.3.2. Implementation Best Practices

1. Feature Engineering:

- **Temporal features:** Transaction velocity, recency, frequency patterns
- **Aggregations:** Customer spending patterns over multiple time windows (1h, 24h, 7d, 30d)
- **Deviations:** Distance from historical behavior (amount, location, time)
- **Network features:** Merchant risk scores, user-merchant relationship patterns
- **Impact:** Well-engineered features often more important than algorithm choice

2. Handling Class Imbalance:

- Apply Borderline-SMOTE for most scenarios (outperforms standard SMOTE and ADASYN)
- Use stratified splitting to ensure fraud representation in train and test sets
- Tune classification thresholds based on business costs (investigation vs. fraud loss)
- Monitor precision-recall trade-off continuously

3. Model Training and Validation:

- Use temporal train-test splits (train on past, test on future)
- Validate on recent data to assess production performance
- Implement time-series cross-validation for robust performance estimates

- Reserve recent data for final evaluation (avoid overfitting to validation set)

4. Production Deployment:

- Start with A/B testing (5-10% of traffic) before full rollout
- Monitor key metrics continuously: AUC-ROC, precision, recall, false positive rate
- Implement automated alerts for performance degradation
- Maintain rollback capabilities for rapid model reversion
- Document model versions, training data, and hyperparameters

5. Handling Concept Drift:

- Retrain models weekly or monthly depending on drift rate
- Monitor performance on recent data to detect drift early
- Consider online learning for continuous adaptation
- Maintain ensemble of models trained on different time periods
- Investigate continual learning as research matures

6. Operational Considerations:

- Balance automated decisions vs. human review based on confidence thresholds
- Integrate with fraud analyst workflows (provide explanations, investigation support)
- Collect feedback from analysts to improve models
- Measure and optimize investigation workload
- Calculate ROI: Fraud losses prevented vs. system costs

13.3.3. Infrastructure Requirements

Minimal Infrastructure (Small-Medium Organizations):

- Cloud-based ML platform (AWS SageMaker, Google Vertex AI, Azure ML)
- Random Forest or XGBoost (CPU-only inference adequate)
- Batch processing acceptable (hourly or daily)
- Basic monitoring and alerting

- Cost: \$5K-20K/year infrastructure

Production Infrastructure (Large Organizations):

- Real-time feature computation (Apache Kafka, Flink)
- Low-latency model serving (¡50ms, TensorFlow Serving, MLflow)
- GPU acceleration for deep learning models
- Comprehensive MLOps: Versioning, monitoring, A/B testing, automated retraining
- Feature store for consistent training-serving features
- Cost: \$100K-500K/year infrastructure

Advanced Infrastructure (Cutting-Edge Organizations):

- Edge computing for ultra-low latency (¡10ms)
- Federated learning for cross-institutional collaboration
- AutoML for continuous model optimization
- Real-time concept drift detection and adaptation
- Advanced explainability dashboards for analysts
- Cost: \$500K-2M/year infrastructure

13.4. Broader Impact and Ethical Considerations

13.4.1. Societal Benefits

Effective fraud detection systems provide substantial societal benefits:

- **Financial Loss Prevention:** Billions of dollars in fraud losses prevented annually
- **Consumer Protection:** Individuals protected from identity theft and financial fraud
- **Market Integrity:** Fair marketplaces free from fraudulent actors
- **Trust Enhancement:** Increased confidence in digital payment systems
- **Economic Efficiency:** Resources allocated productively rather than lost to fraud

13.4.2. Ethical Concerns and Mitigation

Fraud detection systems also raise ethical concerns requiring careful attention:

1. Fairness and Bias:

- **Concern:** Discriminatory fraud detection harming protected groups
- **Mitigation:** Regular bias audits, fairness constraints during training, demographic parity monitoring
- **Challenge:** Legitimate fraud rate differences across demographics complicate fairness definitions

2. Privacy:

- **Concern:** Extensive data collection and behavioral monitoring
- **Mitigation:** Data minimization, differential privacy, federated learning, transparent data practices
- **Challenge:** Balancing fraud prevention effectiveness against privacy preservation

3. False Positives:

- **Concern:** Legitimate customers inconvenienced or harmed by false fraud alerts
- **Mitigation:** Careful threshold tuning, human review for high-impact decisions, transparent appeals process
- **Challenge:** Optimizing precision-recall trade-off varies by context

4. Transparency and Accountability:

- **Concern:** Black-box models making consequential decisions without explanation
- **Mitigation:** Explainable AI techniques, interpretable models for high-stakes decisions, clear accountability
- **Challenge:** Accuracy-interpretability trade-off

5. Adversarial Arms Race:

- **Concern:** Increasingly sophisticated fraud detection drives more sophisticated fraud
- **Mitigation:** Responsible disclosure practices, security through obscurity where appropriate, defensive design
- **Challenge:** Balancing research openness against security concerns

13.4.3. Responsible AI Development

We advocate for responsible AI development in fraud detection:

- **Fairness:** Proactive bias detection and mitigation throughout development lifecycle
- **Privacy:** Privacy-by-design principles, minimizing data collection and retention
- **Transparency:** Clear communication about fraud detection systems to customers
- **Accountability:** Clear lines of responsibility for automated decisions
- **Human Oversight:** Maintaining meaningful human control over high-stakes decisions
- **Security:** Robust defenses against adversarial attacks and model extraction
- **Continuous Monitoring:** Ongoing evaluation of fairness, accuracy, and societal impact

13.5. Limitations of This Survey

This survey has several limitations:

- **Public Data Focus:** Analysis primarily based on public benchmarks; proprietary industry data and systems remain largely undocumented
- **Snapshot in Time:** Fraud detection evolves rapidly; findings reflect state-of-the-art as of 2025
- **Domain Coverage:** Primary focus on credit card, e-commerce, and supply chain fraud; other domains (healthcare fraud, tax fraud) less thoroughly covered
- **Deployment Details:** Limited information on production system architectures and operational practices due to proprietary nature
- **Geographic Scope:** Research primarily from North America, Europe, and East Asia; fraud patterns in other regions less represented

13.6. Future Outlook

The fraud detection field stands at an inflection point with several converging trends:

Technological Advances:

- Ensemble methods and gradient boosting have matured, achieving 90-95% AUC-ROC consistently
- Deep learning demonstrates promise for sequential and multi-modal data
- Emerging technologies (federated learning, GNNs, XAI) transitioning from research to practice
- AutoML beginning to democratize advanced fraud detection

Methodological Evolution:

- Shift from purely supervised to semi-supervised and self-supervised learning
- Growing emphasis on interpretability and fairness alongside accuracy
- Recognition that concept drift adaptation is critical for sustained performance
- Increasing focus on privacy-preserving techniques

Industry Maturation:

- Fraud detection transitioning from artisanal to systematic engineering discipline
- MLOps practices enabling rapid iteration and deployment
- Cross-institutional collaboration becoming feasible through federated learning
- Regulatory frameworks evolving to address AI-driven fraud detection

Remaining Challenges:

- Concept drift remains incompletely solved; continual learning underexplored
- Privacy-utility trade-off still significant (5-15% accuracy loss)
- Limited public benchmarks hinder reproducible research
- Adversarial robustness requires continued attention as fraud sophistication increases

13.7. Closing Remarks

Fraud detection represents a critical application of machine learning with substantial societal and economic impact. This survey has provided a comprehensive analysis of the field, synthesizing findings from traditional machine learning, deep learning, ensemble methods, semi-supervised learning, and emerging technologies across financial and supply chain domains.

Several key insights emerge:

1. Ensemble Methods Dominate: XGBoost, LightGBM, and stacking ensembles consistently achieve state-of-the-art performance across diverse fraud detection tasks, representing the current best practice for most applications.

2. Context Matters: No universally optimal approach exists; algorithm selection depends on data characteristics, latency requirements, interpretability needs, and resource constraints. Practitioners must carefully match methods to their specific context.

3. Challenges Persist: Despite significant progress, extreme class imbalance, concept drift, data scarcity, and the interpretability-accuracy trade-off remain fundamental challenges requiring continued research attention.

4. Emerging Technologies Promise Further Advances: Federated learning, graph neural networks, explainable AI, and continual learning show promise for addressing current limitations and enabling next-generation fraud detection systems.

5. Responsible Development is Critical: As fraud detection systems become more powerful and pervasive, careful attention to fairness, privacy, transparency, and accountability becomes increasingly important.

The field continues to evolve rapidly, driven by advancing machine learning techniques, increasing data availability, and escalating fraud sophistication. Success requires collaboration among researchers advancing the state-of-the-art, practitioners deploying robust systems, and policymakers creating appropriate regulatory frameworks. By combining technical innovation with responsible development practices, the fraud detection community can build systems that effectively prevent fraud while respecting privacy, fairness, and human values.

This survey provides a foundation for researchers and practitioners navigating the complex landscape of fraud detection, offering both a comprehensive view of current capabilities and a roadmap for future advances. As fraud tactics continue to evolve, so too must our detection methods, demanding sustained innovation, rigorous evaluation, and thoughtful deployment of machine learning techniques in service of more secure and trustworthy financial and commercial systems.

Acknowledgments

We thank the fraud detection research community for their continued efforts advancing the field through rigorous research, open-source contributions, and thoughtful discussions of technical and ethical challenges. We acknowledge the financial institutions and technology companies that have shared datasets, insights, and best practices enabling academic research. Finally, we recognize the

fraud analysts and security professionals whose domain expertise and practical experience ground theoretical advances in operational reality.

14. References

References

[1] P. Rebentrost, M. Mohseni, and S. Lloyd, "Quantum support vector machine for big data classification," *Physical Review Letters*, vol. 113, no. 13, p. 130503, 2014.

[2] E. Farhi and H. Neven, "Classification with quantum neural networks on near term processors," *arXiv preprint arXiv:1802.06002*, 2018.

[3] S. Lloyd, M. Mohseni, and P. Rebentrost, "Quantum principal component analysis," *Nature Physics*, vol. 10, no. 9, pp. 631–633, 2014.

[4] Q. Yang, Y. Liu, T. Chen, and Y. Tong, "Federated machine learning: Concept and applications," *ACM Transactions on Intelligent Systems and Technology*, vol. 10, no. 2, pp. 1–19, 2019.

[5] K. Wei, J. Li, M. Ding, C. Ma, H. H. Yang, F. Farokhi, S. Jin, T. Q. S. Quek, and H. V. Poor, "Federated learning with differential privacy: Algorithms and performance analysis," *IEEE Transactions on Information Forensics and Security*, vol. 15, pp. 3454–3469, 2020.

[6] P. Veličković, G. Cucurull, A. Casanova, A. Romero, P. Liò, and Y. Bengio, "Graph attention networks," *arXiv preprint arXiv:1710.10903*, 2018.

[7] H. Zeng, H. Zhou, A. Srivastava, R. Kannan, and V. Prasanna, "GraphSAINT: Graph sampling based inductive learning method," in *International Conference on Learning Representations*, 2020.

[8] D. Wang, J. Lin, P. Cui, Q. Jia, Z. Wang, Y. Fang, Q. Yu, J. Zhou, S. Yang, and Y. Qi, "A semi-supervised graph attentive network for financial fraud detection," in *2019 IEEE International Conference on Data Mining (ICDM)*, IEEE, 2019, pp. 598–607.

[9] S. M. Lundberg and S.-I. Lee, "A unified approach to interpreting model predictions," in *Advances in Neural Information Processing Systems*, 2017, pp. 4765–4774.

[10] M. T. Ribeiro, S. Singh, and C. Guestrin, "Why should I trust you?: Explaining the predictions of any classifier," in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016, pp. 1135–1144.

[11] M. T. Ribeiro, S. Singh, and C. Guestrin, "Anchors: High-precision model-agnostic explanations," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 32, 2018.

[12] M. N. Hasan, R. N. Toma, A. A. Nahid, M. M. M. Islam, and J. M. Kim, "Explainable AI for fraud detection in financial transactions," *IEEE Access*, vol. 11, pp. 23549–23567, 2023.

[13] C. Rudin, "Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead," *Nature Machine Intelligence*, vol. 1, no. 5, pp. 206–215, 2019.

[14] J. Kirkpatrick, R. Pascanu, N. Rabinowitz, J. Veness, G. Desjardins, A. A. Rusu, K. Milan, J. Quan, T. Ramalho, A. Grabska-Barwinska et al., "Overcoming catastrophic forgetting in neural networks," *Proceedings of the National Academy of Sciences*, vol. 114, no. 13, pp. 3521–3526, 2017.

[15] J. Lu, A. Liu, F. Dong, F. Gu, J. Gama, and G. Zhang, "Learning under concept drift: A review," *IEEE Transactions on Knowledge and Data Engineering*, vol. 31, no. 12, pp. 2346–2363, 2018.

[16] C. Finn, P. Abbeel, and S. Levine, "Model-agnostic meta-learning for fast adaptation of deep networks," in *International Conference on Machine Learning*, PMLR, 2017, pp. 1126–1135.

[17] Q. Feng, D. He, S. Zeadally, M. K. Khan, and N. Kumar, "A survey on privacy protection in blockchain system," *Journal of Network and Computer Applications*, vol. 126, pp. 45–58, 2019.

[18] F. Moradi, M. Tarif, and M. Homaei, "A systematic review of machine learning in credit card fraud detection," *Preprint, MDPI AG*, July 2025.

[19] N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer, "SMOTE: synthetic minority over-sampling technique," *Journal of Artificial Intelligence Research*, vol. 16, pp. 321–357, 2002.

[20] H. He, Y. Bai, E. A. Garcia, and S. Li, "ADASYN: Adaptive synthetic sampling approach for imbalanced learning," in *2008 IEEE International Joint Conference on Neural Networks*, IEEE, 2008, pp. 1322–1328.

[21] H. Han, W.-Y. Wang, and B.-H. Mao, "Borderline-SMOTE: a new over-sampling method in imbalanced data sets learning," in *International Conference on Intelligent Computing*, Springer, 2005, pp. 878–887.

[22] T.-Y. Lin, P. Goyal, R. Girshick, K. He, and P. Dollár, "Focal loss for dense object detection," in *Proceedings of the IEEE International Conference on Computer Vision*, 2017, pp. 2980–2988.

[23] Y. Lucas, P. Jurgovsky, M. Vaněk, P. Grançay, M. Kaytoue, R. Soule, and L. He-Guelton, "Credit card fraud detection using machine learning: A survey," *ACM Computing Surveys*, vol. 53, no. 4, pp. 1–38, 2020.

[24] F. Moradi, M. Tarif, and M. Homaei, "Semi-supervised supply chain fraud detection with unsupervised pre-filtering," *arXiv preprint arXiv:2508.06574*, 2025.

[25] F. Moradi, M. Tarif, and M. Homaei, "Robust fraud detection with ensemble learning: A case study on the IEEE-CIS dataset," *Preprint*, July 2025.

[26] F. Moradi, M. Tarif, and M. Homaei, "Ensemble-based fraud detection: A robust approach evaluated on IEEE-CIS," *Preprints.org*, 2025.

[27] S. Bhattacharyya, S. Jha, K. Tharakunnel, and J. C. Westland, "Data mining for credit card fraud: A comparative study," *Decision Support Systems*, vol. 50, no. 3, pp. 602–613, 2011.

[28] J. West and M. Bhattacharya, "Intelligent financial fraud detection: A comprehensive review," *Computers & Security*, vol. 57, pp. 47–66, 2016.

[29] C. Phua, V. Lee, K. Smith, and R. Gayler, "A comprehensive survey of data mining-based fraud detection research," *arXiv preprint arXiv:1009.6119*, 2010.

[30] M. Rtayli and N. Enneya, "Enhanced credit card fraud detection based on SVM-recursive feature elimination and hyper-parameters optimization," *Journal of Information Security and Applications*, vol. 55, p. 102596, 2021.

[31] N. S. Halvaiee and M. K. Akbari, "A novel model for credit card fraud detection using artificial immune systems," *Applied Soft Computing*, vol. 24, pp. 40–49, 2014.

[32] V. Chandola, A. Banerjee, and V. Kumar, “Anomaly detection: A survey,” *ACM Computing Surveys*, vol. 41, no. 3, pp. 1–58, 2009.

[33] W. Jiang, “A survey of fraud detection in financial transactions,” *International Journal of Computer Applications*, vol. 182, no. 17, pp. 1–5, 2018.

[34] A. Abdallah, M. A. Maarof, and A. Zainal, “Fraud detection system: A survey,” *Journal of Network and Computer Applications*, vol. 68, pp. 90–113, 2016.

[35] X. Zhang, Y. Han, W. Xu, and Q. Wang, “HOBA: A novel feature engineering methodology for credit card fraud detection with a deep learning architecture,” *Information Sciences*, vol. 557, pp. 302–316, 2021.

[36] K. Fu, D. Cheng, Y. Tu, and L. Zhang, “Credit card fraud detection using convolutional neural networks,” in *International Conference on Neural Information Processing*, Springer, 2016, pp. 483–490.

[37] J. Jurgovsky, M. Granitzer, K. Ziegler, S. Calabretto, P. E. Portier, L. He-Guelton, and O. Caelen, “Sequence classification for credit-card fraud detection,” *Expert Systems with Applications*, vol. 100, pp. 234–245, 2018.

[38] A. Pumsirirat and L. Yan, “Credit card fraud detection using deep learning based on auto-encoder and restricted Boltzmann machine,” *International Journal of Advanced Computer Science and Applications*, vol. 9, no. 1, pp. 18–25, 2018.

[39] M. Schreyer, T. Sattarov, D. Borth, A. Dengel, and B. Reimer, “Detection of anomalies in large scale accounting data using deep autoencoder networks,” *arXiv preprint arXiv:1709.05254*, 2017.

[40] U. Fiore, A. De Santis, F. Perla, P. Zanetti, and F. Palmieri, “Using generative adversarial networks for improving classification effectiveness in credit card fraud detection,” *Information Sciences*, vol. 479, pp. 448–455, 2019.

[41] T. Pourhabibi, K. L. Ong, B. H. Kam, and Y. L. Boo, “Fraud detection: A systematic literature review of graph-based anomaly detection approaches,” *Decision Support Systems*, vol. 133, p. 113303, 2020.

[42] D. Wang, “A survey on applications of deep learning in credit card fraud detection,” *International Journal of Machine Learning and Computing*, vol. 11, no. 2, pp. 101–108, 2021.

[43] Z. Xu, Y. Zhang, and C. K. Leung, “Credit card fraud detection using deep learning: A review,” *IEEE Access*, vol. 9, pp. 93257–93273, 2021.

[44] M. Zareapoor and P. Shamsolmoali, “Application of credit card fraud detection: Based on bagging ensemble classifier,” *Procedia Computer Science*, vol. 48, pp. 679–685, 2015.

[45] B. Lebichot, Y.-A. Le Borgne, L. He-Guelton, F. Oblé, and G. Bontempi, “Deep-learning domain adaptation techniques for credit cards fraud detection,” in *INNSBDDL 2019: Recent Advances in Big Data and Deep Learning*, Springer, 2021, pp. 78–88.

[46] T. Chen and C. Guestrin, “XGBoost: A scalable tree boosting system,” in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016, pp. 785–794.

[47] A. C. Van Vlasselaer, C. Bravo, O. Caelen, T. Eliassi-Rad, L. Akoglu, M. Snoeck, and B. Baesens, “APATE: A novel approach for automated credit card transaction fraud detection using network-based extensions,” *Decision Support Systems*, vol. 75, pp. 38–48, 2015.

[48] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T. Y. Liu, “LightGBM: A highly efficient gradient boosting decision tree,” in *Advances in Neural Information Processing Systems*, 2017, pp. 3146–3154.

[49] L. Prokhorenkova, G. Gusev, A. Vorobev, A. V. Dorogush, and A. Gulin, "CatBoost: unbiased boosting with categorical features," in *Advances in Neural Information Processing Systems*, 2018, pp. 6638–6648.

[50] X. Dastile, T. Celik, and M. Potsane, "Statistical and machine learning models in credit scoring: A systematic literature survey," *Applied Soft Computing*, vol. 91, p. 106263, 2020.

[51] T. N. Tran, T. A. Nguyen, T. H. Nguyen, and D. T. Nguyen, "Real-time fraud detection using ensemble learning," in *2018 International Conference on Advanced Technologies for Communications (ATC)*, IEEE, 2018, pp. 1–6.

[52] A. Dal Pozzolo, O. Caelen, R. A. Johnson, and G. Bontempi, "Calibrating probability with undersampling for unbalanced classification," in *Proc. IEEE Symposium on Computational Intelligence and Data Mining (CIDM)*, 2015, pp. 159–166.

[53] F. T. Liu, K. M. Ting, and Z.-H. Zhou, "Isolation forest," in *2008 Eighth IEEE International Conference on Data Mining*, IEEE, 2008, pp. 413–422.

[54] D. M. J. Tax and R. P. W. Duin, "Support vector data description," *Machine Learning*, vol. 54, no. 1, pp. 45–66, 2004.

[55] M. M. Breunig, H.-P. Kriegel, R. T. Ng, and J. Sander, "LOF: identifying density-based local outliers," in *ACM SIGMOD Record*, vol. 29, ACM, 2000, pp. 93–104.

[56] Y. Kou, C.-T. Lu, S. Sirwongwattana, and Y.-P. Huang, "Survey of fraud detection techniques," in *IEEE International Conference on Networking, Sensing and Control*, vol. 2, IEEE, 2004, pp. 749–754.

[57] A. Dal Pozzolo, O. Caelen, Y.-A. Le Borgne, S. Waterschoot, and G. Bontempi, "Learned lessons in credit card fraud detection from a practitioner perspective," *Expert Systems with Applications*, vol. 41, no. 10, pp. 4915–4928, 2014.

[58] A. Van Vlasselaer, C. Bravo, O. Caelen, T. Eliassi-Rad, L. Akoglu, M. Snoeck, and B. Baesens, "APATE: A novel approach for automated credit card transaction fraud detection using network-based extensions," *Decision Support Systems*, vol. 75, pp. 38–48, 2015.

[59] Y. Li, J. T. Kwok, and Z. H. Zhou, "Semi-supervised learning using label mean," in *Proceedings of the 26th Annual International Conference on Machine Learning*, 2009, pp. 633–640.

[60] L. Akoglu, H. Tong, and D. Koutra, "Graph based anomaly detection and description: a survey," *Data Mining and Knowledge Discovery*, vol. 29, no. 3, pp. 626–688, 2015.

[61] S. Jain, M. White, and P. Radivojac, "Recovering true classifier performance in positive-unlabeled learning," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 31, 2017.

[62] S. Das, W.-K. Wong, T. Dietterich, A. Fern, and A. Emmott, "Incorporating expert feedback into active anomaly discovery," in *2016 IEEE 16th International Conference on Data Mining (ICDM)*, IEEE, 2016, pp. 853–858.

[63] H. Abroshan, J. Devos, G. Poels, and E. Laermans, "Phishing happens beyond technology: The effects of human behaviors and demographics on each step of a phishing attack," *IEEE Access*, vol. 9, pp. 44044–44072, 2021.

[64] L. Šubelj, S. Furlan, and M. Bajec, "An expert system for detecting automobile insurance fraud using social network analysis," *Expert Systems with Applications*, vol. 38, no. 1, pp. 1039–1052, 2011.

[65] A. Van Vlasselaer, J. Eliassi-Rad, L. Akoglu, M. Snoeck, and B. Baesens, "Gotcha! Network-based fraud detection for social security fraud," *Management Science*, vol. 63, no. 9, pp. 3090–3110, 2017.

[66] R. A. Bauder and T. M. Khoshgoftaar, "Medicare fraud detection using machine learning methods," in *2017 16th IEEE International Conference on Machine Learning and Applications (ICMLA)*, IEEE, 2017, pp. 858–865.

[67] C. S. Hilas, "Designing an expert system for fraud detection in private telecommunications networks," *Expert Systems with Applications*, vol. 36, no. 9, pp. 11559–11569, 2009.

[68] Y. Wang, S. Wei, and Z. Wang, "Session-based fraud detection in online e-commerce transactions using recurrent neural networks," in *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*, Springer, 2018, pp. 241–252.

[69] L. Zheng, G. Liu, C. Yan, C. Jiang, M. Zhou, and M. Li, "Improved TrAdaBoost and its application to transaction fraud detection," *IEEE Transactions on Computational Social Systems*, vol. 7, no. 5, pp. 1304–1316, 2020.

[70] F. Tian, "An agri-food supply chain traceability system for China based on RFID & blockchain technology," in *2016 13th International Conference on Service Systems and Service Management (ICSSSM)*, IEEE, 2016, pp. 1–6.

[71] M. S. Islam and N. H. Rony, "Credit card fraud detection using machine learning and deep learning techniques," in *2020 23rd International Conference on Computer and Information Technology (ICCIT)*, IEEE, 2020, pp. 1–6.

[72] V. López, A. Fernández, S. García, V. Palade, and F. Herrera, "An insight into classification with imbalanced data: Empirical results and current trends on using data intrinsic characteristics," *Information Sciences*, vol. 250, pp. 113–141, 2013.

[73] M. Galar, A. Fernandez, E. Barrenechea, H. Bustince, and F. Herrera, "A review on ensembles for the class imbalance problem: bagging-, boosting-, and hybrid-based approaches," *IEEE Transactions on Systems, Man, and Cybernetics, Part C*, vol. 42, no. 4, pp. 463–484, 2011.

[74] L. A. Jeni, J. F. Cohn, and F. De La Torre, "Facing imbalanced data–recommendations for the use of performance metrics," in *2013 Humaine Association Conference on Affective Computing and Intelligent Interaction*, IEEE, 2013, pp. 245–251.

[75] D. Bahdanau, K. Cho, and Y. Bengio, "Neural machine translation by jointly learning to align and translate," *arXiv preprint arXiv:1409.0473*, 2014.

[76] N. Barakat and A. P. Bradley, "Rule extraction from support vector machines: A review," *Neurocomputing*, vol. 74, no. 1-3, pp. 178–190, 2010.

[77] S. Wachter, B. Mittelstadt, and C. Russell, "Counterfactual explanations without opening the black box: Automated decisions and the GDPR," *Harv. JL & Tech.*, vol. 31, p. 841, 2017.

[78] J. Bien and R. Tibshirani, "Prototype selection for interpretable classification," *The Annals of Applied Statistics*, vol. 5, no. 4, pp. 2403–2424, 2011.

[79] L. L. Minku, A. P. White, and X. Yao, "The impact of diversity on online ensemble learning in the presence of concept drift," *IEEE Transactions on Knowledge and Data Engineering*, vol. 22, no. 5, pp. 730–742, 2010.

[80] J. Gama, I. Žliobaitė, A. Bifet, M. Pechenizkiy, and A. Bouchachia, "A survey on concept drift adaptation," *ACM Computing Surveys*, vol. 46, no. 4, pp. 1–37, 2014.

[81] V. Losing, B. Hammer, and H. Wersing, "Incremental on-line learning: A review and comparison of state of the art algorithms," *Neurocomputing*, vol. 275, pp. 1261–1274, 2018.

[82] J. Gama, P. Medas, G. Castillo, and P. Rodrigues, "Learning with drift detection," in *Brazilian Symposium on Artificial Intelligence*, Springer, 2004, pp. 286–295.

[83] M. Baena-Garcia, J. del Campo-Ávila, R. Fidalgo, A. Bifet, R. Gavalda, and R. Morales-Bueno, "Early drift detection method," in *Fourth International Workshop on Knowledge Discovery from Data Streams*, 2006, vol. 6, pp. 77–86.

[84] R. Elwell and R. Polikar, "Incremental learning of concept drift in nonstationary environments," *IEEE Transactions on Neural Networks*, vol. 22, no. 10, pp. 1517–1531, 2011.

[85] C. Dwork and A. Roth, "The algorithmic foundations of differential privacy," *Foundations and Trends in Theoretical Computer Science*, vol. 9, no. 3–4, pp. 211–407, 2014.

[86] A. Acar, H. Aksu, A. S. Uluagac, and M. Conti, "A survey on homomorphic encryption schemes: Theory and implementation," *ACM Computing Surveys*, vol. 51, no. 4, pp. 1–35, 2018.

[87] P. Mohassel and Y. Zhang, "SecureML: A system for scalable privacy-preserving machine learning," in *2017 IEEE Symposium on Security and Privacy (SP)*, IEEE, 2017, pp. 19–38.

[88] J. Jordon, J. Yoon, and M. van der Schaar, "PATE-GAN: Generating synthetic data with differential privacy guarantees," in *International Conference on Learning Representations*, 2019.

[89] M. J. Page, J. E. McKenzie, P. M. Bossuyt, I. Boutron, T. C. Hoffmann, C. D. Mulrow, L. Shamseer, J. M. Tetzlaff, E. A. Akl, S. E. Brennan et al., "The PRISMA 2020 statement: an updated guideline for reporting systematic reviews," *BMJ*, vol. 372, 2021.

[90] R. M. Cruz, L. G. Hafemann, R. Sabourin, and G. D. Cavalcanti, "Measuring the quality of research in machine learning," *Journal of Machine Learning Research*, vol. 22, no. 1, pp. 1–48, 2021.

[91] E. W. T. Ngai, Y. Hu, Y. H. Wong, Y. Chen, and X. Sun, "The application of data mining techniques in financial fraud detection: A classification framework and an academic review of literature," *Decision Support Systems*, vol. 50, no. 3, pp. 559–569, 2011.

[92] A. Dal Pozzolo, O. Caelen, Y.-A. Le Borgne, S. Waterschoot, and G. Bontempi, "Learned lessons in credit card fraud detection from a practitioner perspective," *Expert Systems with Applications*, vol. 41, no. 10, pp. 4915–4928, 2014.

[93] V. López, A. Fernández, J. G. Moreno-Torres, and F. Herrera, "Analysis of preprocessing vs. cost-sensitive learning for imbalanced classification. Open problems on intrinsic data characteristics," *Expert Systems with Applications*, vol. 39, no. 7, pp. 6585–6608, 2012.

[94] A. Andoni and P. Indyk, "Near-optimal hashing algorithms for approximate nearest neighbor in high dimensions," in *47th Annual IEEE Symposium on Foundations of Computer Science (FOCS'06)*, IEEE, 2006, pp. 459–468.

[95] A. Roy, J. Sun, R. Mahoney, L. Alonzi, S. Adams, and P. Beling, "Deep learning detecting fraud in credit card transactions," in *2018 Systems and Information Engineering Design Symposium (SIEDS)*, IEEE, 2018, pp. 129–134.

[96] F. Itoo, M. Meenakshi, and S. Singh, "Comparison and analysis of logistic regression, Naïve Bayes and KNN machine learning algorithms for credit card fraud detection," *International Journal of Information Technology*, vol. 13, no. 4, pp. 1503–1511, 2021.

[97] A. C. Bahnsen, D. Aouada, A. Stojanovic, and B. Ottersten, "Feature engineering strategies for credit card fraud detection," *Expert Systems with Applications*, vol. 51, pp. 134–142, 2016.

[98] X. Zhang, Y. Han, W. Xu, and Q. Wang, "HOBA: A novel feature engineering methodology for credit card fraud detection with a deep learning architecture," *Information Sciences*, vol. 557, pp. 302–316, 2021.

[99] L. Zheng, G. Liu, C. Yan, C. Jiang, M. Zhou, and M. Li, "Improved TrAdaBoost and its application to transaction fraud detection," *IEEE Transactions on Computational Social Systems*, vol. 7, no. 5, pp. 1304–1316, 2020.

[100] S. Ando and C. Y. Huang, "Deep over-sampling framework for classifying imbalanced data," in *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*, Springer, 2018, pp. 770–785.

[101] Z. Kazemi and H. Zarrabi, "Using deep networks for fraud detection in the credit card transactions," in *2017 IEEE 4th International Conference on Knowledge-Based Engineering and Innovation (KBEI)*, IEEE, 2017, pp. 0630–0633.

[102] D. Cheng, S. Xiang, C. Shang, Y. Zhang, F. Yang, and L. Zhang, "Spatio-temporal attention-based neural network for credit card fraud detection," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 34, 2020, pp. 362–369.

[103] K. Randhawa, C. K. Loo, M. Seera, C. P. Lim, and A. K. Nandi, "Credit card fraud detection using AdaBoost and majority voting," *IEEE Access*, vol. 6, pp. 14277–14284, 2018.

[104] C. Seiffert, T. M. Khoshgoftaar, J. Van Hulse, and A. Napolitano, "RUSBoost: A hybrid approach to alleviating class imbalance," *IEEE Transactions on Systems, Man, and Cybernetics-Part A: Systems and Humans*, vol. 40, no. 1, pp. 185–197, 2010.

[105] M. Ahmed, A. N. Mahmood, and J. Hu, "A survey of network anomaly detection techniques," *Journal of Network and Computer Applications*, vol. 60, pp. 19–31, 2016.

[106] M. Ester, H.-P. Kriegel, J. Sander, and X. Xu, "A density-based algorithm for discovering clusters in large spatial databases with noise," in *Proceedings of the Second International Conference on Knowledge Discovery and Data Mining (KDD-96)*, 1996, pp. 226–231.

[107] J. An and S. Cho, "Variational autoencoder based anomaly detection using reconstruction probability," *Special Lecture on IE*, vol. 2, no. 1, pp. 1–18, 2015.

[108] A. Van Vlasselaer, C. Bravo, O. Caelen, T. Eliassi-Rad, L. Akoglu, M. Snoeck, and B. Baesens, "APATE: A novel approach for automated credit card transaction fraud detection using network-based extensions," *Decision Support Systems*, vol. 75, pp. 38–48, 2015.